

Faculty of Informatics
Engineering Department of
Software Engineering &
Information System

كلية هندسة المعلوماتية
قسم هندسة البرمجيات و نظم
المعلومات

AutoTest & DocGen Manager – Intelligent System for Automated Testing and Documentation

Prepared by

Areej nooraldeen Kamar aldiab Ola najibeh Wiaam alouni

Supervised by

Dr. Riad Sonbol Eng.Raghad Alhossny

2025-2026

SUPERVISION CERTIFICATION

I certify that the preparation of this project entitled

.....,prepared by

.....

.....,was made under my supervision at
Department of Software Engineering & Information System / Faculty of
Informatics Engineering in partial fulfillment of the Requirements for the
Degree of Bachelors of Software Engineering & Information System /
Faculty of Informatics Engineering

Name: Signature: Date:

إقرار المشرف

أشهد بأن هذا المشروع الموسوم.....

.....

والمعد من قبل الطلاب.....

.....

قد تم تحت إشرافي في قسم هندسة البرمجيات ونظم المعلومات / كلية هندسة المعلوماتية، وهو جزء من
متطلبات نيل شهادة الإجازة في هندسة البرمجيات ونظم المعلومات.

الاسم: التوقيع: التاريخ:

ABSTRACT

Software documentation and testing are fundamental pillars of maintainable, high-quality code.

However, modern development teams face persistent challenges: documentation becomes outdated within months of creation, critical APIs and algorithms remain undocumented, and developers spend up to 40% of their time on documentation-related tasks rather than core development activities.

This thesis presents AutoTest & DocGen Manager, an intelligent web-based platform that automates software testing and documentation generation end-to-end using advanced artificial intelligence techniques. The system is built on a novel multi-agent architecture comprising three specialized agents: the HighLevelAgent, which produces executive-level documentation describing overall system purpose and key capabilities; the LowLevelAgent, which performs deep technical analysis of methods, class relationships, design patterns, and architectural recommendations; and the VerifierAgent, which implements a self-correction mechanism to validate output accuracy against the actual code.

The system analyzes source code using Python's native AST module and the Tree-sitter parser for Java and JavaScript, generates UML class diagrams via Mermaid.js integration, and exports documentation in PDF, Markdown, HTML, and XML formats. An intelligent 48-hour caching strategy reduces AI API costs by up to 80%, while a seven-model fallback chain over OpenRouter ensures continuous service availability during rate-limiting periods. Internal testing demonstrates that the multi-agent architecture achieves a 73% reduction in technical error rates compared to single-LLM approaches, with format consistency improved to 95%.

The system integrates Django as the backend framework, React for the frontend, MongoDB for flexible code and documentation storage, and MySQL for structured user and session management.

This project fills a critical market gap: no existing commercial solution provides end-to-end automation combining structural code analysis, test generation and execution, and comprehensive documentation generation within a unified, collaborative web platform.

ملخص

يعد توثيق البرمجيات واختبارها ركيزتين أساسيتين لضمان جودة الكود وقابليته للصيانة على المدى الطويل.

غير أن فرق التطوير الحديثة تعاني من تحديات متزايدة تتمثل في قدم التوثيق وانعدام تغطيته لأجزاء جوهرية من النظام، إضافة إلى الجهد اليدوي الذي يستهلك ما يصل إلى 40% من وقت المطور.

وهو منصة ويب ذكية تعتمد على الذكاء الاصطناعي AutoTest & DocGen Manager تقدم هذه الأطروحة نظام لأتمتة توليد التوثيق و اختبارات البرمجيات بصورة متكاملة من البداية إلى النهاية.

يرتكز النظام على معمارية متعددة العوامل (Multi-Agent Architecture) تضم وكلاء متخصصين: وكيل المستوى الرفيع (HighLevelAgent) الذي يولد توثيقاً تنفيذياً وصفيًا، ووكيل منخفض المستوى (LowLevelAgent) الذي ينتج تحليلًا تقنيًا تفصيليًا للأساليب و العلاقات، ووكيل التحقق (VerifierAgent) الذي يعتمد آلية التصحيح الذاتي للتأكد من دقة المخرجات.

يُحلّل النظام الكود البرمجي باستخدام محلل AST لـ Python ومحلل Tree-sitter للغات Java و JavaScript، وينشئ مخططات UML للصفوف عبر Mermaid.js، ويُصدّر التوثيق بصيغ متعددة تشمل PDF و Markdown و HTML و XML. كما يعتمد نظام تخزين مؤقت ذكي يبلغ عمره 48 ساعة يُخفّض تكاليف استدعاء الذكاء الاصطناعي بنسبة 80%، ويدعم سلسلة احتياطية من سبعة نماذج مجانية عبر OpenRouter لضمان استمرارية الخدمة عند تجاوز حدود الاستخدام.

أثبتت نتائج الاختبار الداخلي أن المعمارية متعددة العوامل تُحقّق انخفاضاً بنسبة 73% في معدل الأخطاء مقارنةً بالنماذج أحادية الوكيل، فضلاً عن رفع اتساق التنسيق إلى 95%. يتكامل النظام مع Django كإطار خلفي، و React كواجهة أمامية، و MongoDB لتخزين الكود والتوثيق، و MySQL لإدارة المستخدمين.

يسد هذا المشروع فجوة حقيقية في السوق؛ إذ لا يوجد نظام تجاري يجمع في آن واحد بين تحليل الكود البنيوي وتوليد الاختبارات وتنفيذها وتوليد التوثيق الشامل ضمن منصة ويب واحدة متكاملة.

Table of Contents

SUPERVISION CERTIFICATION.....	2
ملخص.....	4
List of abbreviations.....	11
Abbreviations.....	11
Chapter 1 Introduction.....	13
Chapter1:Introduction.....	14
1.1 General Overview and Project Background:.....	14
1.2 Problem Statement and Practical/Academic Significance:.....	15
1.2.1 The Core Problem:.....	15
1.2.2 Practical Significance:.....	15
1.2.3 Academic Significance:.....	16
1.3 Project Objectives (Primary and Secondary):.....	16
1.3.1 Primary Objectives:.....	16
1.3.2 Secondary Objectives:.....	16
1.4 Project Scope, Constraints, and Assumptions:.....	17
1.4.1 Project Scope:.....	17
1.4.2 Constraints:.....	17
1.4.3 Assumptions:.....	18
1.5 Project Contribution to Knowledge and the Field:.....	19
1.6 Research Methodology and Scientific Approaches:.....	19
1.6.1 Engineering Research Method (Primary Method):.....	20
1.6.2 Experimental Method (Comparative Evaluation):.....	20
1.6.3 Descriptive-Analytical Method (Literature Review and Market Analysis):.....	20
1.7 Document Organization:.....	20
1.8 Definitions and Key Concepts.....	22
Chapter 2: Development Methodology and Project Management.....	24
Chapter 2: Introduction:.....	25
2.1 Development Methodology: Agile Scrum:.....	25
2.1.1 Why Agile Scrum?.....	25
2.1.2 Scrum Framework Implementation:.....	25
2.1.3 Sprint Structure:.....	26
2.2 Project Timeline Plan (Gantt Chart):.....	28
2.3 Risk Management:.....	28
2.5 Quality Management in the Project:.....	31
2.5.1 Code Quality Assurance:.....	31
2.5.2 Testing Quality Assurance:.....	31
2.5.3 Documentation Quality Assurance:.....	31
2.6 Project Management Tools:.....	31
2.7 Monitoring and Continuous Evaluation Mechanisms:.....	32
2.7.1 Weekly Team Meetings:.....	33
2.7.2 Supervisor Reviews:.....	33
2.8 Software Configuration Management (SCM):.....	33
2.8.1 Git Repository Structure:.....	33
2.8.2 Branching and Merging Strategy:.....	34
2.8.3 Team Members and Responsibilities:.....	36
2.8.4 Collaborative Workflow — Pull Requests:.....	37
2.8.5 Development Workflow:.....	39
2.8.6 Tags and Versioning:.....	39

2.8.7 Summary:	41
Chapter 3: Literature Review	42
3.1 Introduction:	43
3.2 Modern Technologies and Frameworks in the Field:	43
3.3 Similar Systems Analysis:	44
3.3.1 DocuWriter.ai:	45
3.3.2 Swimm.io:	46
3.3.3 Sourcegraph:	47
3.3.4 Qodo.ai (formerly Codium):	48
3.3.5 CodeSee Maps:	49
3.3.6 Mintlify:	50
3.4 Feature Comparison & Gap Analysis:	51
3.4.1 System Limitations Analysis:	52
3.5 Scientific Foundation: Research Papers Review:	52
3.5.1 Key Insights and Contributions of the Reviewed Papers:	54
3.6 Current and Future Trends in the Field:	55
Chapter 4: Requirements Modeling and Analysis	57
4.1 Introduction:	58
4.2 Feasibility Study:	58
4.2.1 Technical Feasibility:	58
4.2.2 Economic Feasibility:	59
4.2.3 Operational Feasibility:	59
4.2.4 Schedule Feasibility:	60
4.3 Requirements Gathering:	60
4.4 Stakeholder Analysis:	61
4.5 Functional Requirements:	61
4.6 Non-Functional Requirements:	64
4.8 UML Diagrams:	66
4.8.1 Main Use Case Diagram:	66
4.8.2 Use Case Descriptions:	67
4.8.3 Sequence Diagrams:	70
4.8.4 Initial Test Case Descriptions:	72
4.9 Modeling According to Methodology (Agile Scrum):	74
4.9.1 User Stories:	74
4.9.2 Story Mapping:	75
4.9.3 Product Backlog:	76
4.10 Requirements Analysis and Documentation:	76
4.10.1 Requirements Traceability Matrix (RTM):	77
4.10.2 Initial Test Cases:	79
Chapter 5 : System Architecture Design	82
5.1 High-Level Design:	83
5.1.1 System Block Diagram:	83
5.1.2 Architectural Design Pattern:	84
5.1.3 System Architecture Design:	85
5.2 Database Design:	86
5.2.1 Entity-Relationship Diagram (ERD Diagram):	86
5.2.2 Relational Schema:	87
5.2.3 Constraints and Keys:	89
5.3 External API Design:	90

5.3.1 API Gateway Overview:	90
5.4.1 Authentication Mechanism:	92
5.4.2 User Roles and Permissions:	92
Additional Security Measures:	93
Chapter 6 : Implementation Environment and Tools	94
6.1 Programming Languages and Frameworks:	95
6.1.1 Frontend: JavaScript and React.js:	95
6.1.2 Backend Services: Python and Django 4.2:	95
6.1.3 API Gateway: Nginx:	95
6.2 Development Environment and Project Structure:	96
6.2.1 Operating System and Hardware:	96
6.2.2 Integrated Development Environments:	97
6.2.3 Containerisation with Docker and Docker Compose:	97
6.2.4 Project Structure (Monorepo):	97
6.2.4.1 UPM_Project (User & Project Management Service):	97
6.2.4.2 Ai_project (Code Intelligence Service):	98
6.2.4.3 Notification_Service (Notification Service):	98
6.3 Database Servers:	99
6.3.1 MySQL 8.0 — User & Project Management Service:	99
6.3.2 MongoDB 7.0 — Code Intelligence Service:	99
6.3.3 Redis 7.x — Celery Message Broker and Cache:	99
6.4 Frontend Dependencies (package.json):	100
6.4.2 Shared Backend Dependencies (requirements.txt):	100
6.4.3 Service-Specific Dependencies:	101
6.5 Continuous Integration and Deployment (CI/CD):	102
6.5.2 Continuous Deployment Pipeline:	102
6.5.3 Docker Containerisation Strategy:	102
6.6 Testing Environment:	103
6.6.2 Staging and Integration Testing:	103
6.6.3 Performance and Load Testing:	103
Chapter 7: Program Implementation	104
7.1 File Structure and Project:	105
7.1.1 UPM_Project Directory Tree:	105
7.1.2 Ai_project Directory Tree:	106
7.1.3 Notification_Service Directory Tree:	106
7.1.4 Common Structural Patterns:	107
7.2 Module Implementation:	107
7.2.1 User and Project Management Service (UPM_Project):	107
7.2.2 Code Intelligence Service (Ai_project):	108
7.2.3 Notification Service (Notification_Service):	109
7.3 API Implementation:	109
7.3.1 Authentication API — JWT with Role-Aware Tokens:	109
7.3.2 Project Management API:	110
7.3.3 Code Intelligence API:	111
7.3.4 Notification API:	111
7.4 Data Access Layer Implementation:	112
7.4.1 Django ORM with MySQL (UPM_Project):	112
7.4.2 Pydantic Models with MongoDB (Ai_project):	112
7.4.3 Django ORM with SQLite (Notification_Service):	113
7.4.4 Cross-Service Data Access:	113

7.5 Business Logic Layer Implementation:.....	114
7.5.1 User Management Logic:.....	115
7.5.2 Project and Artifact Management:.....	115
7.5.3 Code Analysis Pipeline:.....	116
7.5.4 AI Documentation Generation:.....	116
7.5.5 Notification Processing:.....	117
7.6 User Interface Implementation:.....	117
7.7 Intelligent Algorithms Implementation:.....	120
7.7.1 Python AST Analysis Engine:.....	120
7.7.2 Multi-Language Parsing via Tree-sitter:.....	121
7.7.3 UML Diagram Generation Algorithm:.....	122
7.7.4 Topological Code Processing:.....	123
7.7.5 LLM-Powered Documentation Generation:.....	125
7.7.6 Intelligent Caching Strategy:.....	126
7.7.7 LangChain and LangGraph Integration:.....	128
7.7.8 Multi-Agent Documentation System:.....	129
7.7.9 Conflict Detection Algorithm:.....	130
Chapter 7.5:AI Methodology: Prompt Engineering & Evaluation.....	133
7.5.1. Introduction:.....	134
7.5.2. Phase 1: Code Analysis and Feature Extraction:.....	134
The first phase transforms source code into structured representations through Abstract Syntax Tree (AST) analysis. The system implements a factory pattern (processor_factory.py) that routes code files to appropriate processors: PythonProcessor (Python's ast module), TreeSitterProcessor (Java, C#, TypeScript via tree-sitter), and GenericProcessor (fallback). Each processor extracts class names, method signatures, attributes, inheritance relationships, and route definitions.....	134
7.5.2.1 AST Unification and Class Diagram Generation:.....	135
7.5.3 Phase 2: Documentation Generation Pipeline:.....	135
7.5.3.1 Agent Design and Specialization:.....	136
7.5.3.2 LLM Client and Model Management:.....	136
7.5.4 Phase 3: Four-Layer Evaluation Framework:.....	136
7.5.4.1 Layer 1: Factual Accuracy (35%);.....	137
7.5.4.2 Layer 2: Completeness (25%):.....	138
7.5.4.3 Layer 3: LLM Judge (40%):.....	138
7.5.4.4 Layer 4: Human Review (20%, Optional):.....	138
7.5.4.5 Final Score and Verdict:.....	138
7.5.5 Phase 4: Conflict Detection System:.....	139
7.5.5.1 FingerprintBuilder and Rename Detection:.....	140
7.5.5.2 Fusion Reasoning Semantic Analysis:.....	140
7.5.5.3 Report Generation and Grade System:.....	140
7.5.6 Conflict Detection Evaluation (LLM-as-a-Judge):.....	140
7.5.6.1 JUDGE_PROMPT Design:.....	140
7.5.6.2 Conflict Evaluation Results:.....	140
7.5.6.3 Strengths and Weaknesses:.....	141
7.5.7 System Architecture Overview:.....	141
7.5.8 Summary:.....	142
Chapter 8:Testing and Evaluation.....	143
8.1 Test Plan:.....	144
8.1.1 Test Scope:.....	144
8.1.2 What Will Be Tested:.....	144
8.1.3 Testing Tools:.....	144

8.1.4 Test Schedule:.....	144
8.2 Types of Testing:.....	144
8.2.1 Unit Testing:.....	144
8.2.2 Integration Testing:.....	145
8.2.3 System Testing:.....	145
8.2.4 Acceptance Testing:.....	145
8.2.5 Security Testing:.....	145
8.2.6 Performance, Load & Stress Testing:.....	146
8.3 Test Diagrams:.....	146
8.3.1 Test Case Diagrams:.....	146
8.3.2 Test Result Diagrams:.....	146
8.4 Test Results and Analysis:.....	152
8.4.1 UPM Service — Test Summaries:.....	152
8.4.2 UPM Performance Test — Endpoint Detail:.....	152
8.4.3 UPM Load Test — Endpoint Detail:.....	153
8.4.4 UPM Stress Test — Endpoint Detail:.....	153
8.4.5 AI Service — Test Summaries:.....	153
8.4.6 Conflict Detection Endpoints — Cross-Test Results ★ :.....	154
8.4.7 Result Analysis:.....	154
8.5 Quality Metrics:.....	154
8.5.1 Test Coverage:.....	154
8.5.2 Defect Density:.....	155
8.5.3 Response Time Metrics:.....	155
8.5.4 Requests Per Second (RPS):.....	155
8.6 Usability Testing:.....	155
8.6.1 Manual Functional Test Scenarios:.....	155
8.6.2 Strengths and Improvement Suggestions:.....	156
Chapter 9: Results and Analysis.....	157
9.1 System Implementation Results:.....	158
9.1.1 Functional Results — Core Features:.....	158
9.1.2 Interface and Usage Results:.....	158
9.3 Objectives Achievement Verification:.....	159
9.4 Strengths and Weaknesses Analysis:.....	160
9.4.1 Key Strengths:.....	160
9.4.2 Limitations and Areas for Improvement:.....	161
Chapter 10: Conclusion and Recommendations.....	162
10.1 Project Summary and Key Achievements:.....	163
10.2 Challenges and Difficulties:.....	163
10.4 Recommendations:.....	165
10.4.1 Recommendations for Future Student Projects:.....	165
10.4.2 Recommendations for Organizations Adopting the Platform:.....	165
References.....	167

Table of Figures

Figure 1: Gantt chart.....	28
Figure 2: Use Case Diagram.....	67
Figure 3: Code Comparison Sequence Diagram.....	70
Figure 4: Import Project Sequence Diagram.....	71
Figure 5: System Block Diagram.....	83
Figure 6: Strategy Pattern — Multi-Language Processing Architecture.....	84
Figure 7: Template Method Pattern — Document Export Architecture.....	85
Figure 8: System Architecture.....	86
Figure 9: Entity-Relationship Diagram (ERD Diagram).....	87
Figure 10: Class Diagram Interface.....	118
Figure 11: Project File Tree Interface.....	118
Figure 12: Conflict Detection Interface.....	119
Figure 13: CodVSDoc Detection.....	119
Figure 14: CodVSCod Detection.....	119
Figure 15: Full Analysis Detection.....	119
Figure 16: Research Methodology-4Phase Pipeline.....	134
Figure 17: AI Documentation System Architecture.....	135
Figure 18: Layer Evaluation Framework.....	137
Figure 19: Conflict Detection Pipeline.....	139
Figure 20: UPM Service: Performance Test (10 concurrent users) Response Times per Endpoint	147
Figure 21: UPM Service: Performance Test (10 Users) — Response Time & Active Clients.....	147
Figure 22: UPM Service: Load Test (50 concurrent users) Response Times per Endpoint.....	148
Figure 23: UPM Service: Load Test (50 Users) — Response Time & Active Clients.....	148
Figure 24: UPM Service: Stress Test (100 concurrent users) Response Times per Endpoint.....	149
Figure 25: UPM Service: Stress Test (100 Users) — Response Time & Active Clients.....	149
Figure 26: AI Service: Performance Test (10 concurrent users) Response Times per Endpoint (yellow = Conflict Detection).....	150
Figure 27: AI Service: Performance Test (10 Users) — Response Time & Active Clients.....	150
Figure 28: AI Service: Load Test (50 concurrent users) Response Times per Endpoint (yellow = Conflict Detection).....	151
Figure 29: AI Service: Load Test (50 Users) — Spike indicates Docker memory constraint at ~0:30	151
Figure 30: AI Service: Stress Test (100 concurrent users) Response Times per Endpoint (yellow = Conflict Detection).....	152
Figure 31: AI Service: Stress Test (100 Users) — Response time reaches Locust 90s timeout ceiling	152

Index of Tables

Table 1: Risk Management.....	28
Table 2: Project Management Tools.....	32
Table 3: Branching and Merging Strategy.....	34
Table 4: Feature Comparison.....	51
Table 5: Scientific Foundation: Research Papers Review.....	53
Table 6: Schedule Feasibility.....	60

Table 7: Functional Requirements.....	62
Table 8: Non-Functional Requirements.....	64
Table 9: Import Project DESCRIPTION.....	67
Table 10: Compare Code Version Description.....	69
Table 11: Initial Test Case Descriptions.....	72
Table 12: User Stories.....	74
Table 13: Product Backlog.....	76
Table 14: Requirements Traceability Matrix (RTM).....	77
Table 15: Initial Test Cases.....	80
Table 16: Relational Schema.....	87
Table 17: API Gateway.....	90
Table 18: Programming Languages and Frameworks.....	96
Table 19: Libraries and Dependencies.....	101
Table 20: File Structure and Project.....	105
Table 21: API Implementation.....	111
Table 22: Available Free Models on OpenRouter.....	136
Table 23: LLM Judge Evaluation Results.....	138
Table 24: Conflict Evaluation Results (LLM-as-a-Judge).....	140
Table 25: Missed Conflicts Identified by the Judge.....	141
Table 26: System Technology Stack.....	141
Table 27: Test Plan.....	144
Table 28: Types of Testing.....	145
Table 29: Test Diagrams.....	146
Table 30: Test Results and Analysis.....	152
Table 31: Quality Metrics.....	154
Table 32: Usability Testing.....	155
Table 33: System Implementation Results.....	158
Table 34: Objectives Achievement Verification.....	159
Table 35: Performance KPIs.....	161
Table 36: Key Achievements.....	163

List of abbreviations

Abbreviations	Definition
API	Application Programming Interface
AST	Abstract Syntax Tree
AI	Artificial Intelligence
CI/CD	Continuous Integration / Continuous Deployment
ERD	Entity-Relationship Diagram
LLM	Large Language Model
MVC	Model-View-Controller
MVP	Minimum Viable Product
NLP	Natural Language Processing
PO	Primary Objective
RAG	Retrieval-Augmented Generation
RBAC	Role-Based Access Control

RTM	Requirements Traceability Matrix
SCM	Software Configuration Management
SMART	Specific, Measurable, Achievable, Realistic, Time-bound
SO	Secondary Objective
SSO	Single Sign-On
UML	Unified Modeling Language
UPM	User & Project Management
URL	Uniform Resource Locator
ZIP	Compressed File Archive Format

Chapter 1 Introduction

Chapter 1: Introduction

This chapter introduces the **AutoTest & DocGen Manager** project — an intelligent, AI-powered platform designed to revolutionize software development workflows through automated testing and comprehensive documentation generation. The chapter begins by establishing the project context and background, then identifies the core documentation and testing challenges that motivated this work and explains their practical and academic significance. Following this, the chapter defines the project objectives using SMART criteria, delineates the system scope alongside its constraints and assumptions, and articulates the novel contributions this project makes to the field of software engineering. The chapter concludes by describing the scientific research methodology applied throughout the thesis and providing a structured overview of the remaining chapters.

1.1 General Overview and Project Background:

The AutoTest & DocGen Manager is an intelligent, AI-powered platform designed to revolutionize software development workflows through automated testing and comprehensive documentation generation. This project emerges from the growing need to address the persistent challenges that development teams face in maintaining high-quality code documentation and testing coverage while meeting aggressive delivery timelines. The system leverages advanced artificial intelligence techniques, including multi-agent architectures and topological code processing, to automate traditionally manual, time-intensive processes that often become bottlenecks in the software development lifecycle.

In the modern software development landscape, the complexity of software systems has increased exponentially, making manual documentation and testing increasingly impractical. Development teams face a critical trade-off: investing time in comprehensive documentation and testing versus accelerating feature delivery. According to industry research, developers spend up to 40% of their time on documentation-related tasks, and over 60% of code documentation becomes outdated within six months of creation. The AutoTest & DocGen Manager addresses this challenge by providing a unified platform that integrates code analysis, intelligent test generation, automated test execution, and comprehensive documentation generation into a seamless, production-ready workflow.

The system is designed as a web-based application that allows development teams to upload their codebases, automatically analyze project structure and dependencies, generate comprehensive unit and integration tests, execute those tests with immediate feedback, and produce accurate, up-to-date documentation including API references, UML diagrams, and natural language code explanations. The platform supports multiple programming languages including Python, Java, and JavaScript, and employs a novel multi-agent AI architecture with specialized agents — HighLevelAgent, LowLevelAgent, and VerifierAgent — that collaborate to produce accurate and comprehensive outputs, achieving a 73% error reduction compared to single-LLM approaches.

1.2 Problem Statement and Practical/Academic Significance:

The core problem addressed by this project is the absence of an end-to-end automated system that simultaneously handles software testing and documentation generation. Development teams today are forced to choose between velocity and quality: comprehensive documentation and thorough testing both consume enormous amounts of developer time, yet no existing commercial or academic solution unifies these two critical tasks under an intelligent, automated framework. This gap leads directly to technical debt, reduced maintainability, and elevated risk of defects reaching production.

1.2.1 The Core Problem:

Software documentation and testing are fundamental pillars of maintainable, high-quality code. Comprehensive documentation enables knowledge transfer, facilitates onboarding, and ensures long-term system maintainability. Similarly, thorough testing validates functionality, prevents regressions, and builds confidence in code changes. However, modern development teams face persistent challenges in both areas that create a compounding negative effect on software quality, team productivity, and project sustainability.

The documentation challenges are multifaceted and deeply interconnected:

- **Outdated Documentation:** Research indicates that over 60% of code documentation becomes outdated within six months of creation. When code evolves, developers frequently neglect to update corresponding documentation, creating a dangerous disconnect between implementation and documentation.
- **Incomplete Coverage:** Many projects lack documentation for critical APIs, complex algorithms, or architectural decisions. This knowledge gap forces developers to reverse-engineer code logic, significantly increasing development time and error rates.
- **Missing Visual Representations:** The absence of UML class diagrams, and dependency visualizations makes understanding system architecture unnecessarily difficult, particularly for large-scale projects or new team members.
- **Manual Effort and Time Waste:** Developers spend up to 40% of their time on documentation-related tasks, diverting focus from core development activities and delaying feature delivery.

These challenges compound over time, leading to:

- Reduced software quality and increased technical debt
- Elevated maintenance costs and slower development cycles
- Difficult onboarding for new team members
- Higher risk of critical bugs reaching production
- Decreased developer productivity and team morale

1.2.2 Practical Significance:

From a practical perspective, the AutoTest & DocGen Manager addresses a genuine and urgent need in the software development industry. By automating documentation and testing end-to-end, the system enables teams to maintain high-quality documentation and thorough testing coverage

without sacrificing development velocity. The practical impact includes: reducing the time developers spend on documentation from 40% to near zero for documentation generation tasks; ensuring documentation stays synchronized with code through automatic outdated detection and regeneration; providing immediate test feedback with coverage analysis to identify untested code paths; and enabling team collaboration through a web-based interface accessible to all stakeholders.

1.2.3 Academic Significance:

From an academic perspective, this project contributes to the growing body of research on applying multi-agent AI architectures to software engineering tasks. While existing research has explored individual components such as automated test generation (Qodo.ai, DocAgent) and documentation generation (DocuWriter.ai, Mintlify), no academic or commercial system has successfully unified these capabilities under a single intelligent framework with a multi-agent architecture. The project validates the hypothesis that specialized, collaborating AI agents can achieve significantly higher accuracy than single-model approaches in code understanding tasks. The 73% error reduction achieved by the multi-agent architecture compared to single-LLM approaches provides empirical evidence supporting the effectiveness of this paradigm. Additionally, the topological code processing approach contributes a novel methodology for overcoming LLM token limitations in large codebase analysis.

1.3 Project Objectives (Primary and Secondary):

All objectives are defined using the SMART methodology — Specific, Measurable, Achievable, Realistic, and Time-bound — to ensure clarity and verifiability throughout the project lifecycle (February–May 2026).

1.3.1 Primary Objectives:

- PO1: Build an intelligent system for automated code analysis that extracts complete project architecture — including classes, functions, methods, parameters, return types, and hierarchical relationships — with at least 95% structural accuracy, achievable within the project timeline.
- PO2: Implement a comprehensive documentation generation module that produces API documentation in Swagger/OpenAPI format, UML class diagrams via Mermaid.js, and natural language code explanations — all generated automatically from analyzed code with measurable reduction in manual documentation effort.

1.3.2 Secondary Objectives:

- SO1: Implement a multi-agent AI architecture with specialized agents (HighLevelAgent, LowLevelAgent, VerifierAgent) that collaborate to produce accurate outputs, targeting a 70% or greater error reduction compared to single-LLM approaches — measured through internal comparative testing.

- SO2: Develop a topological code processing system that analyzes code based on dependency relationships rather than linear order, overcoming LLM token limitations and improving understanding of large codebases — verified by successful analysis of projects exceeding standard context window limits.
- SO3: Create an interactive web-based dashboard for project uploads, test review, analysis monitoring, and report downloads with real-time progress tracking, response time under 3 seconds, and at least 90% user satisfaction rate — measured through usability testing surveys.
- SO4: Integrate with CI/CD pipelines (GitHub, GitLab, Jenkins) to automatically trigger documentation generation on each commit — measured by the number of commits that successfully trigger the automated pipeline during the evaluation period, targeting 100% trigger success rate.
- SO5: Implement intelligent caching to store generated explanations, reducing AI API costs by up to 80% while maintaining output quality — measured by comparing API call counts with and without caching across identical analysis tasks.
- SO6: Ensure robust security with BCrypt password hashing, optional SSO support, code encryption, and multi-level access control — verified through security testing including penetration testing of authentication endpoints and access control validation, targeting zero critical vulnerabilities in the final release.

1.4 Project Scope, Constraints, and Assumptions:

1.4.1 Project Scope:

The scope of the AutoTest & DocGen Manager encompasses the following key areas and deliverables. The system will be a web-based platform that automates the complete pipeline from code analysis through execution to comprehensive documentation. Specifically, the project scope includes:

- Comprehensive codebase analysis with project tree extraction and dependency mapping
- Intelligent documentation generation including API documentation (Swagger/OpenAPI), UML diagrams (via Mermaid.js), and natural language explanations
- Outdated documentation detection and automatic regeneration
- An interactive web dashboard for project management and report downloads in multiple formats (PDF, Markdown, HTML, XML)
- CI/CD integration for continuous testing and documentation updates (GitHub, GitLab, Jenkins)
- Multi-language support for Python, Java, and JavaScript with extensible architecture for future languages

1.4.2 Constraints:

Documenting project constraints is essential as it protects the project from unfair criticism and clarifies the boundaries within which the system was designed and evaluated. The project operates under the following constraints:

- **Technical Constraints:** The system relies on LLM APIs (GPT-3.5/4 via OpenRouter) for code understanding and generation, which introduces latency and cost considerations. API rate limits may affect processing speed for large codebases. Multi-language support is initially limited to Python, Java, and JavaScript, with other languages requiring additional parser development.
- **Time Constraints:** The project timeline spans from February 21, 2026 to May 30, 2026, providing approximately 14 weeks for complete development, testing, and deployment. This timeline requires careful prioritization of features and iterative delivery through three development sprints.
- **Resource Constraints:** The development team consists of four undergraduate students working part-time, which limits the scope of features that can be fully implemented within the timeline. Cloud hosting and LLM API costs must be managed within a student project budget.
- **Platform Constraints:** The system is designed as a web-based application requiring server infrastructure for deployment. Desktop and mobile applications are outside the current scope. The platform targets modern web browsers and does not support legacy browser versions.

1.4.3 Assumptions:

The following assumptions underpin the project design and development decisions. If any assumption proves invalid, the project plan may require adjustment:

- Users have basic programming knowledge and can navigate a web-based interface without extensive training.
- Uploaded codebases are syntactically valid and can be parsed by the supported language parsers (Tree-sitter for Java, AST for Python). Code with syntax errors may not be fully analyzable.
- Users have legitimate access to the code they upload for analysis. The platform is not responsible for any intellectual property violations committed by users, as it assumes submitted code is owned or licensed by the submitting party.
- Internet connectivity is available for accessing LLM APIs, and users accept the latency implications of cloud-based AI processing.
- The project will be deployed on a server with sufficient resources (CPU, RAM, storage) to handle concurrent analysis requests within acceptable response times.
- Users possess a basic working knowledge of English, as the platform interface and generated documentation are primarily in English. Bilingual (Arabic/English) support is outside the current project scope.
- The platform will not be used to analyze malicious code or code containing intentional security vulnerabilities designed to compromise the server environment. The system assumes all uploaded content is intended for legitimate analysis purposes.

1.5 Project Contribution to Knowledge and the Field:

The AutoTest & DocGen Manager makes several significant contributions to both software engineering practice and the academic research community. These contributions represent novel advances that go beyond existing solutions and address identified gaps in the current state of the art.

- **Novel Multi-Agent Architecture for Software Engineering:** The project presents a specialized multi-agent AI architecture — comprising HighLevelAgent, LowLevelAgent, and VerifierAgent — working collaboratively on code understanding tasks. This specific configuration has not been previously demonstrated in academic literature or commercial products. The empirical 73% error reduction compared to single-LLM approaches provides quantifiable evidence for the effectiveness of this paradigm in software engineering contexts.
- **Topological Code Processing Methodology:** The project introduces topological code processing as a practical method for overcoming LLM token limitations when analyzing large codebases. Rather than processing code linearly, the system analyzes code based on dependency relationships, building context incrementally from foundational modules to dependent ones. This methodology contributes a novel solution to the well-documented challenge of applying LLMs to large-scale code analysis tasks.
- **End-to-End Automation Pipeline:** The project demonstrates the feasibility of integrating automated testing and documentation generation into a unified end-to-end pipeline. While individual tools exist for test generation (Qodo.ai) and documentation generation (DocuWriter.ai, Mintlify), no system has successfully combined these capabilities with code analysis, outdated detection, and CI/CD integration under a single intelligent framework.
- **Self-Correction Mechanisms for AI-Generated Artifacts:** The system implements verification-based quality assurance inspired by academic research (CCISOLVER), where verification agents validate output accuracy against actual code structure and automatically detect and correct inconsistencies. This self-correcting capability significantly improves output reliability compared to single-pass generation approaches, contributing a practical implementation pattern applicable to other AI-generated software engineering artifacts.
- **Cost-Efficient AI Integration:** The 48-hour intelligent caching strategy combined with a seven-model fallback chain over OpenRouter contributes a practical pattern for reducing AI API costs by up to 80% in production systems, offering a replicable cost-optimization approach for AI-powered developer tools.

1.6 Research Methodology and Scientific Approaches:

This section describes the scientific research methodology applied in conducting and writing this thesis. The software development methodology (Agile with iterative sprints) is discussed in detail in Chapter 2. The thesis employs a combination of three complementary research approaches, each selected for its suitability to a specific aspect of the investigation:

1.6.1 Engineering Research Method (Primary Method):

The primary research methodology applied is the Engineering Research Method, in which the investigation is carried out through the systematic design, construction, and evaluation of a working artifact — in this case, the AutoTest & DocGen Manager platform. This method is appropriate because the research questions are answered by demonstrating that a novel system can be built and that it performs at a measurable level of quality. The research process follows the standard engineering cycle: problem identification, requirements analysis, design, implementation, testing, and evaluation against defined success criteria. Conclusions are drawn from the measured behavior of the implemented system rather than from purely theoretical reasoning.

1.6.2 Experimental Method (Comparative Evaluation):

The Experimental Method is applied in Chapter 9 (Results and Analysis) to validate the key architectural hypothesis of this thesis: that a multi-agent AI architecture achieves higher accuracy and lower error rates than single-LLM approaches for code documentation and testing tasks. Controlled experiments are designed with clearly defined independent variables (architecture type: multi-agent vs. single-LLM), dependent variables (error rate, format consistency, documentation completeness), and controlled variables (input codebase, output format requirements, model version). Results are reported quantitatively, enabling reproducibility and peer validation. This method provides the empirical foundation for the 73% error reduction claim.

1.6.3 Descriptive-Analytical Method (Literature Review and Market Analysis):

The Descriptive-Analytical Method is applied in Chapter 3 (Literature Review) to systematically survey, describe, and critically analyze existing systems and academic research in the domain of automated code documentation and testing. Six production-ready commercial systems (DocuWriter.ai, Swimm.io, Sourcegraph, Qodo.ai, CodeSee Maps, Mintlify) and eight academic research papers are described in terms of their features, methodologies, strengths, and limitations. The analysis is then synthesized into a feature comparison matrix and gap analysis that directly informs and justifies the design decisions made in this project. This method ensures that the proposed system is grounded in a thorough understanding of the current state of the art rather than developed in isolation.

1.7 Document Organization:

This report is structured to provide a comprehensive view of the AutoTest & DocGen Manager project, from initial concept through implementation and evaluation. Each chapter builds upon the previous sections to deliver a complete understanding of the system's design, development, and capabilities:

- **Chapter 1:** Project Introduction (Current) — Establishes the project context, defines the problems being addressed, outlines core objectives using SMART criteria, specifies the project scope with constraints and assumptions, articulates the project's contributions to the field, describes the scientific research methodology applied in the thesis, and previews the report structure.

- **Chapter 2:** Development Methodology and Project Management — Details the Agile development methodology adopted with iterative sprint-based delivery, sprint planning and Gantt chart scheduling, risk management strategies with a risk register, team roles and task distribution, quality management processes, project management tools, monitoring and evaluation mechanisms, and Git-based software configuration management.
- **Chapter 3:** Literature Review and Theoretical Framework — Provides comprehensive analysis of existing systems through feature comparison tables, reviews modern techniques and frameworks in the domain (TensorFlow, PyTorch, Django, React), presents comparative analysis identifying strengths and weaknesses, identifies the research gap and proposed innovation, establishes the theoretical foundation through key research papers, and surveys current trends including LLMs, MLOps, and Generative AI.
- **Chapter 4:** System Analysis and Modeling — Presents comprehensive requirements engineering, feasibility study (technical and financial), requirements elicitation through interviews and surveys, stakeholder analysis, functional and non-functional requirements, UML modeling (use case, sequence, activity, class diagrams), initial test case descriptions, and Requirements Traceability Matrix (RTM).
- **Chapter 5:** System Design — Describes the high-level system architecture using microservices with Django backend, architectural patterns employed (MVC, State Pattern, Template Method), database design with ER diagrams and relational schemas for both MySQL and MongoDB, API design with endpoint specifications, and security design with JWT authentication and multi-level access control.
- **Chapter 6:** Implementation Environment and Tools — Documents the technical implementation environment including programming languages and frameworks, development tools and libraries, database management systems, version control and CI/CD tools, and testing environment setup.
- **Chapter 7:** Practical Implementation — Presents the AI architecture in depth, covers code structure and project organization, module implementation with code samples, API implementation, data access layer, business logic layer, user interface implementation with screenshots, and AI algorithm implementation including the multi-agent architecture and topological processing.
- **Chapter 7.5 : AI Methodology: Prompt Engineering & Evaluation:** Provides an in-depth exposition of the AI methodology underpinning the system's intelligence. It details the AI technology stack, presents the prompt engineering strategies for the three specialized agents (HighLevelAgent, LowLevelAgent, and VerifierAgent), explains the agent routing logic and the complete prompt engineering pipeline, describes the semantic conflict detection methodology including pipeline architecture and severity assessment, introduces the LLM-as-a-Judge evaluation framework with its 4-layer scoring architecture for explanation evaluation and conflict evaluation, and presents the quantitative evaluation results demonstrating the effectiveness of the multi-agent approach.
- **Chapter 8:** Testing and Evaluation — Covers the test plan, testing types (unit, integration, system, performance, security, usability, acceptance), test case diagrams and results, quality metrics including coverage and defect density, and usability testing results from real user sessions.

- **Chapter 9: Results and Analysis** — Presents system implementation results with screenshots, comparative analysis with existing systems, verification of objectives achievement against Chapter 1 goals, strengths and weaknesses analysis, and performance KPI measurements.
- **Chapter 10: Conclusions and Recommendations** — Summarizes the project and key achievements, discusses challenges and difficulties encountered, proposes future development directions (e.g., additional language support, Edge AI integration), and provides recommendations for practitioners and future student projects.

1.8 Definitions and Key Concepts

The following definitions clarify key terms and concepts used throughout this report. Each term is defined concisely to ensure consistent understanding across all readers:

- **Multi-Agent Architecture:** A system design pattern where multiple specialized AI agents collaborate to accomplish complex tasks, with each agent responsible for a specific aspect of the overall workflow.
- **Topological Code Processing:** A method of analyzing source code based on dependency relationships rather than linear file order, building context incrementally from foundational modules to dependent ones.
- **LLM (Large Language Model):** An AI model trained on vast text corpora capable of understanding and generating human-like text, such as GPT-3.5 and GPT-4 used in this project.
- **RAG (Retrieval-Augmented Generation):** A technique that enhances LLM outputs by retrieving relevant context from a knowledge base before generating responses, improving accuracy and relevance.
- **NLP (Natural Language Processing):** A branch of AI focused on enabling computers to understand, interpret, and generate human language, used here for code summarization and explanation generation.
- **Microservices:** An architectural pattern that structures an application as a collection of loosely coupled, independently deployable services, each responsible for a specific business capability.
- **CI/CD (Continuous Integration/Continuous Deployment):** A set of practices that automate the integration, testing, and deployment of code changes, enabling rapid and reliable software delivery.
- **AST (Abstract Syntax Tree):** A tree representation of the syntactic structure of source code, used by the system to parse and analyze Python codebases for structural understanding.
- **Tree-sitter:** An incremental parsing library used by the system to parse Java and JavaScript codebases, providing fast and accurate syntax tree generation for code analysis.

- **Mermaid.js:** A JavaScript-based diagramming tool that renders UML diagrams from text descriptions, used by the system to generate class and sequence diagrams from analyzed code.
- **Swagger/OpenAPI:** A specification for machine-readable interface files for describing, producing, consuming, and visualizing RESTful web services, used as the standard format for API documentation output.
- **Self-Correction Mechanism:** A quality assurance approach where a verification agent validates AI-generated outputs against actual code structure, automatically detecting and correcting inconsistencies.
- **SMART Objectives:** A goal-setting framework ensuring objectives are Specific, Measurable, Achievable, Realistic, and Time-bound, applied throughout this project's objective definitions.
- **Requirements Traceability Matrix (RTM):** A document that maps requirements to their corresponding design elements, implementation components, and test cases, ensuring complete traceability throughout the development lifecycle.
- **BCrypt:** A password hashing function designed for secure storage of user passwords, incorporating a salt to protect against rainbow table attacks and being computationally adaptive to resist brute-force attacks.
- **Quality Gates:** Automated pass/fail criteria based on test coverage and documentation completeness metrics, enforced in the CI/CD pipeline to maintain quality standards.
- **HighLevelAgent:** A specialized AI agent in the system's multi-agent architecture responsible for generating executive-level documentation describing the overall system purpose and key capabilities.
- **LowLevelAgent:** A specialized AI agent responsible for deep technical analysis of methods, class relationships, design patterns, and generating detailed low-level documentation.
- **VerifierAgent:** A specialized AI agent that implements a self-correction mechanism to validate output accuracy against the actual code, detecting and correcting errors produced by the other agents.

Chapter 2: Development Methodology and Project Management

Chapter 2: Introduction:

This chapter details the development methodology adopted for the AutoTest & DocGen Manager project, demonstrating that the project was managed through a structured, methodological approach rather than ad hoc practices. The chapter covers the selected development methodology and its justification, the project timeline with Gantt chart scheduling, risk management strategies, team structure and task distribution, quality management processes, project management tools utilized, monitoring and continuous evaluation mechanisms, and software configuration management practices.

2.1 Development Methodology: Agile Scrum:

The AutoTest & DocGen Manager project was developed using the Agile Scrum methodology throughout its entire lifecycle. Scrum is an iterative, incremental framework for managing complex projects that emphasizes flexibility, continuous feedback, and adaptive planning. This methodology was applied from the project's inception in February 2026 through its completion in May 2026, governing all aspects of planning, development, testing, and delivery.

2.1.1 Why Agile Scrum?

The nature of the AutoTest & DocGen Manager project inherently demands an Agile approach. The system is an AI-powered platform whose behavior and accuracy depend on iterative refinement of machine learning models and multi-agent architectures. Unlike traditional software projects where requirements can be fully specified upfront, the development of AI-driven features requires continuous experimentation, evaluation, and adjustment based on real-world results. The accuracy of the multi-agent architecture, and the relevance of documentation outputs cannot be predicted with certainty at the project's outset; they must be validated empirically through iterative development cycles.

Furthermore, the integration of multiple complex components, including static code analysis, LLM-based generation, and documentation rendering, demands continuous integration and frequent validation to ensure compatibility and correctness. A sequential Waterfall approach would defer integration testing to late in the project lifecycle, creating significant risk of discovering fundamental incompatibilities when there is insufficient time to address them. Scrum's iterative sprint structure naturally enforces regular integration and testing, catching issues early when they are least expensive to fix.

The four-member team structure also aligns perfectly with Scrum's small, cross-functional team model. With clearly defined roles, the team can operate with the autonomy and self-organization that Scrum promotes, enabling rapid decision-making and efficient collaboration without the overhead of heavy management structures.

2.1.2 Scrum Framework Implementation:

The project was executed through four development sprints (Sprint 0 through Sprint 3), preceded by an initial scope determination phase, following the Scrum framework's core ceremonies and artifacts. The following table summarizes the Scrum elements applied in the project:

Scrum Element	Implementation in Project	Frequency/Timing
Product Backlog	Prioritized list of all features, user stories, and tasks managed in GitHub Projects	Continuously refined
Sprint Planning	Team meeting to define sprint goals, select backlog items, and estimate effort	Start of each sprint
Sprint Backlog	Selected items for the current sprint with task breakdown and assignments	Updated daily
Daily Standup	Quick sync: what was done, what will be done, any blockers	Weekly (Sunday meetings)
Sprint Review	Demo of completed features to supervisor; feedback collection	End of each sprint
Sprint Retrospective	Reflection on process: what worked, what to improve, action items	End of each sprint
Increment	Functional, tested increment of the system delivered each sprint	Each sprint
Burndown Chart	Visual tracking of remaining work vs. time in each sprint	Updated weekly

2.1.3 Sprint Structure:

The project was organized into four sprints, preceded by an initial scope determination phase, following an incremental delivery approach where each sprint builds upon the deliverables of the previous one while introducing new capabilities. This incremental delivery approach ensures that a functional portion of the system is always available for demonstration, testing, and feedback, reducing the risk of discovering major issues late in the project lifecycle.

Initial Scope Determination (February 21-23, 2026) was a brief 3-day phase at the project's inception where the team defined the overall project scope, identified key stakeholders, and established the high-level objectives and boundaries of the AutoTest & DocGen Manager. This phase produced the initial project charter and scope statement that guided all subsequent sprint planning activities.

Sprint 0: Literature Review (February 24 - March 7, 2026) was a 10-day sprint dedicated to conducting a comprehensive literature review of existing systems and research papers. This sprint produced the foundational knowledge base that informed all subsequent design and development decisions. The primary deliverables included the review of seven similar production systems (DocuWriter.ai, Swimm.io, Sourcegraph, Qodo.ai, Mintlify), the analysis of five academic research papers on code analysis and AI-driven documentation, and the identification of market gaps that the AutoTest & DocGen Manager would address. Concurrently, the team performed system analysis at a high level (February 27 - March 5) and defined the system architecture and high-level design

(March 4-7), ensuring that the literature findings were immediately translated into architectural decisions.

Sprint 1: Inconsistency Detection & Integration CI/CD (March 8-24, 2026) was a 15-day sprint focused on building the core capabilities for detecting inconsistencies in code documentation and integrating continuous integration and deployment pipelines. The primary deliverables included the development of the core inconsistency detection and CI/CD integration backend module (March 8-24), which implemented the fundamental code analysis engine and the CI/CD webhook integration framework; and the building and assembly of the frontend interface (March 12-24), which created the initial user-facing dashboard for project management and analysis monitoring. This sprint established the foundation upon which all subsequent feature development would build, ensuring that the basic analysis pipeline and user interface were functional and validated before adding more complex AI-driven capabilities.

Sprint 2: Complex Project Handling & Integration CI (March 25 - April 11, 2026) was a 15-day sprint focused on extending the system's capabilities to handle complex, multi-file projects and deepening the CI/CD integration. The primary deliverables included the development of the complex project handling and CI/CD integration backend module (March 28 - April 11), which implemented advanced code parsing for multi-file project structures, cross-file dependency analysis, and enhanced CI/CD integration with GitHub and GitLab; and the building of the complex project handling and CI/CD frontend interface (March 25 - April 11), which extended the dashboard to support multi-project management, advanced analysis configuration, and CI/CD pipeline visualization. This sprint significantly expanded the system's capability from basic single-file analysis to comprehensive multi-project support.

Sprint 3: Code Coverage & Analysis & Notification (April 12-27, 2026) was a 14-day sprint focused on implementing code coverage analysis, advanced code quality metrics, and a notification system. The primary deliverables included the development of the code coverage and notification backend module (April 12-27), which implemented code coverage calculation, quality metrics reporting, and real-time notification delivery; and the building of the code coverage and notification frontend interface (April 15-27), which created the coverage visualization dashboard, quality metrics display, and notification center. This sprint completed the core feature set of the AutoTest & DocGen Manager, providing users with comprehensive visibility into code quality and testing completeness.

Following Sprint 3, the project entered a finalization phase consisting of Documentation (April 28-30), Full Website Integration testing (April 28-30), and Report Finishing (May 1). These final tasks ensured that all project artifacts were complete, the system was fully integrated and tested, and the final project report was prepared for submission.

2.2 Project Timeline Plan (Gantt Chart):

The project timeline spans from February 21, 2026 to May 1, 2026, providing approximately 10 weeks for complete development, testing, and delivery. The Gantt chart below outlines all tasks, their durations, dependencies, and sprint assignments as managed in MS Project. The timeline is structured to ensure that each sprint delivers a functional increment of the system, with overlapping tasks enabling parallel frontend and backend development within each sprint.




Figure 1: Gantt chart

The Gantt chart reveals several important scheduling characteristics. Within each sprint, frontend and backend development tasks are executed in parallel, with the frontend tasks typically starting 2-4 days after the backend tasks to allow the backend API contracts to be established first. This parallel development approach maximizes team productivity by enabling frontend and backend developers to work concurrently rather than sequentially. Additionally, the system analysis and architecture design tasks in Sprint 0 overlap with the literature review, reflecting the team's practice of translating research insights into architectural decisions in real-time rather than waiting for the literature review to be fully completed.

2.3 Risk Management:

Risk management is an essential process for identifying, assessing, and mitigating potential threats to the project. Effective risk management ensures that potential problems are anticipated and addressed proactively rather than reactively, minimizing their impact on project timeline, quality, and deliverables. The following risk register documents all identified risks, their category, probability of occurrence, potential impact level, and corresponding mitigation strategies.

Table 1: Risk Management

ID	Category	Risk	Score	Mitigation Plan	Owner	KPI
T-02	Technical	LLM API failure (OpenRouter) halts AI documentation and analysis generation	16 (Critical)	Auto-fallback between LLM providers (5 retries); 48h response cache; backup models: Llama 3.3, Mistral,	AI Team	API success rate & response time

Gemma						
S-01	Security	API key leakage in source code (OPENROUTER_API_KEY, SERVICE_API_KEY)	16 (Critical)	Move all secrets to env vars / Vault; GitLeaks scan on every commit; rotate keys every 90 days	Security Team	Key leak scan results & exposed key count
S-02	Security	Open CORS on UPM service allows XSS and CSRF attacks	16 (Critical)	Restrict CORS to allowed domains; enable CSRF Protection with Trusted Origins; apply CSP headers	Security Team	Active CORS policy & vulnerability scan results
T-01	Technical	Primary server failure causing complete service outage for all users	12 (High)	Daily automated backups; hot standby server; automated health checks every 5 min; DRP with RTO < 1 hour	Infra Team	Uptime % & mean recovery time
S-03	Security	JWT access tokens with 30-day lifespan widen exploitation window	12 (High)	Shorten access token to 15 min, refresh to 7 days; implement Token Blacklisting; add 2FA for sensitive operations	Security Team	Avg token lifetime & suspicious session count
O-04	Operational	No file upload size limit exposes server to resource exhaustion and DoS attacks	12 (High)	Set limits: 50 MB/file, 200 MB/project; apply rate limiting on upload endpoint; monitor storage usage	Infra Team	Upload sizes & storage consumption
O-05	Operational	Exceeding LLM API rate limit (50 req/day) triggers 60-min block	12 (High)	Smart queue distributing requests throughout the day; cache repeated responses; API usage dashboard with alerts	AI Team	Daily request count & cache hit rate
P-01	Performance	Slow LLM response times cause significant delays in documentation generation	12 (High)	Optimize prompts; set request timeout with auto model switching; parallel processing for multiple files	AI Team	Avg response time & timeout rate
T-04	Technical	MongoDB data loss due to missing referential integrity or replication failure	10 (High)	Enable Replica Set with 3+ nodes; daily backup with weekly restore test; apply DB-level validation rules	DB Admin	Replication status & restore test results

Each identified risk is monitored throughout the project lifecycle, with the mitigation strategies implemented proactively. The risk register is reviewed and updated during each sprint retrospective, ensuring new risks are captured and existing risk assessments are validated against actual project experience. The team maintains a risk escalation procedure where any risk that transitions from potential to actual is immediately communicated to all team members and the project supervisor, with a corresponding adjustment to the sprint plan if necessary.

2.4 Team Structure and Task Distribution:

The AutoTest & DocGen Manager project is developed by a team of four undergraduate students from the Department of Software Engineering and Information Systems at the Faculty of Informatics Engineering. Each team member assumes specific roles and responsibilities while maintaining cross-functional awareness through regular communication and code reviews. The project is supervised by Dr. Riad Sonbol as Project Manager and Eng. Raghad Alhossny as Work Monitoring Supervisor.

Project title	AutoTest & DocGen Manager – Intelligent System for Automated Testing and Documentation		
Project start date	February _ 21 _ 2026		
Project finish date	May _ 1 _ 2026		
Project manager	Dr. Riad sonbol		
Project objectives	Develop an intelligent web-based platform for automated software testing and documentation generation. The system will analyze codebases, generate comprehensive tests, execute them with coverage reports, produce technical documentation and UML diagrams, detect outdated documentation, and support Arabic and English languages with CI/CD integration.		
Approach	1. Define project scope and objectives. 2. Analyze system requirements and specification. 3. Design system architecture and decomposition. 4. Develop the first increment, focusing on core features. 5. Test and validate functionalities of the first increment. 6. Develop the second increment: testing and documentation. 7. Integrate CI/CD and perform comprehensive testing. 8. Document all development phases and outcomes.		
Roles and responsibilities:	Name	Role	Responsibility
	Dr.Riad sonbol	Supervisor	Highly project management and work monitoring
	Eng. raghad al Hossny	Supervisor	Work monitoring
	Areej noor aldeen	Engineer	Frontend/Backend Development- Documentation & AI
	Kamar aldiab	Engineer	Frontend/Backend Development Documentation & AI
	Wiaam alouni	Engineer	Frontend/Backend Development - Documentation & AI
	Ola najibeh	Engineer	Frontend/Backend Development - Documentation & AI

2.5 Quality Management in the Project:

Quality management is a critical aspect of the AutoTest & DocGen Manager project, ensuring that the delivered system meets the defined requirements and provides reliable, accurate outputs. Given that the system itself is designed to improve software quality through automated testing and documentation, it is essential that the platform demonstrates exemplary quality standards in its own development and operation. The quality management approach encompasses multiple dimensions including code quality, testing thoroughness, documentation accuracy, and user experience.

2.5.1 Code Quality Assurance:

Code quality is maintained through a combination of automated tools and manual review processes. All code is subject to mandatory peer review before merging into the main branch, with at least one team member reviewing each pull request for correctness, readability, and adherence to coding standards. Code reviews are conducted using GitHub's pull request review system, where reviewers can comment on specific lines, suggest changes, and approve or request modifications before the code is merged. Additionally, static analysis tools are integrated into the CI pipeline to automatically detect potential issues such as code style violations, unused imports, potential bugs, and security vulnerabilities.

2.5.2 Testing Quality Assurance:

The project employs a multi-layered testing strategy to ensure comprehensive quality coverage. Unit tests are written for all core modules, targeting a minimum of 80% code coverage. Integration tests validate the interaction between components, particularly the interface between the AI processing pipeline and the web application backend. System tests verify end-to-end functionality, including code upload, analysis, execution, documentation generation, and report download workflows.

2.5.3 Documentation Quality Assurance:

Since the AutoTest & DocGen Manager is a documentation generation platform, the quality of its own output serves as a direct measure of system effectiveness. Documentation quality is evaluated through several metrics: accuracy of generated API documentation compared to actual code behavior, completeness of UML diagrams in representing class relationships and interactions, readability and correctness of natural language code explanations, and currency of documentation (ensuring outdated detection mechanisms function correctly). The VerifierAgent in the multi-agent architecture provides automated quality assurance by cross-referencing generated documentation against actual code structure, detecting and flagging inconsistencies for correction.

2.6 Project Management Tools:

Effective project management requires appropriate tooling to support planning, tracking, communication, and collaboration. The AutoTest & DocGen Manager team utilizes a carefully selected set of tools that support the Agile Scrum workflow while addressing the specific needs of a distributed development team working on an AI-powered software platform.

Table 2: Project Management Tools

Tool	Category	Purpose	Usage Frequency
MS Project	Project Planning	Gantt chart creation; task scheduling; dependency management; resource allocation	Sprint planning
GitHub Projects	Project Management	Sprint backlog management; task tracking; milestone tracking	Daily
GitHub	Version Control	Source code repository; pull requests; code reviews; issue tracking; release management	Daily
GitHub Actions	CI/CD	Automated testing on push; linting; build verification; deployment automation	Per commit
ClickUp	Task Management	Personal task boards; meeting action items; quick note capture	Weekly
WhatsApp Group	Communication	Quick team communication; urgent notifications; informal coordination	Daily

The selection of tools prioritizes integration and accessibility. MS Project is used for detailed Gantt chart planning and timeline visualization, while GitHub Projects serves as the day-to-day project management tool, providing native integration with the code repository and CI/CD pipeline. This integration enables automatic status updates when pull requests are created, code is reviewed, or deployments succeed, reducing manual tracking overhead and ensuring the project board always reflects the current state of development.

2.7 Monitoring and Continuous Evaluation Mechanisms:

Continuous monitoring and evaluation are essential for ensuring the project remains on track and delivers quality outcomes. The AutoTest & DocGen Manager implements several mechanisms for ongoing assessment and course correction throughout the development lifecycle, ensuring that potential issues are identified early and addressed before they escalate into significant problems.

2.7.1 Weekly Team Meetings:

The team conducts weekly meetings every Sunday, lasting approximately 60-90 minutes, to review progress, discuss challenges, and plan the upcoming week's work. These meetings follow a structured agenda that includes: progress review against sprint goals and timeline, discussion of technical challenges and collaborative problem-solving, task prioritization and assignment for the coming week, risk assessment and mitigation review, and preparation for upcoming milestone reviews. ensuring all decisions and action items are recorded and accessible to all team members.

2.7.2 Supervisor Reviews:

Regular supervisor reviews are conducted with Dr. Riad Sonbol and Eng. Raghad Alhossny at key milestone points throughout the project. These reviews provide an opportunity for external evaluation of project progress, technical approach, and deliverable quality. Supervisor feedback is incorporated into sprint planning, with specific action items derived from review comments tracked in the project management system. The sprint reviews at the end of each sprint serve as major review checkpoints, while informal reviews are conducted as needed for critical technical decisions.

2.8 Software Configuration Management (SCM):


Software Configuration Management (SCM) is the discipline of managing the evolution of software systems throughout their lifecycle, ensuring that changes are controlled, traceable, and reproducible. For the AutoTest & DocGen Manager project, SCM practices are implemented using Git as the version control system, hosted on GitHub, with a well-defined branching strategy and release management process. The SCM approach ensures that the main branch always remains in a stable, deployable state while enabling four team members to develop multiple features in parallel without conflicts.

2.8.1 Git Repository Structure:

The project source code was managed using a single GitHub repository containing both frontend and backend implementations. The repository was organized using multiple branches to separate frontend and backend development and to allow parallel work between team members. This structure helped maintain clarity, reduce merge conflicts, and ensure stable production releases.

Repository Name: Senior-Project1-ATDG

Repository URL: <https://github.com/olanajibah-ENG/Senior-project1-ATDG-.git>




2.8.2 Branching and Merging Strategy:

The project adopts a structured branching strategy based on the Feature Branch Workflow model, combined with elements of the Git Flow methodology. Each feature is developed in an isolated branch and merged into the integration branches (Feature/Backend or Feature/Frontend) via a mandatory Pull Request review before reaching the main production branch. The table below documents all 15 branches used throughout the project lifecycle:

Table 3: Branching and Merging Strategy

#	Branch Name	Purpose & Description
1	main	The stable production branch. Contains only fully tested, deployment-ready code. All final releases are tagged here.
2	development	Primary integration branch. All completed backend and frontend features are merged here before moving to production.
3	Feature/Backend	Integrates all backend-related features and changes before merging into the development branch.
4	Feature/Frontend	Integrates all frontend-related features and changes before merging into the development branch.
5	Feature/Large Project Processing	Handles end-to-end processing of complete codebases: ZIP ingestion, static analysis, dependency mapping, UML generation, and dual-level AI documentation (high-level and low-level).
6	Feature/Conflict Detection	Implements logic to detect inconsistencies between code segments and discrepancies between the codebase and existing

		documentation using AI and static analysis.
7	Feature/GitHub Integration	Handles backend logic for fetching and importing complete project source code directly via a GitHub repository URL, eliminating manual file uploads.
8	Feature/Reviewer	Implements backend logic and RBAC for the Reviewer role: secure API endpoints for accessing analysis results, evaluating documentation, and submitting approvals or feedback.
9	Feature/Role	Implements core RBAC backend logic: defines user privileges (Admin, Developer, Reviewer) and middleware to restrict access based on assigned roles.
10	Feature/Backend-version2	Implements core backend for ZIP-based codebase ingestion and GitHub URL fetching, plus Admin & Reviewer Management Service with RBAC.
11	Feature/Admin UI	Develops the frontend dashboard for the Administrator role: user management, role assignment, system-wide project monitoring.
12	Feature/Reviewer UI	Develops frontend interfaces for the Reviewer role: viewing analysis results, UML diagrams, documentation, and submitting approvals or feedback.
13	Feature/Frontend-version2	Builds frontend for enhanced code ingestion (ZIP upload + GitHub URL) and the comprehensive Admin & Reviewer management dashboard.
14	Feature/Conflict UI	Develops frontend interfaces for presenting project inconsistencies: interactive screens highlighting code contradictions and documentation discrepancies.
15	Feature/Product/testing	Final integration and staging branch. Consolidates all frontend and backend features into a complete release candidate for QA before production deployment.



2.8.3 Team Members and Responsibilities:

The project was developed by a team of four undergraduate students, each assuming full-stack responsibilities across both frontend and backend layers. The team followed a strict peer-

review process using Pull Requests to maintain code quality and shared ownership. The table below summarizes each member's primary contributions:

Team Member	Role	Responsibilities
Areej	Backend & Frontend Developer	Backend core for Role management; Conflict Detection logic; Frontend: Reviewer dashboard, complex project processing, GitHub integration.
Kamar	Backend & Frontend Developer	Frontend: Admin dashboard, complex project processing, GitHub integration; Backend: Conflict Detection feature.
Ola	Backend & Frontend Developer	Backend core for complex project processing; Frontend: Conflict Detection interfaces.
Wiaam	Backend & Frontend Developer	Backend: codebase ingestion via ZIP upload and GitHub URL integration; core backend for Conflict Detection.

2.8.4 Collaborative Workflow — Pull Requests:

The repository was managed with high efficiency by assigning frontend and backend responsibilities to the respective developers. By implementing a Pull Request workflow for merging features into the development environment, the team successfully automated change tracking and ensured code integrity and seamless integration prior to deployment. The table below documents the key Pull Requests executed during the project:

#	Source Branch	Target Branch	Purpose
PR-1	Feature/Role	Feature/Backend	Merge RBAC role logic into backend integration branch
PR-2	Feature/Admin UI	Feature/Frontend	Merge Admin dashboard into frontend integration branch
PR-3	Feature/Graph-cross-analyze	Feature/Large Project Processing	Merge cross-file dependency graph analysis module
PR-4	Feature/Upload-folder	Feature/Large Project Processing	Merge ZIP upload and folder ingestion pipeline

1) Pull Request of merge **Feature/Role** To **Feature/Backend** :

Merge feature/Role into feature/Backend #6

Merged Areejno0r merged 1 commit into feature/Backend from feature/Role on Mar 20

Conversation 0 Commits 1 Checks 0 Files changed 41

Areejno0r commented on Mar 20

No description provided.

finish feature Role

Areejno0r requested review from WiaamAlouni, kamaral diab3-hash and olanajibah-ENG 2 months ago

Areejno0r self-assigned this on Mar 20

Areejno0r merged commit 45d988e into feature/Backend on Mar 20

Revert

Pull request successfully merged and closed

You're all set — the feature/Role branch can be safely deleted.

Delete branch

Reviewers: WiaamAlouni, kamaral diab3-hash, olanajibah-ENG

Assignees: Areejno0r

Labels: None yet

Projects: None yet

Milestone: No milestone

Development: Successfully merging this pull request may close these issues.

2) Pull Request of merge **Feature/Admin UI** To **Feature/Frontend** :

Admin UI done #9

Merged kamaral diab3-hash merged 2 commits into feature/Frontend from AdminUI on Mar 25

Conversation 0 Commits 2 Checks 0 Files changed 4

kamaral diab3-hash commented on Mar 25

No description provided.

Admin UI done

kamaral diab3-hash requested a review from Areejno0r 2 months ago

kamaral diab3-hash self-assigned this on Mar 25

Merge branch 'feature/Frontend' into AdminUI

kamaral diab3-hash merged commit 6494939 into feature/Frontend on Mar 25

Revert

Pull request successfully merged and closed

You're all set — the AdminUI branch can be safely deleted.

Delete branch

Reviewers: Areejno0r

Assignees: kamaral diab3-hash

Labels: None yet

Projects: None yet

Milestone: No milestone

Development: Successfully merging this pull request may close these issues.

Notifications: Subscribe

You're not receiving notifications from this thread.

3) Pull Request of merge **Feature/Graph-cross-analyze** TO **Feature/Large Project Processing** :

Feature/large project #12

Merged Areejno0r merged 6 commits into feature/Backend from feature/large-project on Mar 26

Conversation 0 Commits 6 Checks 0 Files changed 65

olanajibah-ENG commented on Mar 26

No description provided.

olanajibah-ENG added 6 commits 2 months ago

add all file

add all file

add all file

add all file

add all file

add all file

olanajibah-ENG requested review from Areejno0r, WiaamAlouni and kamaral diab3-hash 2 months ago

olanajibah-ENG self-assigned this on Mar 26

Areejno0r merged commit eeb86b into feature/Backend on Mar 26

Revert

Pull request successfully merged and closed

You're all set — the feature/large-project branch can be safely deleted.

Delete branch

Reviewers: WiaamAlouni, kamaral diab3-hash, Areejno0r

Assignees: olanajibah-ENG

Labels: None yet

Projects: None yet

Milestone: No milestone

Development: Successfully merging this pull request may close these issues.

Notifications: Unsubscribe

You're receiving notifications because you authored the thread.

2 participants

2.8.5 Development Workflow:

The development workflow defines the end-to-end lifecycle of every feature, from initial task assignment through final production release. All team members followed this process consistently throughout the project to maintain quality and traceability:


1 — Task Assignment: Tasks are divided into features (V1, V2, V3) and assigned to specific team members through the GitHub Projects sprint backlog.

2 — Feature Development: Each member works on a dedicated feature branch (e.g., Feature/Backend-version2) in isolation from the main codebase.

3 — Pull Request & Review: Once a feature is complete, a Pull Request is opened. A designated reviewer checks the code for correctness, logic, and style compliance before approval.

4 — Merging: After approval, the feature branch is merged into the corresponding development branch (Feature/Backend or Feature/Frontend).

5 — Final Release: Once all features are stable and integration-tested, the development branches are merged into the main production branch and tagged with a version release.



2.8.6 Tags and Versioning:

Git Tags were utilized to mark specific milestones and stable points in the project's development history. Each tag represents a completed, tested version of the system with specific functionalities fully integrated. Tags follow a semantic versioning convention (MAJOR.MINOR.PATCH) with separate tags for backend and frontend components, enabling independent rollback and deployment of each layer. The six releases produced during the project are documented in the table below:

Tag	Type	Description
R:V1.1.0-Backend	Backend	First stable backend release: core code analysis engine and initial CI/CD webhook integration.
R:V1.1.0-Frontend	Frontend	First stable frontend release: initial project management dashboard and analysis monitoring interface.
R:V1.2.0-Backend	Backend	Second backend release: complex multi-file project handling, cross-file dependency analysis, enhanced CI/CD integration.
R:V1.2.0-Frontend	Frontend	Second frontend release: multi-project management, advanced analysis configuration, CI/CD pipeline visualization.
R:V1.3.0-Backend	Backend	Third backend release: code coverage calculation, quality metrics reporting, real-time notification delivery.
R:V1.3.0-Frontend	Frontend	Third frontend release: coverage visualization dashboard, quality metrics display, notification centre.


olanajibah-ENG / Senior-project1-ATDG-

[Code](#)
[Issues](#)
[Pull requests](#)
[4](#)
[Agents](#)
[Actions](#)
[Projects](#)
[Wiki](#)
[Security and quality](#)
[Insights](#)
[Settings](#)

[Releases](#)
[Tags](#)

Tags

V1.3.0-Frontend

3 hours ago
187467e
zip
tar.gz
Notes

V1.3.0-Backend

3 weeks ago
3b154df
zip
tar.gz
Notes

V1.2.0-Frontend

on Apr 1
00710d5
zip
tar.gz
Notes

V1.2.0-Backend

on Apr 1
9d78cd6
zip
tar.gz
Notes

V1.1.0-Backend

on Mar 26
4f8cd29
zip
tar.gz
Notes

V1.1.0-Frontend

on Mar 25
6494939
zip
tar.gz
Notes

R:V1.1.0-Backend

Latest

Compare

olanaibah-ENG released this now

V1.1.0-Backend

449c29

This Release is include components of UPM & AI-Project & Notification :

- Account User Service
- Project (CRUD) Service
- API Gateway for request routing
- Upload or copy code to analysis
- Create UML Diagram (class diagram)
- A logical Explanation on two mode high and low level
- Export a document containing all of the above

R:V1.1.0-Frontend

Latest

Compare

olanaibah-ENG released this 4 minutes ago

V1.1.0-Frontend

6494939

This Release is include components UI of UPM & AI-Project & Notification :

- Account User Service
- Project (CRUD) Service
- API Gateway for request routing
- Upload or copy code to analysis
- Create UML Diagram (class diagram)
- A logical Explanation on two mode high and low level
- Export a document containing all of the above

R:V1.2.0-Backend

Latest

Compare

olanaibah-ENG released this now

V1.2.0-Backend

9d78cd6

This Release is include components of ZIP & Github :

- Upload codebase via ZIP file or import via GitHub repository URL for analysis.
- Admin & Reviewer Management Service with custom access privileges.

R:V1.2.0-Frontend

Latest

Compare

olanaibah-ENG released this now

V1.2.0-Frontend

00718d5

This Release is include components UI of ZIP & Github :

- Upload codebase via ZIP file or import via GitHub repository URL for analysis.
- Admin & Reviewer Management Service with custom access privileges.

R:V1.3.0-Backend

Latest

Compare

olanaibah-ENG released this now

V1.3.0-Backend

3b154d7

This Release is include components of conflict-detection-foundation:

- Code VS Code.
- Code VS Documents.
- Fully.

R:V1.3.0-Frontend

Latest

Compare

olanaibah-ENG released this now

V1.3.0-Frontend

187467e

This Release is include components UI of conflict-detection-foundation:

- Code VS Code.
- Code VS Documents.
- Fully.

2.8.7 Summary:

The SCM practices adopted by the AutoTest & DocGen Manager team — structured branching, mandatory Pull Request reviews, semantic versioning, and continuous integration — provided a robust framework for coordinating parallel development across four team members over four sprints. The strict separation of concerns between feature branches, integration branches, and the production branch ensured that the main branch remained stable at all times, while the tagging strategy provided clear traceability between delivered releases and their corresponding sprint milestones. These practices collectively contributed to the high code quality and on-schedule delivery of the project.

Chapter 3: Literature Review

3.1 Introduction:

This chapter provides a comprehensive review of existing literature, commercial systems, and academic research related to automated code documentation, testing, and intelligent code analysis. The review is structured to fulfill the project requirements by covering:

1. Similar commercial and practical systems.
2. Recent academic research papers (2020–2025).
3. Modern technologies and frameworks used in the field.
4. Comparative analysis.
5. Research gap and proposed innovation.

3.2 Modern Technologies and Frameworks in the Field:

The landscape of automated code documentation and testing has been profoundly transformed by recent advancements in artificial intelligence, web development frameworks, and code analysis tools. This section provides an overview of the key technologies and frameworks that form the technical foundation of the AutoTest & DocGen Manager project, placing them in the context of current industry practices and academic research.

Large Language Models (LLMs): Large Language Models represent the most significant technological shift in code understanding and generation. Models such as OpenAI's GPT-3.5 and GPT-4, Google's PaLM 2, and Meta's Code Llama have demonstrated remarkable capabilities in code comprehension, summarization, and generation tasks. These models are trained on vast corpora of source code and natural language, enabling them to understand complex programming patterns, infer developer intent, and produce human-readable documentation from code artifacts. In the context of automated documentation, LLMs serve as the core reasoning engine, capable of transforming abstract syntax trees and code structures into coherent technical explanations. Recent research has shown that combining multiple specialized LLM prompts in a multi-agent architecture significantly outperforms single-model approaches, achieving higher accuracy in code summarization tasks while reducing hallucination rates.

Django Web Framework: Django is a high-level Python web framework that encourages rapid development and clean, pragmatic design. It follows the Model-View-Template (MVT) architectural pattern and provides built-in features including an Object-Relational Mapping (ORM) system, authentication middleware, and administrative interfaces. For the AutoTest & DocGen Manager, Django serves as the backend framework, handling API endpoint management, user authentication, file upload processing, and coordination between the frontend interface and the AI processing pipeline. Django's robust security features, including built-in protection against CSRF, XSS, and SQL injection, make it particularly suitable for enterprise-grade applications handling sensitive code repositories. Its extensibility through middleware and signals allows seamless integration with external AI services and asynchronous task processing.

React Frontend Library: React is a declarative, component-based JavaScript library developed by Meta for building dynamic user interfaces. Its virtual DOM implementation enables efficient rendering of complex, data-intensive interfaces with real-time updates. In the AutoTest & DocGen Manager, React powers the interactive dashboard where users upload code repositories, monitor analysis progress, review generated documentation, and visualize UML diagrams. React's

component architecture allows for modular UI development, where each element — from the project management panel to the code coverage visualization chart — is implemented as an independent, reusable component. State management solutions such as React Context and hooks facilitate the real-time data flow required for displaying live analysis progress and test execution results.

Tree-sitter and Abstract Syntax Tree (AST) Parsing: Tree-sitter is an incremental parsing system that generates concrete syntax trees for source code with error recovery capabilities. It supports multiple programming languages including Python, Java, and JavaScript, making it the ideal tool for multi-language code analysis. Combined with native AST modules (such as Python's `ast` module), Tree-sitter enables the AutoTest & DocGen Manager to extract detailed structural information from source files, including class hierarchies, function signatures, variable scoping, and control flow patterns. This structural extraction forms the foundation upon which documentation generation and test case identification are built, as accurate code understanding requires precise parsing of syntax and semantics.

Mermaid.js Diagramming Library: Mermaid.js is a JavaScript-based diagramming and charting tool that uses text descriptions to generate diagrams dynamically. It supports class diagrams, sequence diagrams, flowcharts, and entity-relationship diagrams through a simple markdown-like syntax. The AutoTest & DocGen Manager integrates Mermaid.js to automatically render UML class diagrams and sequence diagrams from analyzed code structures, providing visual representations of system architecture that complement the generated text documentation. This integration allows developers to visualize code relationships and interaction flows without requiring manual diagram creation in external tools.


MongoDB and MySQL Dual Database Architecture: The project employs a dual database strategy, leveraging MySQL for structured relational data (user accounts, project metadata, configuration settings) and MongoDB as a document-oriented store for unstructured content (source code files, generated documentation, JSON analysis results). This hybrid approach optimizes query performance for structured operations while providing the flexibility needed to store variable-length code artifacts and AI-generated documentation of different formats.

3.3 Similar Systems Analysis:

We examined seven production systems that represent the spectrum of approaches to automated code documentation, intelligent testing, and code intelligence. These systems were categorized into three groups based on their primary focus: documentation-focused platforms, code intelligence platforms, and AI-powered testing tools. Each system was evaluated for its feature completeness, technological approach, strengths, and critical gaps relative to our project objectives. The following subsections present detailed analyses of each system.

3.3.1 DocuWriter.ai:

[DocuWriter.ai - AI Code Documentation Platform | Generate Docs from Source Code](#)




DocuWriter.ai is an AI-powered code documentation platform designed to automatically generate and maintain comprehensive documentation from source code files. The platform specializes in continuous documentation refresh, ensuring that documentation remains synchronized with code changes over time. It supports automatic generation of technical documentation for classes, functions, APIs, and code modules, and includes native support for Swagger/OpenAPI documentation generation in JSON format, making it particularly valuable for RESTful API projects.

The platform also provides AI-powered code test suite generation, intelligent code refactoring suggestions, and UML diagram generation capabilities. DocuWriter.ai integrates with popular workflow automation tools such as n8n and supports Git repository integration for continuous documentation updates. Its knowledge base management system allows teams to organize and export documentation in multiple formats including Markdown and PDF.

Despite its comprehensive feature set, DocuWriter.ai exhibits several notable limitations that our system addresses. It lacks deep project structure analysis capabilities, meaning it cannot extract and visualize the hierarchical relationships between code modules, classes, and functions. Its dependency mapping is limited, providing only superficial connections between code elements rather than comprehensive dependency graphs. The test generation capability is basic and does not include test execution, coverage reporting, or automated fixing of failing tests. Furthermore, the platform does not provide explicit outdated documentation detection mechanisms, relying instead on periodic manual or automated refresh cycles rather than intelligent change detection.

3.3.2 Swimm.io:

Swimm: Application Understanding Platform




Swimm.io is a code-coupled documentation platform that uniquely addresses the challenge of documentation staleness through automatic synchronization between documentation and code changes. Unlike traditional documentation tools that generate static documents, Swimm maintains a living link between documentation elements and their corresponding code artifacts, automatically detecting when code changes render documentation outdated and alerting the development team.

The platform supports direct documentation-to-code linking, continuous documentation workflows, team collaboration features, and deep IDE integration with Visual Studio Code and JetBrains suite. Swimm's approach to documentation treats documentation as a first-class citizen in the development workflow, integrating it directly into the code review process and ensuring that documentation changes are reviewed alongside code changes.

However, Swimm.io's focus on documentation synchronization comes at the expense of other critical capabilities. It does not provide automated test generation or execution, limiting its utility for teams that need comprehensive testing support alongside documentation. Its UML diagram generation is minimal, focusing primarily on textual documentation rather than visual representations of system architecture. The platform lacks AI-driven code summarization capabilities, meaning complex modules must be documented manually. Additionally, Swimm requires manual initial documentation creation rather than auto-generating documentation from existing code, which increases the initial setup burden for teams adopting the platform.

3.3.3 Sourcegraph:

[Sourcegraph](#) | The code intelligence platform for enterprises





Sourcegraph is an enterprise-grade code intelligence platform that provides advanced code search, navigation, and dependency analysis capabilities across large-scale codebases. It supports over 40 programming languages and excels in understanding complex project structures, making it an invaluable tool for organizations managing large monorepositories or distributed codebases across multiple teams and repositories.

Key capabilities include dependency graph visualization that maps the relationships between modules, packages, and external dependencies; multi-repository code search with advanced filtering and cross-reference detection; project tree analysis and exploration tools; batch code changes across multiple repositories; and enterprise-scale codebase search with indexing support for millions of lines of code. Sourcegraph's code intelligence engine provides deep insights into code relationships, including symbol navigation, definition references, and usage patterns across the entire codebase.

Sourcegraph's primary limitation from our project's perspective is its exclusive focus on code intelligence and exploration. It does not generate documentation automatically, requiring development teams to manually create and maintain documentation based on the insights it provides. Similarly, it lacks any testing capabilities, meaning teams must use separate tools for test generation and execution. While its dependency mapping and code structure analysis are exceptional, these capabilities serve human exploration rather than automated documentation generation, representing a fundamentally different use case from what our system aims to deliver.

3.3.4 Qodo.ai (formerly Codium):

AI Code Review – Qodo: Deploy with confidence





Qodo.ai is an agentic code integrity platform that leverages advanced AI techniques, including Retrieval-Augmented Generation (RAG) and agentic AI architectures, to provide intelligent test generation, code quality analysis, and automated code review capabilities. The platform represents the cutting edge of AI-powered development tools, employing sophisticated context-awareness mechanisms to understand the broader codebase context when generating tests and analyzing code quality.

Qodo.ai's key capabilities include intelligent test generation covering happy paths, edge cases, and rare execution scenarios; automated test execution with self-healing capabilities that can automatically fix failing tests when code changes; RAG-based codebase awareness that maintains deep contextual understanding of the project; agentic AI for autonomous decision-making in test strategy and code quality assessment; intelligent code refactoring that adheres to team coding standards; Git agent integration (Qodo Merge) for automated pull request reviews; and native IDE plugins for both Visual Studio Code and JetBrains environments.

Despite its advanced capabilities in testing and code quality, Qodo.ai has notable gaps in documentation-related features. It does not generate UML diagrams from code structure, limiting its ability to help teams visualize system architecture. Its documentation generation is limited to inline comments and docstrings, rather than comprehensive technical documentation. The platform also lacks explicit outdated documentation detection capabilities and is primarily IDE-focused, which constrains its utility for web-based team collaboration and CI/CD pipeline integration scenarios.

3.3.5 CodeSee Maps:

CodeSee – Bring visibility to your codebase




CodeSee Maps is a code visualization and exploration platform that automatically generates interactive project maps and dependency graphs for codebases. The platform enables developers to understand complex project architecture at a glance through visual codebase exploration, making it particularly valuable for onboarding new team members and navigating large, unfamiliar codebases.

The platform's key features include interactive codebase maps with project tree visualization that provides a bird's-eye view of code organization; automated dependency graph generation and exploration that reveals both explicit and implicit code dependencies; code tour creation for team onboarding documentation that guides new members through critical code paths; pull request impact analysis and change visualization that helps reviewers understand the full scope of code changes; multi-repository exploration with cross-project linking capabilities; and seamless integration with GitHub for continuous map updates as code evolves.

CodeSee Maps' limitations are significant from our project's perspective. It provides no automated documentation generation, meaning teams must manually create all project documentation. It lacks any test generation or execution capabilities, requiring integration with separate testing tools. The visualization is static in nature, meaning it does not include AI-driven code summarization or intelligent analysis of code behavior. Additionally, code tours require manual annotation rather than being AI-automated, which limits scalability for large projects with frequent code changes.

3.3.6 Mintlify:

Mintlify - The Intelligent Documentation Platform



Mintlify is an AI-native documentation platform that represents a modern approach to developer documentation by combining automatic documentation generation with beautiful, customizable portal hosting. The platform specializes in creating well-structured, searchable documentation portals from code repositories with minimal configuration, making it an attractive choice for teams that want to provide polished API documentation to external consumers.

Mintlify's feature set includes AI-powered documentation generation from source code that produces clean, well-formatted documentation; beautiful and customizable documentation portal hosting with modern UI design; built-in search and analytics for tracking documentation usage patterns; automatic API reference generation for REST and GraphQL endpoints; continuous sync with GitHub and GitLab repositories for automatic documentation updates; and multi-language SDK documentation support that caters to diverse development communities.

Despite its strengths in documentation presentation and hosting, Mintlify lacks several capabilities that are central to our project. It does not provide automated test generation or execution, focusing exclusively on the documentation domain. Its code structure analysis and dependency mapping capabilities are limited, providing insufficient depth for generating comprehensive architectural documentation. The platform does not generate UML diagrams from code structure, and it lacks the multi-agent AI architecture that our system employs for improved accuracy and self-correction in documentation generation.

3.4 Feature Comparison & Gap Analysis:

This section presents a comprehensive feature comparison across all seven reviewed systems, organized into four major capability categories: code structure analysis, documentation intelligence, testing capabilities, and AI integration. The comparison reveals distinct patterns in how existing systems address different aspects of the automated documentation and testing landscape, and highlights the specific gaps that the AutoTest & DocGen Manager is designed to fill.

Legend: $\triangle$ Full Support | $\triangle$ Partial Support | $\square$ Not Available | * Features to be completed this semester

Table 4: Feature Comparison

Feature / Capability	DocuWriter	Swimm.io	Sourcegraph	Qodo.ai	CodeSee	Mintlify	Our System
Project Tree Extraction	✗	$\triangle$	✓	✗	✓	✗	✓ *
Dependency Mapping	$\triangle$	✗	✓	✓	✓	✗	✓ *
Cross-file References	✗	✓	✓	$\triangle$	$\triangle$	✗	✓ *
Multi-language Parsing	$\triangle$	$\triangle$	✓	✓	$\triangle$	$\triangle$	✓ *
Auto Documentation	✓	✓	✗	✓	✗	✓	✓ *
Outdated Doc Detection*	$\triangle$	✓	✗	✗	✗	$\triangle$	✓ *
API Docs (Swagger)	✓	$\triangle$	✗	✓	✗	✓	✓ *
UML Diagram Generation	✓	✗	✗	✗	✗	✗	✓ *
Test Generation	✓	✗	✗	✓	✗	✗	✓
Test Execution & Fix	✗	✗	✗	✓	✗	✗	✓
Code Coverage Reports	✗	✗	✗	✓	✗	✗	✓
RAG / Context Awareness	✗	$\triangle$	✓	✓	$\triangle$	$\triangle$	✓ *
AI Code Summarization	$\triangle$	$\triangle$	✗	✓	✗	✓	✓ *
Multi-Agent Architecture	✗	✗	✗	$\triangle$	✗	✗	✓ *
CI/CD Integration	✓	✓	✓	✓	✓	✓	✓ *
Deployment Type	Cloud (SaaS)	Cloud + IDE Plugins (VS Code, JetBrains)	Self-Hosted / Cloud	IDE Plugins + Cloud	Cloud (SaaS) + GitHub App	Cloud (SaaS)	Web/Cloud

* indicates critical differentiating features

3.4.1 System Limitations Analysis:

The following table provides a consolidated analysis of each system's main strengths and critical gaps, along with an assessment of their relevance to the AutoTest & DocGen Manager project. This analysis informed our architectural decisions and feature prioritization strategy.

System	Main Strengths	Critical Gaps	Relevance
DocuWriter.ai	Continuous doc refresh, Swagger generation, UML diagrams	No project structure analysis, limited test generation, no outdated doc detection	Medium (60%) - Validates doc generation approach
Swimm.io	Outdated doc detection, code-doc sync, continuous workflow	No test generation, limited UML, manual doc creation required	Medium-High (65%) - Validates outdated detection approach
Sourcegraph	Deep code intelligence, dependency graphs, 40+ languages	No doc generation, no testing, manual documentation only	High (80%) - Complements our analysis module design
Qodo.ai	Advanced test generation, RAG technology, agentic AI	No UML diagrams, no outdated doc detection, IDE-focused only	High (75%) - Validates multi-agent approach
CodeSee Maps	Visual dependency mapping, project tours, GitHub integration	No automated doc/test generation, static visualization only	Medium-High (65%) - Informs mapping module design
Mintlify	AI-native doc generation, modern UI, API auto-generation	No test generation, limited code analysis, no multi-agent AI	Medium (60%) - Benchmark for documentation quality

3.5 Scientific Foundation: Research Papers Review:

Beyond the practical comparison of existing commercial systems, the architectural decisions of the AutoTest & DocGen Manager are firmly grounded in rigorous academic research. This section presents a systematic review of eight foundational research papers published between 2020 and 2025 in reputable venues. These papers tackle key technical challenges directly relevant to our system, including large-scale code differencing, graph-based program analysis, code-comment consistency detection, semantic merge conflict resolution, repository-level code understanding and editing, long-context code modeling, and cross-modal pre-trained representations for code.

AI-based or critical features that do not have a straightforward solution (papers-based comparison).

Table 5: Scientific Foundation: Research Papers Review

Paper ID	Paper Title	Publis her	Year	Methodology	Datase t	Findings	Input	Output	Approa ch	Evaluatio n	Result
1	HyperDiff: Computing Source Code Diffs at Scale	ACM ESEC /FSE	2023	Lazy AST decomposition + HyperAST DAG + Lazyfied GumTree	19 large softwa re project s	Achieves order-of-magnitude speedups with 99.3% identical diffs	HyperAST DAG, commit ASTs, metadata	Scalable diff/edit script	Lazy top-down & diff bottom-up matching	CPU time, memory, diff equivalence	$\times 12.7$ faster, $\times 226$ faster phases, $\times 4.5$ less memory
2	Detecting Code-Comment Consistency	arXiv	2024	Feature engineering + skip-gram + Random Forest	22 Java project s, 83k changes	Detects outdated comments with $>90\%$ precision	Old/new code & comments	Outdated comment classification	Code/comment relation features	Precision, recall, F1	$>90\%$ precision, works on Python
3	ConflictLens: An LLM-Based Method for Detecting Semantic Merge Conflicts	SEKE	2025	Two-stage LLM pipeline: static + dynamic	85 merge scenarios, 31 Java project s	High precision & recall for semantic conflicts	Base/left/right/merged code	Conflict classification + tests	Few-shot + CoT + test generation	Precision, recall, F1, accuracy	Precision 0.91, Recall 0.76, F1=0.82
4	CodePlan: Repository-level Coding using LLMs and Planning	arXiv	2024	LLM-driven planning + dependency propagation	C# migration repos, Python editing repos	First effective multi-file LLM editing	Repo code + seed spec	Multi-step edit plan	Planning + adaptive propagation	Repair success	Solved 5/6 repos
5	RepoBench: Benchmarking Repository-Level Code Auto-Completion Systems	arXiv	2023	Retrieval + long-context modeling	10k Python + 15k Java train, 1k Python test	Long-context models outperform	Multi-file repos	Snippet retrieval + next-line prediction	RepoBench-R/C/P tasks	Accuracy @k, EM, EditSim	UniXcoder best retrieval
6	CodeBERT: Pre-trained Model for Code & NL	EMNLP	2020	MLM + RTD hybrid pretraining	CodeSearchNet	Strong code search & doc generation	NL-PL pairs	Code embeddings	Transformer bimodal pretraining	MRR, BLEU	SOTA on code search
7	UniXcoder: Unified Cross-Modal Pre-training for Code Representation	Microsoft Research	2022	Multi-modal pretraining (code+AST+comments), Mask Attention Matrices with Prefix Adapters for controlling model behavior (encoder/decoder/fine-tuning)	CodeSearchNet + custom	SOTA on many tasks	Code + AST + comments	Unified representation	Mask attention + prefix adapters	Clone detection, search	SOTA performance

3.5.1 Key Insights and Contributions of the Reviewed Papers:

After reviewing the eight papers summarized in Table 3-1, several critical insights were derived that directly shaped the design and implementation of the **AutoTest & DocGen Manager** system:

1. **Repository-Level Analysis and Planning** The paper **CodePlan (2024)** demonstrated the effectiveness of LLM-driven planning and dependency propagation for multi-file projects. This insight led to the adoption of **project-level analysis** and the construction of a **dependency graph** in our system.
2. **Advanced Code Representation** **UniXcoder (2022)** and **CodeBERT (2020)** highlighted the importance of multi-modal representations (code + AST + natural language). This influenced our hybrid parsing strategy using **Python AST** and **Tree-sitter** for Java and JavaScript.
3. **Semantic Verification and Consistency** **ConflictLens (2025)** and **Detecting Code-Comment Consistency (2024)** provided strong evidence for the need of a verification layer. This directly inspired the design of the **VerifierAgent**, which performs semantic validation and self-correction.
4. **Graph-Based Understanding** **Graph Neural Networks on Program Analysis (2021)** reinforced the value of heterogeneous program graphs, supporting our use of dependency graphs for topological code processing.

These papers collectively confirm that while individual techniques exist, **no unified framework** combines deep structural analysis, multi-agent collaboration, automated testing, and comprehensive documentation generation.

3.5.2 Research Gap:

Despite the valuable contributions of the reviewed papers, several important gaps remain:

- Most approaches focus on either code understanding or specific tasks (diffs, conflict detection, completion) but **lack an end-to-end system**.
- Limited support for **multi-agent collaboration** with self-correction mechanisms.
- Insufficient emphasis on **cost optimization** and production readiness (caching, model fallback, multi-language support).
- Few solutions integrate **automated test generation + execution** with high-quality documentation in a single web platform.

3.5.3 Proposed Innovation:

The AutoTest & DocGen Manager addresses these gaps by introducing:

- A novel **three-specialized-agent architecture** (HighLevelAgent, LowLevelAgent, VerifierAgent).
- **Topological Code Processing** based on dependency graphs.
- Hybrid parsing (AST + Tree-sitter) combined with intelligent 48-hour caching and 7-model fallback over OpenRouter.
- Full integration of structural analysis, test generation/execution, and multi-format documentation (PDF, Markdown, UML via Mermaid.js) in one unified web platform.

This integration represents a practical and innovative advancement built upon the strong theoretical foundations provided by the reviewed literature.

3.6 Current and Future Trends in the Field:

The field of automated software documentation and intelligent code analysis is evolving rapidly, driven by several converging technological trends that are reshaping how development teams interact with codebases. Understanding these trends is essential for positioning the AutoTest & DocGen Manager within the broader trajectory of software engineering innovation.

Generative AI for Software Engineering: The emergence of Generative AI has fundamentally altered the landscape of developer tooling. Large language models trained on code — including GitHub Copilot, Amazon CodeWhisperer, and Google's Gemini Code Assist — have moved beyond simple code completion to encompass entire workflow automation, including documentation generation, test creation, and code review. The trend is shifting from assistive tools that require developer oversight to autonomous agents capable of independently analyzing codebases and producing production-ready artifacts. This aligns directly with the AutoTest & DocGen Manager's multi-agent architecture, which represents the next evolutionary step: specialized AI agents that collaboratively process, verify, and refine documentation and test outputs without human intervention. Industry analysts predict that by 2027, over 70% of enterprise documentation will be AI-assisted or AI-generated, making early adoption of multi-agent documentation systems a strategic advantage.

LLM Advancements and Context Window Expansion: One of the most impactful recent developments is the dramatic increase in LLM context windows. Models such as Google's Gemini 1.5 Pro support up to 2 million tokens, while Anthropic's Claude supports up to 200K tokens. This expansion directly addresses the historical limitation of repository-level code understanding, where large codebases exceeded the processing capacity of earlier models. However, despite these improvements, challenges remain in maintaining accuracy and coherence when processing extremely large repositories, which is why techniques such as topological code processing and dependency-based context construction — as implemented in the AutoTest & DocGen Manager — continue to provide significant advantages over naive full-context approaches.

MLOps for AI-Powered Development Tools: Machine Learning Operations (MLOps) represents the standardization of ML model deployment, monitoring, and maintenance in production environments. In the context of AI-powered documentation tools, MLOps practices ensure that the underlying language models maintain consistent output quality, that prompt engineering changes are systematically tested, and that model performance degradations are detected before they impact documentation accuracy. The AutoTest & DocGen Manager incorporates MLOps principles through its intelligent caching system, verification agents, and output validation mechanisms, ensuring that generated documentation meets quality thresholds before being presented to users.

Edge AI and On-Premise Code Intelligence: As organizations increasingly prioritize data privacy and security, there is a growing demand for AI-powered tools that can operate on-premise without sending proprietary code to cloud-based APIs. Edge AI techniques enable language models to run locally on developer machines or private servers, addressing concerns about code confidentiality in industries such as finance, healthcare, and defense. While the current implementation of the

AutoTest & DocGen Manager relies on cloud-based LLM APIs, the modular architecture is designed to support future integration with locally-deployed models such as Code Llama or Mistral, enabling organizations with strict data governance requirements to benefit from automated documentation without compromising code privacy.

RAG-Based Code Understanding: Retrieval-Augmented Generation (RAG) has emerged as a dominant paradigm for enhancing LLM performance in domain-specific tasks. By combining a retrieval system that searches through code repositories and documentation databases with a generative model that produces contextual responses, RAG architectures achieve significantly higher accuracy than standalone LLMs. Tools such as Qodo.ai have demonstrated the effectiveness of RAG in test generation scenarios. The AutoTest & DocGen Manager's architecture is inherently compatible with RAG principles, as its topological code processing and dependency-aware context building serve as a specialized form of retrieval that ensures the LLM receives relevant, structured code context for documentation generation.

Autonomous Multi-Agent Systems: The trend toward autonomous AI agents that can independently plan, execute, and verify complex tasks represents the frontier of AI-powered development tools. Frameworks such as LangChain, AutoGen, and CrewAI enable the orchestration of multiple specialized agents that collaborate on sophisticated workflows. The AutoTest & DocGen Manager's multi-agent architecture — comprising specialized agents for high-level summarization, low-level code explanation, and output verification — is a direct implementation of this trend, positioning the system at the cutting edge of autonomous documentation generation.

Chapter 4: Requirements Modeling and Analysis

4.1 Introduction:

This chapter presents a comprehensive analysis of the requirements for the AutoTest & DocGen Manager system. The requirements engineering process follows a structured approach that includes feasibility assessment, stakeholder identification, requirements elicitation, and formal documentation using UML modeling techniques. The chapter is organized according to the project's Agile Scrum methodology, ensuring that all requirements are traceable, prioritized, and validated against the project objectives defined in Chapter 1. Each requirement has been carefully analyzed to ensure it contributes meaningfully to the system's core mission of automating testing and documentation generation through intelligent multi-agent AI architecture.

4.2 Feasibility Study:

A feasibility study was conducted to determine whether the AutoTest & DocGen Manager project is technically viable, economically justifiable, and operationally practical within the given constraints. The study evaluates four dimensions of feasibility: technical, economic, operational, and schedule feasibility. This assessment ensures that the project team can deliver a functional, high-quality system within the defined timeline and resource constraints, while meeting all stated objectives and success criteria outlined in the project charter.

4.2.1 Technical Feasibility:

The technical feasibility assessment confirms that all required technologies, frameworks, and tools are available and mature enough to support the development of the AutoTest & DocGen Manager system. The backend framework Django (Python) provides robust REST API support, built-in ORM with SQLite for structured data, and comprehensive security middleware including JWT authentication and CSRF protection. The frontend framework React offers a component-based architecture ideal for building interactive dashboards with real-time progress tracking. MongoDB serves as a document-oriented database for storing unstructured content including source code artifacts, AI-generated documentation, and analysis results. This hybrid database strategy (SQLite for relational data, MongoDB for document storage) is a well-established pattern in modern AI-powered web applications.

The AI components leverage OpenRouter as a unified API gateway to access large language models, including free-tier models that eliminate direct API costs. The multi-agent architecture (HighLevelAgent, LowLevelAgent, VerifierAgent) coordinates through a central LLMClient module that interfaces with OpenRouter's API. Asynchronous task processing is handled by Celery with Redis as a message broker, enabling long-running documentation generation and test execution tasks to proceed without blocking the main application. Tree-sitter and Python's native AST module provide multi-language code parsing for Python, Java, and JavaScript, while WeasyPrint, xhtml2pdf, and ReportLab enable server-side PDF generation from documentation content. All core dependencies are open-source, and the development team possesses the necessary technical skills validated through prerequisite coursework and prior project experience.

4.2.2 Economic Feasibility:

The economic feasibility analysis evaluates the total cost of development against the projected benefits. The project uses entirely open-source technologies (Django, React, MySQL, MongoDB, Tree-sitter, Mermaid.js) that carry zero licensing fees. The AI integration layer uses OpenRouter to access free-tier large language models, which means the primary AI processing capability is available at no cost. This eliminates what would otherwise be the single largest operational expense, as comparable AI API services typically charge \$50-100 per month for production-level usage. The only infrastructure costs are web hosting services (estimated at \$20-40 per month for cloud deployment), which are standard for any web application.

The system incorporates multiple cost-optimization mechanisms that further reduce operational expenses. An intelligent caching layer stores generated explanations and documentation summaries, avoiding redundant AI API calls when the same code is analyzed multiple times. The incremental analysis capability ensures that only modified files are re-processed, reducing both computational resources and processing time. The multi-agent pipeline implements a verification step that catches errors early, minimizing the need for expensive re-generation cycles. Together, these optimizations can reduce AI processing costs by an estimated 60-70% compared to a naive implementation without caching or incremental analysis.

The benefits of the system include significant time savings for development teams, with an estimated reduction of documentation effort from 40% of developer time to near zero for routine documentation tasks. Automated test generation can reduce QA cycle time by approximately 50%, while code diff and change detection features reduce manual code review overhead. For a team of 5 developers, the system can save an estimated 40+ hours per month in documentation and testing activities alone. The return on investment is therefore strongly positive, with the system paying for itself within the first month of team-wide adoption.

4.2.3 Operational Feasibility:

The system is designed with operational feasibility as a core consideration. The web-based architecture ensures cross-platform accessibility through standard web browsers without requiring specialized client software. The intuitive dashboard interface is designed to minimize the learning curve, targeting a maximum onboarding time of 30 minutes for non-technical users and 15 minutes for developers familiar with similar tools. The system supports standard project upload formats (ZIP archives, Git repository URLs), requiring no special configuration or setup from end users.

Administrative operations including user management, project organization, and report generation are accessible through a clean, role-based interface. The system's RESTful API design enables future integration with existing CI/CD pipelines and development workflows, ensuring that adoption does not disrupt current team processes but rather enhances them. The multi-language support (Python, Java, JavaScript) covers the most widely used programming languages in industry, maximizing the system's applicability across different development environments.

4.2.4 Schedule Feasibility:

The project timeline spans from February 21, 2026 to May 1, 2026, providing approximately 14 weeks for complete development, testing, and delivery. The work is organized into three two-week sprints following the Agile Scrum methodology adopted in Chapter 2. A pre-sprint literature review period (February 21 - March 7, 2026) ensures the team has a solid theoretical foundation before implementation begins. Each sprint delivers a functional increment, reducing the risk of late-stage integration failures.

Critical path analysis identifies the multi-agent AI pipeline, code parser integration, and test execution engine as the highest-risk schedule items. These components are allocated to the earliest sprint (Sprint 1) to maximize the buffer time available for debugging and refinement. The schedule includes two formal seminar presentations (March 14 and April 25, 2026) that serve as progress checkpoints, ensuring the project remains on track. Contingency time of one week (April 23-30) is reserved for final integration, bug fixes, and documentation polish.

Table 6: Schedule Feasibility

Feasibility Dimension	Assessment	Key Factors	Risk Level
Technical	Feasible	Mature open-source stack, available APIs, team skills validated	Low
Economic	Feasible	Open-source tools, low API costs, high ROI for teams	Low
Operational	Feasible	Web-based, intuitive UI, standard upload formats	Low
Schedule	Feasible with monitoring	14 weeks, 4 sprints, contingency buffer included	Medium

4.3 Requirements Gathering:

Requirements gathering for the AutoTest & DocGen Manager was conducted through a combination of methodologies appropriate for the Agile Scrum development approach adopted by the project team. The primary elicitation techniques included stakeholder interviews, competitive analysis of existing systems (detailed in Chapter 3), and analysis of academic research papers that identified the technical challenges and state-of-the-art solutions in automated code documentation and testing. These approaches were selected to ensure comprehensive coverage of both user needs and technical possibilities.

Semi-structured interviews were conducted with two key stakeholder groups: software development professionals and the project supervisors (Dr. Riad Sonbol and Eng. Raghad Al Hossny). The interviews with developers focused on identifying pain points in current documentation and conflict detection workflows, preferred interaction patterns, and desired output formats. The supervisor interviews provided academic perspective on research methodology,

evaluation criteria, and innovation requirements. With responses documented and analyzed to extract actionable requirements.

The competitive analysis of six existing systems (DocuWriter.ai, Swimm.io, Sourcegraph, Qodo.ai, CodeSee Maps, and Mintlify) provided a structured framework for identifying feature gaps and market opportunities. Each system was evaluated across multiple dimensions including documentation capabilities, testing features, code analysis depth, integration options, and deployment models. The gaps identified through this analysis directly informed the functional requirements, ensuring that the AutoTest & DocGen Manager addresses real deficiencies in the current tooling landscape rather than duplicating existing functionality.

Additionally, a review of eight academic research papers (HyperDiff, GNN on Program Analysis, Code-Comment Consistency, LLM-Based Semantic Merge, CodePlan, RepoBench, CodeBERT, and UniXcoder) provided the theoretical foundation for technical requirements, particularly those related to the multi-agent AI architecture, topological code processing, and self-correction mechanisms. The synthesis of all these sources produced a comprehensive, well-grounded set of requirements that balance practical user needs with academic innovation.

4.4 Stakeholder Analysis:

Stakeholder analysis is essential for ensuring that the AutoTest & DocGen Manager system addresses the needs and expectations of all parties affected by its development and deployment. The identification process followed a systematic approach, mapping all individuals and groups who have a direct or indirect interest in the system. Each stakeholder was categorized by their role, level of influence, and specific requirements, enabling the team to prioritize features and design decisions that maximize stakeholder satisfaction.

Stakeholder	Role	Key Needs & Expectations	Influence Level
Software Developers (Primary Users)	End users who upload code and use generated documentation	Fast code upload, accurate documentation, clear UML diagrams, conflict detection, responsive dashboard	High
Development Team Leads	Managers overseeing team workflows and code quality	Team collaboration, project management, quality metrics, CI/CD integration, progress tracking	High
Project Supervisors	Academic advisors ensuring research quality and methodology rigor	Novel AI architecture, academic contribution, proper documentation, rigorous evaluation methodology	High

4.5 Functional Requirements:

Functional requirements define the specific behaviors, functions, and capabilities that the AutoTest & DocGen Manager system must deliver. These requirements were derived from the stakeholder analysis, competitive gap assessment, and research findings presented in the preceding sections. Each requirement is uniquely identified with an FR (Functional Requirement) prefix and

categorized by the system module it addresses. The requirements are organized into six functional domains that correspond to the system's core capabilities.

Table 7: Functional Requirements

Req_ID	Requirment Title	Actor	Usecase	Priority
RE-FR-01	The system must Generate comprehensive documentation (API docs, explanations, summaries) for all classes/functions	Developer/Reviewer	UC-03	High
RE-FR-02	The system must Produce an intelligent high-level summary of the project (architecture, main modules, dependency graph)	Developer	UC-03	Medium
RE-FR-03	The system must export analysis results in multiple formats (JSON, HTML, PDF, PlantUML, Mermaid).	Developer/Reviewer	UC-03	High
RE-FR-04	The system must allow users to select specific files/folders for analysis (instead of analyzing everything by default).	Developer	UC-04	High
RE-FR-05	The system shall allow the user to view a complete list of all existing versions for a project.	Developer	UC-04	High
RE-FR-06	The system shall allow the user to restore any previous version of the project.	Developer	UC-04	High
RE-FR-07	The system shall allow the user to delete an existing version.	Developer	UC-04	High
RE-FR-08	The system shall allow the user to rename or label versions (e.g., “v1.0”, “Initial Import”, “Post-Refactor”).	Developer	UC-04	High
RE-FR-09	The system should integrate with Continuous Integration tools such as GitHub, GitLab, and Bitbucket.	Github System	UC-05	Medium
RE-FR-10	The system must allow uploading a complete local folder or cloning a repository from GitHub/GitLab (recursive)	Developer	UC-05	High
RE-FR-11	The system must extract all files while preserving folder structure and excluding unsupported files (e.g. .git, node_modules)	Developer	UC-05	High
RE-FR-12	The system shall create a new version automatically whenever new code is uploaded or imported.	Developer	UC-05	High
RE-FR-13	The system shall maintain a continuous integration link with the connected GitHub repository after the initial import.	Github System	UC-05	High
RE-FR-14	The system shall automatically detect new commits pushed to the linked GitHub	Github	UC-05	High

	repository.	System		
RE-FR-15	The system shall automatically pull the updated source code when a new commit is detected.	Github System	UC-05	High
RE-FR-16	The system shall automatically create a new version entry for each synced update.	Developer	UC-05	High
RE-FR-17	The system shall notify the user when new updates are synced from GitHub.	Developer	UC-05	High
RE-FR-18	The system must Apply the Universal Language Processor + Tree-sitter / AST parsing to each file independently	Developer	UC-06	High
RE-FR-19	The system must Extract and aggregate classes, functions, variables, and imports from all files into a global symbol table	Developer	UC-06	High
RE-FR-20	The system must Detect and resolve dependencies across files (function calls, import resolution)	Developer	UC-06	High
RE-FR-21	The system must Merge Class Diagrams from individual files into a unified project-level diagram	Developer	UC-06	High
RE-FR-22	The system must Support re-analysis of only the modified file(s) with update of the global graph	Developer	UC-06	High
RE-FR-23	The system must Handle folders containing multiple programming languages (Python + Java + JavaScript + ...)	Developer	UC-06	High
RE-FR-24	The system must identifies the programming language used in the input code without the need for user intervention.	Developer	UC-06	High
RE-FR-25	The system must support for major programming languages: Python, Java, JavaScript, C++, C#, PHP, Ruby, Go, Rust, Swift.	Developer	UC-06	High
RE-FR-26	The system must converting the code into a syntax tree that represents the logical structure of the code.	Developer	UC-06	High
RE-FR-27	The system must maintain a persistent state of analysis results to enable incremental updates.	Developer	UC-06	High
RE-FR-28	The system must provide a web-based dashboard to visualize analysis results (UML diagrams, dependency graphs, metrics).	Developer	UC-06	High
RE-FR-29	The system must display analysis time and resource usage for performance monitoring.	Developer	UC-06	High
RE-FR-30	The system shall automatically trigger code	Developer	UC-06	High

	re-analysis after syncing the updated code.			
RE-FR-31	The system must compare two different versions of the same source code file.	Developer	UC-07	High
RE-FR-32	The system must identify the lines that were modified, deleted, or added.	Developer	UC-07	High
RE-FR-33	The system must display the differences using a clear and intuitive diff interface.	Developer	UC-07	High
RE-FR-34	The system must highlight additions, deletions, and modifications using different colors.	Developer	UC-07	High
RE-FR-35	The system must suggest documentation updates based on the detected changes.	Developer/ Reviewer	UC-08	High
RE-FR-36	The system must suggest regenerating the documentation when critical inconsistencies are found.	Developer/ Reviewer	UC-08	High
RE-FR-37	The system must Compare changes across files and identify parts of the documentation that require updating	Developer/ Reviewer	UC-08	High
RE-FR-38	The system shall allow the Admin to manage all projects, including viewing, editing, deleting, and accessing detailed project information and associated files.	Admin		High
RE-FR-39	The system shall allow the Admin to manage all projects, including viewing, editing, deleting, and accessing detailed project information and associated files.	Admin		High

4.6 Non-Functional Requirements:

Non-functional requirements specify the quality attributes, constraints, and performance standards that the AutoTest & DocGen Manager system must satisfy. These requirements ensure that the system is not only functionally correct but also reliable, secure, usable, and performant. The non-functional requirements are categorized into six domains: performance, security, usability, reliability, scalability, and compatibility.

Table 8: Non-Functional Requirements

Req ID	Category	Requirement Description
NFR-01	Performance	The system shall respond to user navigation actions within acceptable time
NFR-02	Performance	Code analysis for medium-sized projects

		shall complete within acceptable time
NFR-03	Performance	The dashboard shall provide real-time progress updates during processing
NFR-04	Security	User passwords shall be securely hashed using BCrypt algorithm
NFR-05	Security	The system must support Single Sign-On (SSO) in addition to standard user login and account creation.
NFR-06	Security	API keys and sensitive configuration shall be stored securely
NFR-07	Security	User-uploaded code shall be accessible only to authorized users
NFR-08	Usability	The system shall be accessible through standard web browsers without plugins
NFR-09	Usability	The interface shall be intuitive enough for new users to complete core tasks quickly
NFR-10	Usability	The system shall provide clear error messages with actionable guidance
NFR-11	Maintainability	The system must maintain a complete record of project versions and history, including all completed analysis, documentation, and test results.
NFR-12	Maintainability	The system should automatically contribute to the documentation of its own code and structure through its outputs (API docs, UML diagrams).
NFR-13	Availability	The system should maintain an uptime of 99.5% to ensure developers can access their documentation anytime.

4.7 Requirements Prioritization:

Requirements prioritization was performed using the **MoSCoW method** (Must, Should, Could, Won't) to ensure that the development team focuses on delivering the most critical features first. This prioritization was informed by the stakeholder analysis, the competitive landscape assessment, and the technical dependencies between requirements.

Requirements marked as "**Must**" represent the minimum viable product (MVP) features without which the system cannot fulfill its core purpose. "**Should**" requirements add significant value and are planned for inclusion if development capacity permits. "**Could**" requirements are desirable enhancements that improve the user experience but are not critical for the initial release

Priority Level	Count	Percentage	Sprint Allocation	Description
Must	28	0.72	Sprint 1 & 2	Core features required for system viability and MVP
Should	8	0.21	Sprint 2 & 3	Important enhancements that add significant value
Could	3	0.08	Sprint 3 (if capacity permits)	Desirable features that improve user experience
Won't	0	0	N/A	No requirements deferred

4.8 UML Diagrams:

UML (Unified Modeling Language) diagrams provide visual representations of the system's functionality, structure, and behavior. This section presents the key UML diagrams that model the AutoTest & DocGen Manager system from multiple perspectives: use case diagrams capture functional requirements from the user's viewpoint, and sequence diagrams detail the interactions between system components for critical operations.

4.8.1 Main Use Case Diagram:

The use case diagram provides a high-level view of the system's functional scope by identifying the actors (users and external systems) and their interactions with system functions. The AutoTest & DocGen Manager system has five actors: the Visitor (unauthenticated user who can register and sign in), the Developer (authenticated user who performs core operations such as project management, code analysis, documentation generation, and version comparison), the Administrator (user with elevated privileges for project management and project import), the GitHub System (external version control platform that provides source code for analysis), and the Reviewer (user who participates in documentation generation and evaluates documentation freshness). The diagram identifies eight main use cases that correspond to the system's primary functional requirements.




Figure 2: Use Case Diagram

4.8.2 Use Case Descriptions:

Each use case identified in the use case diagram is described in detail below, following a structured template that specifies the use case ID, name, actor, precondition, main flow, alternative flows, Exception Flows, and postcondition. The descriptions focus on the six core use cases that represent the system's primary value proposition.

Table 9: Import Project DESCRIPTION

Use Case Name Import Project	
Use Case ID	UC-05
Actors	Developer (Primary)
Description	This use case allows the developer to import their source code into the system via uploading a ZIP file, entering a GitHub repository URL, or uploading a Local Folder for further analysis.
Main Flow	1. The Developer selects "Import Project" from the system interface. 2. The system requests the user to choose the source (ZIP File Upload, GitHub Repository URL, or Local Folder). 3. The system processes the files and validates the size. 4. The system preserves the folder structure and excludes unnecessary files (e.g., .git, node_modules). 5. The system displays the project tree for the Developer to confirm.

	6. The system saves the project state and metadata in the Database. 7. The system displays a success message and triggers Code Analysis
Alternative Flows	A1: Upload ZIP File (Starts at Step 2): 1. The Developer uploads a source code ZIP file. 2. The system validates the ZIP file format and size. 3. The system securely extracts the contents of the ZIP file. 4. If the size is OK and extraction is successful, the flow continues to Step 4 in Main Flow. A2: Import via GitHub URL (Starts at Step 2) 1.The Developer inputs the GitHub repository URL into the system. 2.The system validates the URL and checks the repository size and accessibility via the GitHub API. 3.If valid and accessible, the system fetches and clones the source code files. 4.The flow continues to Step 4 in Main Flow A3: Upload Local Folder (Starts at Step 2): 1. The Developer selects and submits a local folder from their device. 2.The system validates the folder size locally. 3.If the size is OK, the flow continues to Step 4 in Main Flow.
Exception Flows	E1: Size Limit Exceeded (Step A1.2, A2.2, or A3.2): The system rejects the upload or fetch process and displays a "Size exceeds limit" error. E2: Invalid GitHub URL or Access Denied (Step A2.2): The system displays an error (e.g., "Repository not found or access denied") if the URL is incorrect or the repository is private without proper access. E3: Invalid File Format (Step A1.2): The system displays "Invalid file format. Please upload a valid ZIP file". E4: Empty Local Folder (Step A3.1): The system displays "The selected folder is empty or contains no supported code files" and asks the user to select again.
Post-conditions	On Success: The project is stored in the database, and the system proceeds to the analysis phase. On Failure: No project data is saved, and the user is notified of the specific error.

Table 10: Compare Code Version Description

Use Case Name Compare Code Version	
Use Case ID	UC-07
Actors	Developer (Primary)
Description	This use case allows the developer to compare two versions of the same file or project to identify differences such as added, deleted, or modified lines. The comparison is based on AST-level structural analysis to ensure accurate semantic diffing.
Main Flow	1. The Developer selects “Compare Code Version” from the system interface. 2. The system retrieves and displays the list of available versions (commits or saved snapshots). 3. The Developer selects Version A and Version B for comparison. 4. The system fetches the content of both versions from the database. 5. The system generates AST representations for both versions (if not already cached). 6. The system performs a structural comparison using the Diff Engine (AST-based diff). 7. The system identifies added, deleted, and modified code segments. 8. The system displays the diff result to the Developer using color-coded highlighting. 9. The system forwards the detected changes to UC-08 Evaluate Documentation Freshness for further processing.
Alternative Flows	A1: No Differences Found • If the two selected versions are identical, → The system displays: “No changes detected.” A2: Only One Version Available • If the project contains only one version, → The system displays: “Comparison requires at least two versions.”
Exception Flows	E1: Version Not Found • If one of the selected versions is missing or corrupted, → The system displays an error and requests the Developer to select another version. E2: AST Generation Failure • If the system fails to generate an AST for either version, → The system displays: “Unable to analyze this version.”
Post-conditions	On Success • The system displays a detailed diff view showing structural and textual changes. • The system sends the extracted changes to UC-08 for documentation freshness evaluation. On Failure • No comparison result is generated. • The system notifies the Developer of the specific error.

4.8.3 Sequence Diagrams:

Sequence diagrams detail the interactions between system components (objects) for specific use cases, showing the chronological order of messages exchanged.




Figure 3: Code Comparison Sequence Diagram




Figure 4: Import Project Sequence Diagram

4.8.4 Initial Test Case Descriptions:

Initial test case descriptions provide an early linkage between functional requirements and verification activities. These test cases will be expanded and detailed in Chapter 8 (Testing and Evaluation). Each test case specifies the linked requirement, a test description, the expected result, and the test type (functional, boundary, security, or performance).

Table 11: Initial Test Case Descriptions

Test ID	Linked Req	Test Description	Expected Result	Type	Category
TC-01	RE-FR-10	Upload a valid ZIP file containing a Python project.	Project created, files extracted, folder structure preserved.	Functional	Project Upload & Import
TC-02	RE-FR-10	Clone a valid GitHub repository URL.	Repository cloned, project created with source code.	Functional	Project Upload & Import
TC-03	RE-FR-11	Upload a ZIP containing .git and node_modules folders.	System excludes .git and node_modules, processes code files only.	Functional	Project Upload & Import
TC-04	RE-FR-04	Select specific files/folders for analysis.	Only selected files are analyzed, others skipped.	Functional	File & Folder Selection
TC-05	RE-FR-18/19	Analyze a Python project with multiple modules.	AST extracted, classes/functions/imports identified correctly.	Functional	Code Analysis & AST Parsing
TC-06	RE-FR-23/25	Analyze a project with cross-file function calls.	Global symbol table includes inter-file dependencies.	Functional	Code Analysis & AST Parsing
TC-07	RE-FR-24	Modify one file in a previously analyzed project and re-analyze.	Only modified file re-processed, global graph updated.	Functional	Code Analysis & AST Parsing
TC-08	RE-FR-25	Upload a project containing Python, Java, and JavaScript files.	All three languages analyzed correctly, results aggregated.	Functional	Multi-Language Support
TC-09	RE-FR-20	Upload a file without specifying language.	System auto-detects the language and processes correctly.	Functional	Multi-Language Support
TC-10	RE-FR-21	Upload source code in each of the 10 supported languages.	All languages parsed successfully with correct AST output.	Functional	Multi-Language Support
TC-11	RE-FR-22	Generate class diagrams from a multi-file project.	Unified project-level class diagram produced.	Functional	UML & Diagram Generation
TC-12	RE-FR-01	Generate documentation for an analyzed project.	Comprehensive API docs and explanations generated for all classes/functions.	Functional	Documentation Generation & Export
TC-13	RE-FR-02	Generate high-level project summary.	Architecture, modules, and dependency graph	Functional	Documentation Generation &

			summarized correctly.		Export
TC-14	RE-FR-03	Export documentation in JSON, HTML, PDF, PlantUML, and Mermaid formats.	All five export formats generated successfully and are valid.	Functional	Documentation Generation & Export
TC-15	RE-FR-31/32	Detect Conflict - Code vs Code / Code vs Documentation / Full Analysis.	Line-by-line differences correctly identified.	Functional	Conflict Detection & Diff Analysis
TC-16	RE-FR-33	View diff results on the diff interface.	Differences displayed with clear, intuitive layout.	Functional	Conflict Detection & Diff Analysis
TC-17	RE-FR-34	View diff results with color highlighting.	Additions in green, deletions in red, modifications in yellow.	Functional	Conflict Detection & Diff Analysis
TC-18	RE-FR-35	Modify code and trigger documentation update check.	System suggests specific documentation sections to update.	Functional	Code-Documentation Consistency
TC-19	RE-FR-36	Introduce critical code inconsistency.	System suggests full documentation regeneration.	Functional	Code-Documentation Consistency
TC-20	RE-FR-12	Upload new code to an existing project.	New version automatically created with timestamp.	Functional	Version Management
TC-21	RE-FR-05	View version history of a project.	Complete list of all versions displayed with labels.	Functional	Version Management
TC-22	RE-FR-06	Restore a previous version of a project.	Project reverted to selected version state.	Functional	Version Management
TC-24	RE-FR-14	Push a new commit to linked GitHub repository.	System detects the new commit within the polling interval.	Functional	GitHub Integration
TC-25	RE-FR-30	GitHub sync triggers re-analysis.	Incremental re-analysis performed on modified files only.	Functional	GitHub Integration
TC-26	RE-FR-38	Admin views list of all users.	All registered users displayed with roles and activity.	Functional	Admin Panel & User Management
TC-27	RE-FR-38	Admin deletes a project.	Project and associated files removed, user notified.	Functional	Admin Panel & User Management
TC-28	RE-FR-29	Admin Creates Reviewer Account.	Reviewer created successfully.	Functional	Admin Panel & User Management
TC-29	RE-FR-28	View analysis results on dashboard.	UML diagrams, dependency graphs, and metrics displayed.	Functional	Dashboard & Results View

4.9 Modeling According to Methodology (Agile Scrum):

As the AutoTest & DocGen Manager project follows the Agile Scrum methodology (defined in Chapter 2, Section 2.1), the requirements modeling approach incorporates Agile-specific artifacts including User Stories, Story Mapping, and a prioritized Product Backlog. These artifacts bridge the gap between traditional requirements documentation and the iterative, incremental delivery model of Scrum, ensuring that each sprint delivers measurable value to stakeholders.

4.9.1 User Stories:

User stories capture requirements from the perspective of end users, following the standard format: "As a [role], I want [feature] so that [benefit]." Each user story is written in simple, non-technical language that all stakeholders can understand, and includes acceptance criteria that define when the story is considered complete. Story points are assigned using a simplified Fibonacci scale (1, 3, 5, 8, 13) based on the estimated complexity and effort required for implementation.

Table 12: User Stories

Story ID	User Story	Points	Sprint	Acceptance Criteria
US-01	As a Developer, I want to upload a ZIP folder or clone a GitHub repository so that I can import my project into the system	5	Sprint 1	- ZIP upload works with progress bar - Git clone works with valid URL - Folder structure preserved - .git, node_modules, etc. are excluded
US-02	As a Developer, I want the system to automatically detect the programming language of each file so that I don't have to specify it manually	3	Sprint 1	- Correct detection for all 10 supported languages - Fallback mechanism when detection fails
US-03	As a Developer, I want the system to parse my code using AST / Tree-sitter so that I can see classes, functions, imports and dependencies	8	Sprint 1	- Global symbol table is built correctly - Cross-file dependencies resolved - AST parsing works for supported languages
US-04	As a Developer, I want to select specific files/folders for analysis so that I can focus only on relevant parts	3	Sprint 1	- Only selected items are analyzed - Clear visual feedback of selection scope
US-05	As a Developer, I want multi-language support so that I can analyze polyglot projects	5	Sprint 1	- Python, Java, JavaScript, C++, C#, PHP, Ruby, Go, Rust, Swift are supported
US-06	As a Developer, I want to view UML class diagrams and dependency graphs on the dashboard so that I can understand the architecture visually	5	Sprint 1	- Unified project-level class diagram - Interactive dependency graph - Dashboard loads within reasonable time
US-07	As a Developer, I want the system to generate comprehensive AI-powered documentation so that I don't have to write it manually	8	Sprint 2	- API documentation for classes & functions - High-level project summary - Explanations in natural language
US-08	As a Developer, I want to export	5	Sprint	- JSON, HTML, PDF, PlantUML,

	documentation in multiple formats so that I can use it in different tools	2		Mermaid - All exports are valid and readable
US-09	As a Developer, I want to compare two versions of a file and see a colored diff so that I can track changes	5	Sprint 2	- Line-by-line diff - Color coding (green/red/yellow) - Side-by-side view
US-10	As a Developer, I want the system to suggest documentation updates when code changes are detected so that documentation stays consistent	8	Sprint 2	- Change impact analysis - Specific suggestions for affected sections - Option for full regeneration on critical changes
US-11	As a Developer, I want the system to automatically sync with my linked GitHub repository so that analysis stays up-to-date	13	Sprint 2	- New commits detected - Incremental re-analysis of changed files only - New version created automatically - User notification
US-12	As a Developer, I want to manage project versions (view, restore, rename, delete) so that I can control project history	5	Sprint 3	- Version history list with timestamps & labels - Restore, rename, and delete functions work
US-13	As an Admin, I want to manage users (view, create, edit, delete, change roles) so that I can control access	5	Sprint 3	- Full CRUD on users - Role management (Developer, Reviewer, Admin) - Activity log
US-14	As an Admin, I want to manage all projects in the system so that I can oversee and maintain the platform	5	Sprint 3	- View all projects - Delete projects - Access project details and files

4.9.2 Story Mapping:

Story mapping organizes user stories into a two-dimensional grid that visualizes the user journey from left to right (representing the workflow sequence) and from top to bottom (representing priority from critical to nice-to-have). This technique helps the team understand the overall system scope, identify the minimum viable product (MVP), and plan release slices that deliver incremental value.

The story map for the AutoTest & DocGen Manager reveals a clear workflow progression. The horizontal axis follows the user journey: Code Upload, Language Detection, Code Analysis, Visualization, Documentation Generation, Code Diff, Smart Updates, GitHub Sync, Version Management, and Administration. The vertical axis separates Must-have stories (MVP) at the top from Should-have and Could-have stories below.

The MVP slice includes stories US-01 through US-07, covering the core value proposition: upload code, auto-detect language, perform AST analysis, select files, support ten languages, and visualize results on the dashboard. The second release slice adds US-08 through US-11, enabling documentation generation, export, code diff, smart update suggestions, and GitHub synchronization. The third release slice covers US-12 through US-15, adding version management,

notifications, and administration. This phased delivery approach ensures that the most valuable features are available earliest while progressively building toward the complete system.

4.9.3 Product Backlog:

The product backlog is the prioritized list of all work items needed to deliver the AutoTest & DocGen Manager system. It serves as the single source of truth for requirements and is continuously refined throughout the project lifecycle. Each backlog item is categorized as an Epic (large feature spanning multiple sprints), a Story (user-facing feature deliverable in one sprint), or a Task (technical work). Items are prioritized using the MoSCoW method and allocated to specific sprints based on dependencies and criticality.

Table 13: Product Backlog

Rank	Backlog Item	Type	Priority	Est. (SP)	Sprint
1	Code Upload and Import (ZIP + Git Clone)	Epic	Must	8	Sprint 1
2	Automatic Language Detection	Story	Must	3	Sprint 1
3	Multi-Language Code Parser (Tree-sitter/AST)	Epic	Must	13	Sprint 1
4	File/Folder Selection for Targeted Analysis	Story	Must	3	Sprint 1
5	Analysis Dashboard with UML & Dependency Graphs	Epic	Must	8	Sprint 1
6	Multi-Agent AI Documentation Pipeline	Epic	Must	8	Sprint 2
7	Multi-Format Documentation Export	Story	Must	5	Sprint 2
8	Code Diff & Change Detection	Epic	Must	5	Sprint 2
9	GitHub Integration & Auto Sync	Epic	Must	13	Sprint 2
10	Project Version Management	Epic	Must	5	Sprint 3
11	Admin User Management Panel	Story	Must	5	Sprint 3
12	Admin Project Management Panel	Story	Must	5	Sprint 3
13	Analysis Performance Metrics Display	Task	Could	3	Sprint 3

4.10 Requirements Analysis and Documentation:

This section presents the formal documentation artifacts that ensure complete traceability and verification of all requirements throughout the development lifecycle. The Requirements Traceability Matrix (RTM) establishes forward and backward traceability between requirements, design components, implementation modules, and test cases. This traceability ensures that every requirement is addressed in the system design and verified through testing, while also enabling impact analysis when requirements change.

4.10.1 Requirements Traceability Matrix (RTM):

The Requirements Traceability Matrix establishes forward and backward traceability for all identified requirements. Forward traceability ensures that each requirement is addressed in the system design, implemented in the appropriate code module, and verified through one or more test cases. Backward traceability allows the team to trace from any test case or implementation module back to its originating requirement, ensuring that all code and test artifacts serve a documented purpose. The RTM below presents a representative subset covering requirements from each functional category.

Table 14: Requirements Traceability Matrix (RTM)

req-id	Use cases	analysis	System design	Detailed design	coding	Test cases
RE-FR-01	UC-03	Generate Document	GD-BLL-03	DD-GDD	C-DG C-LLMDG	TC-12, TC-14
RE-FR-02	UC-03	Generate Document	GD-BLL-03	DD-GDD	C-DG C-LLMDG	TC-13
RE-FR-03	UC-03	Generate Document	GD-BLL-03	DD-GDD	C-DG C-LLMDG	TC-14,TC-15
RE-FR-04	UC-04	Project Management	PM-BLL-02			TC-07, TC-10
RE-FR-05	UC-04	Project Management	PM-BLL-02			TC-21, TC-22
RE-FR-06	UC-04	Project Management	PM-BLL-02			TC-23
RE-FR-07	UC-04	Project Management	PM-BLL-02			TC-24
RE-FR-08	UC-04	Project Management	PM-BLL-02			TC-24
RE-FR-09	UC-05	IP-AN	IP-BLL-02 IP-BLL-03			TC-18, TC-19
RE-FR-10	UC-05	IP-AN	IP-BLL-02			TC-01, TC-02
RE-FR-11	UC-05	IP-AN	IP-BLL-02 IP-BLL-03			TC-03
RE-FR-12	UC-05	IP-AN	IP-BLL-02 IP-BLL-03			TC-21
RE-FR-13	UC-05	IP-AN	IP-BLL-02 IP-BLL-03			TC-18
RE-FR-14	UC-05	IP-AN	IP-BLL-02 IP-BLL-03			TC-18, TC-19
RE-FR-15	UC-05	IP-AN	IP-BLL-02 IP-BLL-03			TC-18
RE-FR-16	UC-05	IP-AN	IP-BLL-02			TC-19

<u>IP-BLL-03</u>						
RE-FR-17	UC-05	<u>IP-AN</u>	<u>IP-BLL-02</u> <u>IP-BLL-03</u>			TC-18
RE-FR-18	UC-06	Genaral Code Analysis	<u>GCA-BLL-03</u>	<u>DD-MLD</u>	C-CA <u>C-PA</u> <u>C-MLA</u>	TC-08
RE-FR-19	UC-06	Genaral Code Analysis	<u>GCA-BLL-03</u>	<u>DD-MLD</u>	<u>C-CA</u> <u>C-PA</u> <u>C-MLA</u>	TC-08, TC-09
RE-FR-20	UC-06	Genaral Code Analysis	<u>GCA-BLL-03</u>	<u>DD-MLD</u>	<u>C-CA</u> <u>C-PA</u> <u>C-MLA</u>	TC-09
RE-FR-21	UC-06	Genaral Code Analysis	<u>GCA-BLL-03</u>	<u>DD-MLD</u>	<u>C-CA</u> <u>C-PA</u> <u>C-MLA</u>	TC-11
RE-FR-22	UC-06	Genaral Code Analysis	<u>GCA-BLL-03</u>	<u>DD-MLD</u>	<u>C-CA</u> <u>C-PA</u> <u>C-MLA</u>	TC-19, TC-20
RE-FR-23	UC-06	Genaral Code Analysis	<u>GCA-BLL-03</u>	<u>DD-MLD</u>	<u>C-CA</u> <u>C-PA</u> <u>C-MLA</u>	TC-06
RE-FR-24	UC-06	Genaral Code Analysis	<u>GCA-BLL-03</u>	<u>DD-MLD</u>	<u>C-CA</u> <u>C-PA</u> <u>C-MLA</u>	TC-05
RE-FR-25	UC-06	Genaral Code Analysis	<u>GCA-BLL-03</u>	<u>DD-MLD</u>	<u>C-CA</u> <u>C-PA</u> <u>C-MLA</u>	TC-06, TC-09
RE-FR-26	UC-06	Genaral Code Analysis	<u>GCA-BLL-03</u>	<u>DD-MLD</u>	<u>C-CA</u> <u>C-PA</u> <u>C-MLA</u>	TC-08
RE-FR-27	UC-06	Genaral Code Analysis	<u>GCA-BLL-03</u>	<u>DD-MLD</u>	<u>C-CA</u> <u>C-PA</u> <u>C-MLA</u>	TC-19, TC-20
RE-FR-28	UC-06	Genaral Code Analysis	<u>GCA-BLL-03</u>	<u>DD-MLD</u>	<u>C-CA</u> <u>C-PA</u> <u>C-MLA</u>	TC-11, TC-12, TC-13, TC-28
RE-FR-29	UC-06	Genaral Code Analysis	<u>GCA-BLL-03</u>	<u>DD-MLD</u>	<u>C-CA</u> <u>C-PA</u> <u>C-MLA</u>	TC-29
RE-FR-30	UC-06	Genaral Code Analysis	<u>GCA-BLL-03</u>	<u>DD-MLD</u>	<u>C-CA</u> <u>C-PA</u> <u>C-MLA</u>	TC-19, TC-20
RE-FR-31	UC-07	<u>CCV-AN</u>	<u>CCV-BLL-03</u>		C-CD	TC-16

RE-FR-32	UC-07	CCV-AN	CCV-BLL-03	C-CD	TC-16
RE-FR-33	UC-07	CCV-AN	CCV-BLL-03	C-CD	TC-16
RE-FR-34	UC-07	CCV-AN	CCV-BLL-03	C-CD	TC-16
RE-FR-35	UC-08	Evaluate Doc Freshness	EDF-BLL-03		TC-17
RE-FR-36	UC-08	Evaluate Doc Freshness	EDF-BLL-03		TC-17
RE-FR-37	UC-08	Evaluate Doc Freshness	EDF-BLL-03		TC-17
RE-FR-38					TC-21
RE-FR-39					TC-20

4.10.2 Initial Test Cases:

The initial test cases defined in Section 4.8.4 provide the foundation for the comprehensive test plan that will be developed in Chapter 8. These initial test cases ensure early validation of the system's core functionality and cover the primary user workflows including code upload, multi-language analysis, documentation generation, code diff, smart updates, GitHub synchronization, version management, and administration.

The initial test suite includes 30 test cases covering functional testing (27 cases), performance testing (1 case), and security testing (1 case). Functional tests verify that each requirement produces the expected behavior under normal conditions. Performance tests ensure that analysis operations complete within acceptable time frames. Security tests validate that the system is protected against common web vulnerabilities. Each test case is linked to one or more functional or non-functional requirements through the RTM, ensuring complete requirement coverage. The test cases will be expanded in Chapter 8 to include additional edge cases, integration tests, and end-to-end test scenarios based on the detailed system design developed in Chapters 5 and 6.

In summary, this chapter has established a rigorous requirements engineering framework for the AutoTest & DocGen Manager project. The feasibility study confirmed the project's viability across technical, economic, operational, and schedule dimensions. A total of 39 functional requirements were identified and organized into six categories, supported by 14 non-functional requirements spanning performance, security, usability, reliability, scalability, and compatibility. The requirements were prioritized using the MoSCoW method, with 72% classified as Must-have for the MVP. UML diagrams, user stories, story mapping, and a prioritized product backlog were developed following the Agile Scrum methodology. Finally, the Requirements Traceability Matrix ensures complete traceability from stakeholder needs through design, implementation, and testing artifacts, providing a solid foundation for the system design and implementation phases that follow.

Test Case Scenario: Sec-01: Project Upload & Import Test Cases

Table 15: Initial Test Cases

Test case id	Test case title	Test steps	Expected result
TC-01	Valid ZIP File Upload	1. Login as Developer 2. Go to Upload Project page 3. Upload a valid ZIP containing Python project	Project is created successfully, files extracted, and folder structure preserved
TC-02	GitHub Repository Clone	1. Login as Developer 2. Enter valid GitHub repository URL 3. Click "Clone Repository"	Repository is cloned, project created with all source code
TC-03	Upload ZIP with Excluded Folders	1. Prepare ZIP containing .git and node_modules 2. Upload the ZIP file	System excludes .git, node_modules and other unsupported folders automatically
TC-04	Code Upload Failure (Invalid File)	1. Attempt to upload invalid file (e.g. .exe or corrupted ZIP)	System rejects the file and shows clear error message

Test Case Scenario: Sec-02: Language Detection & Multi-Language Support Test Cases

Test case id	Test case title	Test steps	Expected result
TC-05	Automatic Language Detection	1. Upload files from different languages without specifying type	System correctly auto-detects Python, Java, JavaScript, etc.
TC-06	Multi-Language Project Analysis	1. Upload project containing Python + Java + JavaScript files 2. Start analysis	All languages parsed successfully, results aggregated
TC-07	Unsupported Language	1. Upload project written in unsupported language	System displays clear error message: "Programming language is currently unsupported"

Test Case Scenario: Sec-03: Code Analysis & AST Parsing Test Cases

Test case id	Test case title	Test steps	Expected result
TC-08	Successful AST Analysis (Python)	1. Upload valid Python project 2. Start analysis	Classes, functions, imports extracted correctly and global symbol table built
TC-09	Cross-File Dependencies	1. Upload multi-file project with inter-file function calls 2. Generate analysis	Cross-file dependencies resolved correctly
TC-10	Targeted File/Folder	1. Select specific files/folders	Only selected files are analyzed,

Analysis	2. Start analysis	others skipped
----------	-------------------	----------------

Test Case Scenario: Sec-04: Visualization & Dashboard Test Cases

Test case id	Test case title	Test steps	Expected result
TC-11	UML Class Diagram Generation	1. Complete project analysis 2. Open Dashboard → Class Diagram	Unified project-level class diagram rendered correctly
TC-12	Dependency Graph Visualization	1. Complete project analysis 2. Open Dependency Graph	Interactive dependency graph displayed with correct relationships
TC-13	Dashboard Overview	1. Complete full analysis 2. Navigate to main dashboard	All results (diagrams, summary, metrics) displayed properly

Test Case Scenario: Sec-05: Documentation Generation & Export Test Cases

Test case id	Test case title	Test steps	Expected result
TC-14	AI Documentation Generation	1. Analyze project 2. Click "Generate Documentation"	Comprehensive API docs, explanations and project summary generated
TC-15	Multi-Format Export	1. Generate documentation 2. Export in all formats	JSON, HTML, PDF, PlantUML, and Mermaid files generated and valid

Test Case Scenario: Sec-06: Diff Analysis & Version Control Test Cases

Test case id	Test case title	Test steps	Expected result
TC-16	Code Diff Between Versions	1. Upload new version of a file 2. Compare with previous version	Line-by-line diff shown with color highlighting (green/red/yellow)
TC-17	Documentation Update Suggestion	1. Modify code that affects documentation 2. Check suggestions	System suggests specific documentation sections to update

Test Case Scenario: Sec-08: Admin Panel Test Cases

Test case id	Test case title	Test steps	Expected result
TC-18	Admin User Management	1. Login as Admin 2. Create new Reviewer account	Reviewer account created successfully
TC-19	Admin Project Management	1. Login as Admin 2. Delete a user project	Project and associated data removed successfully

Chapter 5 : System Architecture Design

5.1 High-Level Design:

5.1.1 System Block Diagram:

The system block diagram provides a high-level visual representation of the major components of the system and their interactions. It illustrates how the Web Application communicates with the backend services through a centralized API Gateway, which routes requests to the appropriate microservices. The diagram clearly delineates the three core services—User & Project Management Service, Code Intelligence Service, and Notification Service—along with their respective data storage solutions. Each service operates independently, following the microservices architectural pattern, and communicates through well-defined API contracts.




Figure 5: System Block Diagram

1. **API Gateway (API-PL-01):**Function Serves as the single entry point for the entire system, responsible for routing all external web requests to the appropriate Microservice.
2. **User & Project Management Service (UPM-BLL-02):**Function A dedicated microservice specializing in managing user accounts, project settings, and project specifications.
3. **Code Intelligence Service (CI-BLL-03):**Function A microservice dedicated to processing and analyzing complex data (like reviews or code) using AI algorithms.
4. **Notification Service (NS-BLL-04):** Function A microservice responsible for processing and sending all types of alerts and notifications to users.
5. **SQL Database Access (SDA-DAL-05):**Responsible for communication with the SQL database of the system
6. **NoSQL Database Access (NSDA-DAL-06):**Responsible for communication with the NoSQL database of the system

As shown in the system block diagram, the architecture is organized into distinct layers. The Web Application layer represents the frontend interface that users interact with directly. Below it, the API Gateway serves as the centralized entry point that receives all incoming requests from the frontend, performs authentication and rate limiting, and routes each request to the appropriate backend microservice. The User & Project Management Service handles all user authentication, profile management, and project CRUD operations, persisting its structured data in a relational SQL database such as PostgreSQL. The Code Intelligence Service is responsible for code analysis, AI-powered code review, and code comparison features, storing semi-structured code analysis results and metadata in a NoSQL database such as MongoDB. The Notification Service manages real-time notifications, email alerts, and in-app messaging, and may rely on external services such as SMTP servers or push notification providers rather than a dedicated database.

5.1.2 Architectural Design Pattern:

The system employs two primary design patterns to ensure extensibility, maintainability, and separation of concerns:

1. Strategy Pattern — Multi-Language Code Processing:

The Strategy Pattern is applied to decouple language-specific parsing logic from the core analysis pipeline. The Project Analyzer component depends on the ILanguageProcessorStrategy interface rather than concrete implementations. At runtime, the Language Detector identifies the programming language based on file extensions and selects the appropriate processor (Java_Processor, Python_Processor, or JavaScript_Processor). This allows new language support to be added by implementing the interface without modifying existing code, adhering to the Open/Closed Principle.




Figure 6: Strategy Pattern — Multi-Language Processing Architecture

2. Template Method Pattern — Document Export:

The Template Method Pattern is used in the documentation export pipeline. The ReportGenerator class defines the common algorithm skeleton — including code structure analysis and AST extraction — while delegating the specific export format logic to its subclasses. Each output subclass (HTMLOutput, MarkdownOutput, PDFOutput, XMLOutput) implements its own export_document() method, enabling format-specific rendering without duplicating the shared analysis logic.




Figure 7: Template Method Pattern — Document Export Architecture

5.1.3 System Architecture Design:

This section provides a detailed textual description of the system architecture based on the microservices design pattern. The architecture is structured into four primary layers: the Presentation Layer (Web Application), the API Gateway Layer, the Service Layer (microservices), and the Data Layer (databases). Each layer has clearly defined responsibilities and communicates with adjacent layers through standardized protocols.




Figure 8: System Architecture

5.2 Database Design:

5.2.1 Entity-Relationship Diagram (ERD Diagram):

The ERD Diagram is a visual representation of the data entities, their attributes, and the relationships between them. It serves as the foundation for the database schema and helps ensure data integrity and efficient query performance. The diagram should identify all primary entities, their key attributes, and the cardinality of relationships (one-to-one, one-to-many, many-to-many).




Figure 9: Entity-Relationship Diagram (ERD Diagram)

5.2.2 Relational Schema:

The relational schema defines the actual table structure for the database. Each table should specify the column names, data types, and constraints. The schema should be normalized to at least the third normal form (3NF) to eliminate data redundancy and ensure data consistency.

Table 16: Relational Schema

Table Name	Column Name	Data Type	Constraints	Description
roles	RoleID	SERIAL	PRIMARY KEY	Role identifier
	RoleName	VARCHAR(50)	NOT NULL, UNIQUE	Admin, Developer, Reviewer
	RoleDescription	TEXT	-	-
	PermissionList	JSONB	-	List of permissions
users	UserID	SERIAL	PRIMARY KEY	User ID
	RoleID	INT	FOREIGN KEY (roles.RoleID)	User role
	Username	VARCHAR(100)	NOT NULL, UNIQUE	-
	PasswordHash	VARCHAR(255)	NOT NULL	Hashed password
	FullName	VARCHAR(150)	-	-
	Email	VARCHAR(150)	NOT NULL, UNIQUE	-
	SignupDate	TIMESTAMP	DEFAULT CURRENT_TIMESTAMP	-

projects	ProjectID	SERIAL	PRIMARY KEY	Project ID
	UserID	INT	FOREIGN KEY (users.UserID)	Project owner
	ProjectName	VARCHAR(200)	NOT NULL	-
	ProjectDescription	TEXT	-	-
	CreationDate	TIMESTAMP	DEFAULT CURRENT_TIMESTAMP	-
repositories	LastModifiedDate	TIMESTAMP	DEFAULT CURRENT_TIMESTAMP	-
	SourceType	VARCHAR(50)	-	ZIP / GitHub
	RepoID	SERIAL	PRIMARY KEY	GitHub Repository
	ProjectID	INT	FOREIGN KEY (projects.ProjectID)	Linked project
	RepoName	VARCHAR(200)	NOT NULL	-
folders	SourceType	VARCHAR(50)	-	GitHub / GitLab
	CreatedAt	TIMESTAMP	DEFAULT CURRENT_TIMESTAMP	-
	FolderID	SERIAL	PRIMARY KEY	Folder in project
	RepoID	INT	FOREIGN KEY (repositories.RepoID)	-
	ProjectID	INT	FOREIGN KEY (projects.ProjectID)	-
code_artifacts	ParentFolderID	INT	FOREIGN KEY (folders.FolderID)	For nested folders
	FolderName	VARCHAR(255)	NOT NULL	-
	FolderPath	VARCHAR(500)	NOT NULL	Full path
	CodeID	SERIAL	PRIMARY KEY	Source file
	FolderID	INT	FOREIGN KEY (folders.FolderID)	-
project_versions	FileName	VARCHAR(255)	NOT NULL	-
	UploadDate	TIMESTAMP	DEFAULT CURRENT_TIMESTAMP	-
	Language	VARCHAR(50)	-	Auto-detected
	Version	VARCHAR(50)	-	-
	FilePath	VARCHAR(500)	NOT NULL	-
	VersionID	SERIAL	PRIMARY KEY	Project versions
	ProjectID	INT	FOREIGN KEY (projects.ProjectID)	-
	VersionLabel	VARCHAR(100)	NOT NULL	v1.0, Initial, etc.
	CommitHash	VARCHAR(40)	-	If from Git
	CreatedAt	TIMESTAMP	DEFAULT CURRENT_TIMESTAMP	-

5.2.3 Constraints and Keys:

Database constraints play a critical role in maintaining data integrity, enforcing business rules, and ensuring referential consistency across the system. The following section details the primary keys, foreign keys, and other constraints implemented in the database design.

Primary Keys

Primary keys are used to uniquely identify each record in a table:

- roles.RoleID → PRIMARY KEY
- users.UserID → PRIMARY KEY
- projects.ProjectID → PRIMARY KEY
- repositories.RepoID → PRIMARY KEY
- folders.FolderID → PRIMARY KEY
- code_artifacts.CodeID → PRIMARY KEY
- project_versions.VersionID → PRIMARY KEY

Foreign Keys

Foreign keys establish relationships between tables and maintain referential integrity:

Foreign Key	References	Description
users.RoleID	roles.RoleID	Links user to their role
projects.UserID	users.UserID	Links project to its owner
repositories.ProjectID	projects.ProjectID	Links Git repository to project
folders.RepoID	repositories.RepoID	Links folder to repository
folders.ProjectID	projects.ProjectID	Links folder to project
folders.ParentFolderID	folders.FolderID	Supports nested folder structure
code_artifacts.FolderID	folders.FolderID	Links source file to its folder
project_versions.ProjectID	projects.ProjectID	Links version to its project

Unique Constraints

- users.Username → UNIQUE
- users.Email → UNIQUE
- roles.RoleName → UNIQUE
- projects.UserID + ProjectName → UNIQUE (one user cannot have two projects with the same name)

Check Constraints & Enumeration

- users.RoleID → Must reference a valid role
- code_artifacts.Language → Must be one of the supported languages (Python, Java, JavaScript, etc.)
- projects.SourceType → CHECK (SourceType IN ('ZIP', 'GitHub', 'GitLab', 'Local'))
- project_versions.VersionLabel → Cannot be empty

Additional Constraints

- NOT NULL constraints applied on all critical fields (e.g., ProjectName, FileName, PasswordHash, FolderPath).

- Default Values:
 - CreationDate, SignupDate, CreatedAt → CURRENT_TIMESTAMP
 - LastModifiedDate → Auto-update on record modification

Indexes:

- Index on projects.UserID
- Index on code_artifacts.FolderID
- Index on project_versions.ProjectID (for fast version retrieval)
- Composite index on folders.ProjectID + FolderPath

5.3 External API Design:

This section documents the external APIs exposed by the system. The architecture follows a microservices approach with a central API Gateway (Nginx) that routes requests to the appropriate backend services.

5.3.1 API Gateway Overview:

The system uses Nginx as an API Gateway to:

- Route requests between services
- Handle CORS and security headers
- Provide unified entry point
- Support load balancing and future scaling

Base URL: http://localhost:80/api/ (or domain in production)

Table 17: API Gateway

Endpoint	Method	Service	Description	Request Parameters	Response	Status Codes
/api/upm/projects/	POST	UPM	Create new project	project_name, project_description	Project object	201, 400, 401
/api/upm/projects/	GET	UPM	Get user projects	user_id (JWT)	List of projects	200, 401
/api/upm/projects/{id}/	GET	UPM	Get project details	-	Project details	200, 404
/api/upm/projects/{id}/	DELETE	UPM	Delete project	-	Success message	204, 403
/api/analysis/upload-folder/	POST	AI	Upload code (ZIP or files)	files, project_name, user_email	Analysis result	202, 400
/api/analysis/{project_id}/	GET	AI	Get analysis results	-	Analysis data + diagrams	200, 404
/api/analysis/{project_id}/high-level/	POST	AI	Generate high-level documentation	-	Documentation (Markdown/PDF)	200, 202

/api/analysis/{project_id}/ low-level/	POST	AI	Generate low-level documentation	-	Detailed documentation	200, 202
/api/analysis/diff/	POST	AI	Compare versions	version1_id, version2_id	Diff result	200, 400
/api/notifications/project/	POST	Notification	Project created/deleted	user_email, action, project_name	Notification status	201
/api/notifications/ documentation/	POST	Notification	Documentation exported	user_email, file_name, file_type	Notification status	201
/api/notifications/create/	POST	Notification	General notification	user_email, title, message	Notification status	201
/api/token/	POST	UPM	Login (JWT)	username, password	Access & Refresh tokens	200, 401
/api/token/refresh/	POST	UPM	Refresh token	refresh	New access token	200, 401

Detailed API Specifications:

1. UPM_Project Service (User & Project Management)

Base Path: /api/upm/

- **Authentication:** JWT Bearer Token
- **Main Endpoints:**
 - POST /projects/ → Create Project
 - GET /projects/ → List User Projects
 - GET /projects/{project_id}/ → Project Details
 - DELETE /projects/{project_id}/ → Delete Project

2. AI_Project Service (Code Analysis & Documentation)

Base Path: /api/analysis/

- **Key Features:**
 - Multi-language code upload and parsing
 - AST-based analysis
 - AI-powered documentation generation
 - Version comparison (Diff)
- **Important Endpoints:**
 - POST /upload-folder/ — Upload ZIP or multiple files
 - POST /{project_id}/high-level/ — Generate High-Level Summary
 - POST /{project_id}/low-level/ — Generate Detailed Documentation
 - POST /diff/ — Code Diff between versions

3. Notification Service

Base Path: /api/notifications/

- **Stateless service** (no heavy storage)
- Supports Email notifications via SMTP
- Uses Celery for background processing

Main Endpoints:

- POST /project/ — Project events
- POST /documentation/ — Export notifications
- POST /create/ — General system notifications
- GET /notifications/stats/ — Statistics

Authentication & Security:

- **Mechanism:** JWT (JSON Web Tokens)
- **Token Lifetime:** Access Token = 30 days, Refresh Token = 1 year
- **Authorization:** Role-Based Access Control (RBAC)
 - Admin
 - Developer
 - Reviewer
 -

5.4 Security and Access Control Design:

This section defines the security architecture of the system, covering authentication mechanisms, user roles, and access control policies. A robust security design protects sensitive source code, analysis results, and documentation while preventing unauthorized access and ensuring compliance with security best practices.

5.4.1 Authentication Mechanism:

Mechanism: JWT (JSON Web Tokens) with Refresh Token Rotation

- **Token Management:**
 - Access Token lifetime: **30 days**
 - Refresh Token lifetime: **365 days**
 - Refresh tokens are stored securely in HttpOnly cookies
 - Token revocation (blacklisting) is supported upon logout
- **Password Policy:**
 - Minimum length: **8 characters**
 - Must contain at least one uppercase letter, one lowercase letter, one number, and one special character
 - Passwords are hashed using **PBKDF2** with high iteration count (100,000+ iterations)
 - Support for Multi-Factor Authentication (MFA) is prepared for future implementation

5.4.2 User Roles and Permissions:

The system implements a **Role-Based Access Control (RBAC)** model. Each role has a clearly defined set of permissions that restrict access to specific features and data.

Role	Permissions	Access Level
Admin	Full system access including user management, project oversight, system configuration, and viewing all projects and logs	Full Access
Developer	Create and manage own projects, upload code (ZIP/GitHub), run analysis,	Standard

	generate documentation, view and compare versions, export reports	
Reviewer	View analysis results, UML diagrams, dependency graphs, generated documentation, and provide feedback/comments	Read + Review
Guest	View public landing page and general documentation only	Limited

Additional Security Measures

- **CORS Protection:** Strictly configured allowed origins
- **CSRF Protection:** Enabled with trusted origins
- **Rate Limiting:** Applied on sensitive endpoints (login, file upload, analysis requests)
- **API Gateway Security:** Nginx handles routing, security headers (HSTS, X-Frame-Options, XSS Protection), and hides internal microservices structure
- **Data Protection:**
 - Source code stored securely in isolated storage
 - Separation of databases (MySQL for users/projects, MongoDB for analysis results)
- **Logging & Monitoring:** All critical actions (login, upload, delete, analysis) are logged for auditing

Chapter 6 : Implementation Environment and Tools

6.1 Programming Languages and Frameworks:

A defining characteristic of microservices architecture is the freedom it affords each service to adopt the language and framework best suited to its domain. In the present system, however, a deliberate decision was made to adopt a unified backend technology stack—Python with Django 4.2—across all three microservices. This choice prioritises team velocity, knowledge sharing, and operational simplicity over polyglot flexibility. The frontend, by contrast, employs JavaScript with React.js, reflecting the distinct requirements of browser-based user interfaces.

6.1.1 Frontend: JavaScript and React.js:

The web application layer is implemented in JavaScript (ECMAScript 2015+) using React.js as the primary UI library. React was selected for its component-based architecture, which encourages the construction of reusable, testable UI elements that compose naturally into complex page layouts. The project leverages React Router v6 for declarative client-side routing, Redux Toolkit for predictable global state management, and Axios for HTTP communication with the backend through the Nginx API Gateway. Styling is handled by Tailwind CSS, a utility-first framework that promotes design consistency while eliminating the overhead of hand-written stylesheets. Vite serves as the build tool, offering sub-second hot-module replacement during development and optimised production bundles through tree-shaking and code splitting.

6.1.2 Backend Services: Python and Django 4.2:

All three backend microservices—the User & Project Management Service (UPM_Project), the Code Intelligence Service (Ai_Project), and the Notification Service (Notification_Service)—are implemented in Python 3.10/3.11 atop Django (versions 4.2 and 5.2). Django was chosen for its mature object-relational mapper (ORM), built-in admin interface, and comprehensive authentication framework, all of which accelerate development of data-intensive web services. Each service exposes its functionality through a Django REST Framework (DRF) API layer, providing automatic serialisation, pagination, and OpenAPI schema generation. The uniform framework choice simplifies cross-service code reviews, reduces the learning curve for new team members, and allows shared utility packages (logging, error handling, authentication middleware) to be reused verbatim across service boundaries.

6.1.3 API Gateway: Nginx:

Unlike the application-layer gateways common in cloud-native deployments, the present system employs Nginx as a reverse-proxy API Gateway. Nginx is configured to route incoming HTTP

requests to the appropriate backend service based on URL path prefixes: requests prefixed with `/api/upm/` are forwarded to the UPM service, `/api/ai/` to the Code Intelligence service, and `/api/notify/` to the Notification service. In addition to routing, Nginx handles SSL termination, request compression (gzip), static-file serving for the React frontend, and rate limiting to protect the backend from burst traffic. This infrastructure-level gateway approach was preferred over an application-layer alternative because it introduces no additional codebase to maintain and benefits from Nginx's battle-tested performance and low memory footprint.

Table 18: Programming Languages and Frameworks

Component	Language	Framework / Tool	Role
Web Application	JavaScript (ES6+)	React.js 18, Vite	Interactive frontend UI and state management
UPM_Project	Python 3.11	Django 4.2/5.2, DRF 3.14/3.16	User authentication, project CRUD, relational data
Ai_Project	Python 3.11	Django 4.2/5.2, DRF 3.14/3.16	Code analysis, AI inference, NoSQL data access
Notification_Service	Python 3.11	Django 4.2, DRF 3.14	Email, in-app, and push notification delivery
API Gateway	—	Nginx 1.24	Reverse proxy, SSL, rate limiting, static files

6.2 Development Environment and Project Structure:

Reproducibility of the development environment is a prerequisite for collaborative software engineering. This section describes the operating systems, integrated development environments, and containerisation tools employed throughout the project, as well as the structural organisation of the multi-repo.

6.2.1 Operating System and Hardware:

Development was carried out on workstations running Ubuntu 22.04 LTS as the primary operating system, supplemented by Windows 11 machines for frontend-focused tasks. Ubuntu was selected for its native compatibility with Docker, Python virtual environments, and server-side tooling. The minimum hardware specification comprised 16 GB of RAM, an Intel Core i7 or AMD Ryzen 7 processor, and 512 GB of solid-state storage—resources sufficient to run multiple Docker containers concurrently and to accommodate the memory demands of AI model inference during code analysis.

6.2.2 Integrated Development Environments:

Two IDEs were adopted across the team. Visual Studio Code (VS Code) served as the primary editor for frontend development and general-purpose scripting, augmented by extensions for ESLint, Prettier, Python, Docker, and Git. PyCharm Professional was used for backend Django development, where its integrated database browser, Django template support, and intelligent refactoring tools proved particularly valuable. Both IDEs shared a project-level `.editorconfig` file to enforce consistent indentation, charset, and end-of-line conventions, and a `.pre-commit` configuration ensured that linting checks were executed automatically before each commit.

6.2.3 Containerisation with Docker and Docker Compose:

Docker Engine 24.x and Docker Compose v2 form the backbone of both the development and deployment workflows. Each microservice is packaged in its own Docker image defined by a multi-stage Dockerfile that separates the dependency installation stage from the minimal runtime stage, thereby reducing image size and attack surface. A single `docker-compose.yml` file at the repository root orchestrates the full stack—frontend, Nginx gateway, three Django services, MySQL, MongoDB, and Redis—allowing any developer to bring the entire system up with a single `docker-compose up` command. Named volumes persist database data across container restarts, and an internal bridge network isolates inter-service communication from the host machine.

6.2.4 Project Structure (Monorepo):

The system comprises three independent Django projects, each maintained in its own repository on GitHub. This multi-repo approach was adopted because each service has distinct dependency requirements, deployment cycles, and team ownership, making separate repositories more practical than a single multi-repo. The three projects are: `UPM_Project` (User & Project Management Service), `Ai_project` (Code Intelligence Service), and `Notification_Service` (Notification Service). Although the services are housed in separate repositories, they share a common architectural template rooted in Django's project-app structure, and they communicate through well-defined REST API contracts exposed via the Nginx reverse-proxy gateway.

6.2.4.1 UPM_Project (User & Project Management Service):

`UPM_Project` is the largest of the three services and follows a layered architecture that separates concerns across six packages. The `models/` package defines the core domain entities—`UserProfile`, `Role`, `Project`, `CodeArtifact`, `Repository`, `Folder`, and `Version`—using Django's ORM backed by MySQL. The `repositories/` package encapsulates data-access logic (`artifact_repo`, `project_repo`, `user_repo`), decoupling query construction from the business layer. The `business_logic/` package orchestrates domain operations such as artifact management, project lifecycle, and user registration. The `serializers/` package contains DRF serializers for each model, handling validation and JSON rendering. The `views/` package exposes REST endpoints grouped by feature: account management, project CRUD, artifact upload, folder operations, GitHub integration,

reviewer actions, and administrative functions. A dedicated `celery_tasks/` package offloads long-running operations—notably file uploads to MongoDB/GridFS—to background workers via Redis. The service also integrates with GitHub through a `services/github_service.py` module that handles repository connection, synchronisation, and webhook processing.

6.2.4.2 `Ai_project` (Code Intelligence Service):

`Ai_project` is the most architecturally complex service, combining a standard Django REST API layer with a dedicated AI engine subsystem. The `core_ai/` Django app is structured similarly to UPM: `models/` (AI task, analysis result, code file, conflict report, evaluation, project analysis result), `repository/` (data access for AI tasks and code analysis), `serializers/`, `services/`, and `views/`. Beyond this conventional layer, the service introduces two specialised subsystems. The `ai_engine/` package contains the core intelligence components: `agents.py` implements multi-agent orchestration for different analysis tasks, `llm_client.py` manages communication with external large-language-model APIs, and `orchestrator.py` coordinates the overall analysis pipeline. The `language_processors/` package provides language-specific code parsers: a `base_processor` abstract class, a `generic_processor` for unknown languages, `java_processor` and `python_processor` for language-specific analysis, `tree_sitter_processor` for syntax-tree-based parsing, and a `processor_factory` that selects the appropriate parser at runtime. The `ai_engine/doc/` sub-package generates analysis documentation in multiple formats (HTML, Markdown, PDF, XML). Celery tasks in `celery_tasks/` handle asynchronous execution of analysis jobs, project-wide analysis, conflict detection, and export operations. Data is persisted primarily in MongoDB via MongoEngine, with collections for code files, analysis jobs, analysis results, AI explanations, documentation files, and project-level analysis results.

6.2.4.3 `Notification_Service` (Notification Service):

`Notification_Service` adopts the simplest structure of the three services, reflecting its narrower functional scope. The `notifications/` Django app consolidates its logic into a flat module layout: `models.py` defines the Notification model with fields for type, recipient, message, read status, and task tracking; `serializers.py` provides DRF serialisation; `views.py` exposes endpoints for creating project, code, documentation, user, and system notifications, as well as list, detail, mark-read, resend, and delete operations; `urls.py` maps these views to URL patterns under the `/api/notifications/` prefix; and `notification_service.py` encapsulates the core notification-dispatch logic. A `tasks.py` module leverages Celery to process notification delivery asynchronously, preventing blocking on external service calls (e.g., email or push-notification providers). The service is configured to use MySQL as its primary database (via `mysqlclient`) and Redis as the Celery broker. Flower is included as a dependency for real-time monitoring of Celery task execution. Despite their differing levels of complexity, all three services share a common structural DNA: a Django project directory containing `settings.py`, `urls.py`, `celery_app.py`, and WSGI/ASGI entry points; a primary Django app (`core_upm`, `core_ai`, or `notifications`) housing models, serializers, and views; a Celery configuration connected to a shared Redis instance; and a

Dockerfile for containerised deployment. This consistency reduces the cognitive overhead of switching between services and allows shared tooling—such as the drf-yasg Swagger documentation generator and the simplejwt authentication library—to be configured identically across the codebase.

6.3 Database Servers:

The microservices architecture mandates the database-per-service pattern: each service owns a private data store, and no service accesses another service's database directly. This pattern decouples data schemas, permits independent schema evolution, and eliminates single points of failure at the persistence layer. The system employs three distinct storage technologies—MySQL, MongoDB, and Redis—each chosen to match the data-access characteristics of the service it serves.

6.3.1 MySQL 8.0 — User & Project Management Service:

The UPM_Project service stores its structured relational data in MySQL 8.0. MySQL was selected for its high-performance InnoDB storage engine, robust transactional support (ACID compliance), and widespread ecosystem of tooling. Django's ORM generates optimised SQL queries automatically, while Django migrations provide a versioned, reversible mechanism for schema changes. The MySQL instance is configured with UTF-8 character encoding (utf8mb4), SSL-encrypted connections, and connection pooling via Django's `CONN_MAX_AGE` setting. Automated daily backups are performed using `mysqldump` and stored in a secure offsite location.

6.3.2 MongoDB 7.0 — Code Intelligence Service:

The Ai_Project service persists its semi-structured analysis results, code metadata, and AI model outputs in MongoDB 7.0. MongoDB's document-oriented storage model accommodates the heterogeneous and evolving schemas of code analysis reports without requiring costly migrations. Each analysis result is stored as a BSON document containing embedded arrays of issues, metrics, and AI-generated suggestions, enabling the entire report to be retrieved in a single read operation. The service accesses MongoDB through MongoEngine, an ODM (Object-Document Mapper) that integrates transparently with Django. The instance runs with authentication enabled and TLS encryption for data in transit.

6.3.3 Redis 7.x — Celery Message Broker and Cache:

Redis serves a dual purpose within the system. First, it acts as the message broker for Celery, the distributed task queue used by the Ai_Project service to offload long-running operations—such as full-repository code analysis and AI model inference—from the synchronous request-response cycle. Celery enqueues task messages into Redis, and worker processes consume them asynchronously, publishing results back to a dedicated result backend (also Redis). Second, Redis functions as a lightweight in-memory cache for frequently accessed data, such as the most recent

analysis summaries and notification counts, reducing repeated database queries and improving response latency

Service	Database	Version	Port	Access Layer	Key Features
UPM_Project	MySQL	8	3306	Django ORM	ACID, utf8mb4, SSL, connection pooling
Ai_Project	MongoDB	7	27017	MongoEngine ODM	Flexible schema, embedded docs, TLS
Ai_Project	Redis	7.x	6379	Celery / django-redis	Task broker, result backend, caching

6.4 Libraries and Dependencies:

Every microservice maintains its own dependency manifest—requirements.txt for the Django services and package.json for the React frontend—ensuring isolated, reproducible environments. This section enumerates the principal libraries upon which each component depends, grouping them by functional role.

6.4.1 Frontend Dependencies (package.json):

The React frontend relies on the following core libraries. React 18 and ReactDOM provide the component rendering engine; Redux Toolkit and react-redux manage global application state with minimal boilerplate; Axios handles HTTP requests to the Nginx gateway; and React Router v6 enables declarative single-page navigation. Styling is achieved through Tailwind CSS and its companion Headless UI library for accessible, unstyled interactive components. Data visualisation of analysis results is rendered by Recharts, while Formik and Yup manage form state and validation schemas respectively. The development toolchain comprises TypeScript for static type checking, Vite for fast builds and hot-module replacement, ESLint for linting, and Prettier for automatic code formatting

6.4.2 Shared Backend Dependencies (requirements.txt):

Because all three backend services share the Django framework, a common set of dependencies appears in each service's requirements.txt. Django (4.2 for Ai_project and Notification_Service; 5.2 for UPM_Project) provides the core web framework; Django REST Framework 3.14 supplies serialisation, view sets, and pagination for the REST API layer; djangorestframework-simplejwt implements JWT-based stateless authentication; django-cors-headers configures Cross-Origin Resource Sharing for browser-based requests; and gunicorn serves as the production WSGI HTTP server. Environment-specific settings are managed through python-decouple, which reads variables from a .env file, keeping secrets out of source control.

6.4.3 Service-Specific Dependencies:

Beyond the shared core, each service declares additional libraries that address its particular domain requirements. The UPM_Project service includes mysqlclient as the MySQL database adapter and Pillow for user-avatar image processing. The Ai_Project service lists pymongo and mongoengine for MongoDB connectivity, scikit-learn and transformers (Hugging Face) for machine-learning inference, tree-sitter for syntactic code parsing, celery and redis for asynchronous task execution, and aiohttp for outbound HTTP calls to external AI APIs. The Notification_Service includes firebase-admin for push notifications, django-celery-email for asynchronous email delivery via Celery, and redis-py for caching notification state.

Table 19: Libraries and Dependencies

Component	Library	Version	Role
Frontend	react, react-dom	18.x	Component rendering engine
Frontend	@reduxjs/toolkit, react-redux	2.x	Global state management
Frontend	axios	1.6	HTTP client for API Gateway communication
Frontend	tailwindcss, headlessui	3.x / 1.x	Utility-first CSS and accessible UI primitives
Frontend	recharts	2.x	Charting library for analysis visualisation
Backend (shared)	django	4.2	Core web framework
Backend (shared)	djangorestframework	3.14	REST API serialisation and view sets
Backend (shared)	djangorestframework-simplejwt	5.3	JWT authentication tokens
Backend (shared)	gunicorn	21.x	Production WSGI HTTP server
Backend (shared)	python-decouple	3.8	Environment variable management
UPM_Project	mysqlclient	2.2	MySQL database adapter
UPM_Project	Pillow	10.x	Image processing for avatars
Ai_Project	mongoengine	0.27	MongoDB ODM for Django
Ai_Project	scikit-learn, transformers	1.3 / 4.x	ML inference and NLP models
Ai_Project	tree-sitter	0.2	Source-code syntax parsing
Ai_Project	celery, redis	5.3 / 5.x	Asynchronous task queue and broker
Notification_Service	firebase-admin	6.x	Push notification delivery
Notification_Service	django-celery-email	3	Async email via Celery

6.5 Continuous Integration and Deployment (CI/CD):

Manual build and deployment processes are error-prone and do not scale to a multi-service architecture. The project therefore automates the entire pipeline—from code commit to production deployment—using GitHub Actions and Docker, ensuring that every change is validated, built, and delivered in a consistent, repeatable manner.

6.5.1 Continuous Integration Pipeline:

The CI pipeline is triggered on every push to dev and on every pull request targeting dev or main. It comprises four sequential stages. First, Code Linting runs ESLint on the frontend and flake8 on each Django service to enforce coding standards. Second, Unit Tests executes the Jest suite for the frontend and the pytest suite for each backend service, producing coverage reports that are uploaded as build artifacts. Third, Build Verification confirms that the frontend production bundle can be built and that each Docker image assembles without error. Fourth, Security Scanning invokes npm audit and pip-audit to detect known vulnerabilities in third-party dependencies. A pull request may not be merged until every stage reports success.

6.5.2 Continuous Deployment Pipeline:

The CD pipeline activates when code is merged into main. The process proceeds as follows: (1) Build—GitHub Actions builds production Docker images for each modified service using multi-stage Dockerfiles that minimise image size; (2) Publish—images are tagged with the Git commit SHA and pushed to GitHub Container Registry (ghcr.io); (3) Staging Deploy—the updated containers are deployed to a staging server that mirrors the production topology, and a suite of integration tests is executed; (4) Production Deploy—if staging tests pass, the production server pulls the new images and restarts the affected services via Docker Compose; (5) Health Check—the pipeline verifies that each updated service responds to its /health endpoint within a configurable timeout. Should a health check fail, an automated rollback restores the previous container versions, and the team is notified via a Slack webhook.

6.5.3 Docker Containerisation Strategy:

Every microservice, the Nginx gateway, and each database run inside dedicated Docker containers, connected through an internal bridge network. Multi-stage Dockerfiles separate the build stage (dependency installation, static asset compilation) from the runtime stage (a slim Python or Nginx image containing only what is necessary to execute), reducing final image sizes by approximately 60 percent compared to single-stage builds. Docker Compose manages the full stack lifecycle, and named volumes ensure that MySQL, MongoDB, and Redis data persist across container restarts. Environment variables injected through .env files allow a single Docker Compose configuration to be reused across development, staging, and production environments with only connection-string and secret differences.

6.6 Testing Environment:

A robust testing infrastructure is indispensable for validating functional correctness, performance, and reliability before code reaches end users. The project employs a three-tier testing strategy—local, staging, and performance—each providing progressively broader coverage and higher fidelity to the production environment.

6.6.1 Local Development Testing:

During routine development, each engineer runs the complete system stack locally via Docker Compose. Unit tests are executed before every commit through a pre-commit hook: pytest for the Django services and Jest for the React frontend. Hot-reloading is enabled for both the frontend (Vite dev server) and the backend (Django's runserver within the container), providing immediate feedback on code changes. Local testing covers individual API endpoints, serializer validation, model methods, and React component rendering.

6.6.2 Staging and Integration Testing:

The staging environment, provisioned automatically by the CD pipeline, replicates the production topology on a dedicated server. It runs the same Docker Compose configuration with separate database instances populated by anonymised seed data. Integration tests verify cross-service interactions—for example, that creating a project through the UPM service triggers a code analysis job in the AI service, which in turn publishes a notification through the Notification service. End-to-end tests, written with Playwright, simulate complete user workflows including login, project creation, analysis initiation, and result visualisation, exercising the full stack from browser to database.

6.6.3 Performance and Load Testing:

Performance testing is conducted using Locust, a Python-based load-generation framework. Test scenarios simulate concurrent users performing typical workflows—project creation, code analysis requests, and notification retrieval—under incrementally increasing loads. Key metrics collected include median and 95th-percentile response times (target: < 2 s for API requests), throughput in requests per second, error rate (target: < 1 %), and per-container resource utilisation (CPU, memory, database connection count). Results are archived alongside each release and compared against baseline measurements to detect regressions before they reach production.

Chapter 7: Program Implementation

7.1 File Structure and Project:

The system is organized as a multi-service architecture consisting of three independent Django microservices, a React.js frontend application, and an Nginx API Gateway that orchestrates communication between the client and the backend services. Each microservice follows a consistent internal structure that promotes separation of concerns, maintainability, and testability. The directory layout of each service reflects the layered architecture pattern, with dedicated packages for models, views, serializers, business logic, repositories, and asynchronous task definitions. This section presents an overview of the project directory tree for each service and identifies the common structural DNA that unifies the codebase.

Table 20: File Structure and Project

Service	Framework	Database	Primary Responsibility
UPM_Project	Django 5.2 + DRF	MySQL 8.0	User & Project Management, Authentication, Artifact Storage
Ai_project	Django 4.2 + DRF	MongoDB 7.0	Code Analysis, AI Documentation, Conflict Detection
Notification_Service	Django 5.2 + DRF	SQLite	Notification Creation, Email Delivery, Status Tracking

7.1.1 UPM_Project Directory Tree:

The UPM_Project service serves as the central management hub for users, projects, and code artifacts. Its directory structure follows a feature-based organization within the core_upm application, where each functional concern is encapsulated in its own submodule. The models package contains entity definitions for User, Project, CodeArtifact, Folder, Repository, and Version. The views package houses the API view controllers, the serializers package manages data transformation between Django models and JSON representations, and the business_logic and repositories packages implement the service and repository patterns respectively. Additionally, dedicated modules handle Celery tasks, permissions, JWT token customization, middleware, and inter-service communication utilities.

```
UPM_Project/
├── UPM_Project/
│   ├── settings.py      # Django project configuration
│   ├── urls.py          # Database, middleware, Celery config
│   ├── celery_app.py    # Root URL routing
│   ├── wsgi.py / asgi.py # Celery application instance
│   └── deployment.py     # Deployment entry points
├── core_upm/
│   ├── models/          # Core application module
│   │   ├── user.py      # Domain entities
│   │   ├── project.py   # Role, UserProfile
│   │   ├── artifact.py  # Project
│   │   ├── folder.py    # CodeArtifact
│   │   ├── repository.py # Folder
│   │   └── version.py   # Repository
│   ├── views/           # Version
│   ├── serializers/     # API view controllers
│   └── business_logic/  # DRF serializers
│                       # Service layer (UserService, etc.)
```

— repositories/	# Repository layer (data access)
— services/	# External service integrations
— celery_tasks/	# Asynchronous task definitions
— permissions.py	# Custom permission classes
— tokens.py	# RoleAwareRefreshToken
— middleware.py	# Request/response middleware
— notification_utils.py	# NotificationClient helper
— external_storage.py	# Cross-service storage proxy

7.1.2 Ai_project Directory Tree:

The `Ai_project` service is the most structurally complex of the three microservices, reflecting its responsibility for code intelligence, AI-powered documentation, and conflict detection. The `core_ai` application is organized into several major subsystems: the `models` package uses Pydantic-based data models mapped to MongoDB collections rather than the traditional Django ORM; the `ai_engine` package contains the documentation orchestration system, the AI agents, and the LLM client; the `language_processors` package implements the Strategy pattern for multi-language code analysis using AST-based processors; and the `services` package provides higher-level analysis workflows including project-level analysis, dependency graph construction, context building, and conflict detection. The conflict detection subsystem itself is a mini-module with its own orchestrator, AST unification, structural diffing, semantic analysis, and report generation components.

Ai_project/	
— Ai_project/	# Django project configuration
— settings.py	# MongoDB, Celery, LLM config
— urls.py	# Root URL routing
— celery_app.py	# Celery application instance
— core_ai/	# Core AI application module
— models/	# Pydantic + MongoDB models
— codefile.py	# CodeFile (GridFS reference)
— analysis.py	# AnalysisResult, AnalysisJob
— ai_explanation.py	# AIExplanation, GeneratedFile
— conflict_report.py	# ConflictReport
— evaluation.py	# EvaluationResult, HumanFeedback
— ai_engine/	# AI documentation subsystem
— orchestrator.py	# DocumentationOrchestrator
— agents.py	# High/Low Level Agents
— llm_client.py	# OpenRouter API client
— doc/	# Documentation generation
— language_processors/	# Strategy-pattern processors
— python_processor.py	# Python AST analysis
— java_processor.py	# Java analysis
— base_processor.py	# ILanguageProcessorStrategy
— processor_factory.py	
— services/	# Business workflows
— project_analysis_service.py	
— conflict_detection/	# Conflict detection subsystem
— repository/	# MongoDB data access
— celery_tasks/	# Async analysis tasks
— views/	# API view controllers

7.1.3 Notification_Service Directory Tree:

The `Notification_Service` adopts a simpler, flat structure compared to the other two services. This simplicity is by design: the service has a single, well-defined responsibility — receiving notification requests, persisting them to the database, and delivering them asynchronously via email. The `notifications` Django app contains the `Notification` model, its corresponding serializer, a

SimpleNotification helper class that encapsulates email-sending logic, Celery task definitions for asynchronous delivery, and a set of REST API views. The flat structure reflects the single-aggregate nature of the service and minimizes unnecessary abstraction layers, which is appropriate for a service with limited domain complexity.

```
Notification_Service/
├── notification_service/      # Django project configuration
│   ├── settings.py          # SQLite, Celery, email config
│   ├── urls.py              # Root URL routing
│   └── celery_app.py         # Celery application instance
├── notifications/           # Notification application
│   ├── models.py            # Notification model
│   ├── views.py             # CRUD + create endpoints
│   ├── serializers.py       # DRF serializers
│   ├── notification_service.py # SimpleNotification (email)
│   ├── tasks.py             # Celery async tasks
│   ├── urls.py              # App URL routing
│   └── admin.py             # Django admin registration
```

7.1.4 Common Structural Patterns:

Despite their differing complexities, all three microservices share a common structural DNA that reflects the overarching architectural decisions of the system. Each service employs Django as the web framework with Django REST Framework for API exposure, Celery for asynchronous task processing, and a consistent separation between the API layer (views and serializers), the business logic layer (services or business_logic packages), and the data access layer (repositories or direct ORM access). Furthermore, each service defines its own Docker container configuration, and inter-service communication is achieved exclusively through synchronous HTTP REST calls with short timeouts for non-critical operations, ensuring loose coupling and independent deployability. This uniformity reduces the cognitive load for developers working across multiple services and facilitates consistent monitoring, logging, and deployment practices.

7.2 Module Implementation:

This section provides a detailed description of the implementation of each major module within the three backend microservices. The discussion is organized by service, and within each service, the modules are described following the layered architecture from data models through business logic to API endpoints. This structure enables the reader to trace the implementation path from persistent storage to the external interface exposed to the client.

7.2.1 User and Project Management Service (UPM_Project):

Models Layer. The UPM_Project service defines its domain entities using the Django ORM mapped to a MySQL 8.0 database. The core models include *Role* and *UserProfile* for role-based access control, *Project* for project management, *CodeArtifact* for versioned code storage, *Folder* for hierarchical file organization, *Repository* for GitHub-linked repositories, and *Version* for version tracking. The Role model establishes the three-tier role hierarchy (Developer, Admin, Reviewer) with a JSONField for flexible permission definitions, while the UserProfile model extends Django's built-in User with a one-to-one relationship that adds role assignment, full name, online status, and activity tracking fields.

```
class Role(models.Model):
    role_id = models.AutoField(primary_key=True)
    role_name = models.CharField(max_length=50, unique=True)
    role_type = models.CharField(max_length=20, choices=[
```

```

        ('DEVELOPER', 'Developer'),
        ('ADMIN', 'Admin'),
        ('REVIEWER', 'Reviewer')
    ], unique=True, default='DEVELOPER')
    description = models.TextField(blank=True)
    permissions_list = models.JSONField(default=dict)

class UserProfile(models.Model):
    user = models.OneToOneField(User, on_delete=models.CASCADE,
                               primary_key=True, related_name='profile')
    role = models.ForeignKey(Role, on_delete=models.SET_NULL,
                             null=True, blank=True)
    full_name = models.CharField(max_length=255)
    signup_date = models.DateTimeField(auto_now_add=True)
    last_seen = models.DateTimeField(null=True, blank=True)
    is_online = models.BooleanField(default=False)

```

Repository Pattern. The service employs the Repository pattern to abstract data access operations away from the business logic layer. Each domain entity has a corresponding repository class — *UserRepository*, *ProjectRepository*, and *ArtifactRepository* — that encapsulates all database query operations. This design decision ensures that the business logic layer remains agnostic of the underlying persistence mechanism and facilitates unit testing through mock injection. Notably, the *ArtifactRepository* also handles cross-service data access by making HTTP requests to the *Ai_project* service for external file content storage and retrieval, thereby acting as an anti-corruption layer between the two services.

Business Logic Layer. The service layer is implemented through dedicated service classes: *UserService*, *ProjectService*, and *ArtifactService*. Each service class receives its corresponding repository through constructor injection, enabling dependency inversion and testability. The *UserService* handles user registration with duplicate-checking logic, authentication with email-based fallback, and profile management. The *ProjectService* manages the full lifecycle of projects including creation, member assignment, and status transitions. The *ArtifactService* orchestrates code upload workflows that span both the *UPM_Project* and *Ai_project* services, coordinating metadata storage in MySQL and content storage in MongoDB through the *ArtifactRepository*'s cross-service methods.

7.2.2 Code Intelligence Service (*Ai_project*):

Models Layer — Pydantic with MongoDB. Unlike the *UPM_Project* service, the *Ai_project* service uses Pydantic models mapped to MongoDB collections rather than the traditional Django ORM. This decision was driven by the semi-structured nature of code analysis results, which exhibit high schema variability across different programming languages and analysis types. The primary models include *CodeFile* (with a GridFS reference for large source files), *AnalysisResult* (containing AST structures, extracted features, and dependencies), *AIExplanation* and *GeneratedFile* (for AI-generated documentation), *ConflictReport* (for detected code-documentation inconsistencies), and *EvaluationResult* with *HumanFeedback* (for quality assessment). Each Pydantic model includes MongoDB-specific configuration such as collection names and ObjectId handling, providing type-safe serialization and validation while maintaining the flexibility of a document database.

AI Engine Subsystem. The AI engine is the centerpiece of the *Ai_project* service and is organized into three cooperating components. The *DocumentationOrchestrator* serves as the entry point, receiving requests for documentation generation and determining whether a cached explanation exists in MongoDB before invoking the appropriate AI agent. The *Agents* module defines a hierarchy of AI agents: *HighLevelAgent* and *LowLevelAgent* for single-file analysis, and

`ProjectHighLevelAgent` and `ProjectLowLevelAgent` for project-wide analysis. Each agent constructs a specialized system prompt and delegates the actual language model interaction to the `LLMClient`. The `LLMClient` encapsulates the OpenRouter API integration, managing model selection, rate limiting, automatic model fallback on 429 errors, and response caching with a 48-hour TTL.

Language Processors — Strategy Pattern. The language processing subsystem implements the Strategy design pattern through the `ILanguageProcessorStrategy` interface, which defines three core operations: `parse_source_code` (AST generation), `extract_features` (metric extraction), and `generate_uml_diagram` (UML data generation). Concrete implementations include `PythonProcessor` (using Python's built-in `ast` module), `JavaProcessor`, and `GenericProcessor` for fallback analysis. The `ProcessorFactory` selects the appropriate processor based on the file type, and a `TreeSitterProcessor` is also available for languages not natively supported by dedicated processors. This pattern allows new language support to be added by implementing a single interface without modifying existing code.

7.2.3 Notification Service (Notification_Service):

The `Notification_Service` implements a streamlined notification pipeline with minimal architectural complexity. The Notification model captures all notification metadata including the recipient email, title, message body, notification type (one of six predefined categories such as `PROJECT_CREATED`, `CODE_ADDED`, and `REVIEWER_ALERT`), delivery status (`PENDING`, `SENT`, or `FAILED`), priority level, read status, and an optional Celery task ID for traceability. The `SimpleNotification` helper class encapsulates the email delivery logic using Django's built-in SMTP backend, providing a clean interface between the notification model and the delivery mechanism. Celery tasks defined in the `tasks` module handle the asynchronous dispatch of notifications, updating the status field upon successful or failed delivery. The flat module structure reflects the single-aggregate nature of this service and avoids unnecessary abstraction layers that would add complexity without proportional benefit.

7.3 API Implementation:

All three microservices expose their functionality through RESTful APIs built with Django REST Framework. The API layer follows a consistent pattern: view classes or function-based views handle HTTP request parsing and response formatting, serializers manage data validation and transformation, and business logic is delegated to the service layer. This section presents representative API implementations from each service, with particular emphasis on the authentication mechanism, project management endpoints, and code analysis triggers.

7.3.1 Authentication API — JWT with Role-Aware Tokens:

The authentication system is implemented as a JWT-based mechanism with role-aware token generation. The standard `SimpleJWT` library is extended through the `RoleAwareRefreshToken` class, which injects the user's role type, full name, and user ID into the token payload during generation. This approach eliminates the need for additional database queries on subsequent authenticated requests, as the role information is embedded directly in the access token and can be verified without server-side state. The following code snippet demonstrates the login endpoint implementation, which authenticates the user, updates the online status and last seen timestamp, generates a role-aware token pair, and returns the tokens along with user profile data.

```
class UserLoginAPIView(APIView):
    permission_classes = [AllowAny]
    user_service = UserService()
```

```

def post(self, request):
    username = request.data.get('username')
    password = request.data.get('password')
    user = self.user_service.authenticate_user(username, password)
    # Update online status
    profile = user.profile
    profile.last_seen = timezone.now()
    profile.is_online = True
    profile.save(update_fields=['last_seen', 'is_online'])
    # Generate role-aware JWT
    refresh = RoleAwareRefreshToken.for_user(user)
    return Response({
        'refresh': str(refresh),
        'access': str(refresh.access_token),
        'role_type': refresh.payload.get('role_type', 'DEVELOPER'),
        'full_name': refresh.payload.get('full_name', ''),
        'user': UserSerializer(user).data
    }, status=status.HTTP_200_OK)

```

The `RoleAwareRefreshToken` class extends the standard `RefreshToken` by overriding the `for_user` class method. Upon token generation, it attempts to retrieve the user's profile and injects the `role_type`, `full_name`, and `user_id` claims into the token payload. If the profile does not exist (which may occur during initial setup or for system accounts), the method gracefully falls back to assigning the default `DEVELOPER` role. This design ensures backward compatibility and system resilience while providing the role information needed for authorization decisions at the API gateway and within individual view handlers.

```

class RoleAwareRefreshToken(RefreshToken):
    @classmethod
    def for_user(cls, user):
        token = super().for_user(user)
        try:
            profile = user.profile
            token["role_type"] = profile.role_type
            token["full_name"] = profile.full_name
            token["user_id"] = user.id
        except Exception:
            token["role_type"] = "DEVELOPER"
        return token

```

7.3.2 Project Management API:

The project management API provides a full set of CRUD endpoints for the Project entity, along with specialized endpoints for project member management, status transitions, and reviewer assignment. The `project_views` module defines view classes that map HTTP methods to service-layer operations: POST requests trigger project creation through `ProjectService`, GET requests retrieve project details with optimized database queries using `select_related` for role and user data, PUT and PATCH requests handle updates and partial updates respectively, and DELETE requests perform soft deletion by updating the project status rather than removing the record from the database. Permission enforcement is achieved through custom permission classes such as `IsAdminRole` and `IsReviewerRole`, which inspect the `role_type` claim embedded in the JWT token to determine access rights.

```

class IsAdminRole(BasePermission):
    message = "Admin access required."
    def has_permission(self, request, view):
        user = request.user

```



```

if not user or not user.is_authenticated:
    return False
try:
    return getattr(user.profile, "role_type", None) == "ADMIN"
except Exception:
    return False

```

7.3.3 Code Intelligence API:

The code intelligence API exposed by the `Ai_project` service provides endpoints for code file upload, analysis triggering, result retrieval, documentation generation, conflict detection, and evaluation feedback. The `codefile` view handles file uploads by storing the source content in GridFS and creating a `CodeFile` metadata document in MongoDB. The `analysis` view accepts POST requests to initiate code analysis, which spawns a Celery task (`analyze_code_file_task`) and returns the task ID for polling. The `explanation_views` module provides endpoints for requesting AI-generated documentation at both high and low levels of abstraction. The `conflict_views` module exposes the conflict detection pipeline, allowing users to submit pairs of code or documentation for structural and semantic comparison. All endpoints follow the same pattern of validating input through serializers, delegating to service classes for business logic, and returning standardized response objects.

7.3.4 Notification API:

The notification API provides endpoints for creating notifications, querying notifications by user email, marking notifications as read, and retrieving notification statistics. The create endpoint accepts notification data including the recipient email, title, message, type, and priority, persists the notification record with a PENDING status, and optionally triggers an asynchronous Celery task for email delivery. The query endpoint supports filtering by user email and read status, with results ordered by creation date in descending order. The mark-as-read endpoint performs a partial update on the `is_read` field. All endpoints use DRF's generic view classes with appropriate serializer classes, ensuring consistent input validation and output formatting across the notification API surface.

Table 21: API Implementation

Service	Method	Endpoint	Description
UPM	POST	<code>/api/users/login/</code>	User authentication with JWT
UPM	POST	<code>/api/users/register/</code>	User registration
UPM	GET	<code>/api/projects/</code>	List user projects
UPM	POST	<code>/api/projects/</code>	Create new project
UPM	POST	<code>/api/artifacts/</code>	Upload code artifact
AI	POST	<code>/api/codefiles/</code>	Upload source file
AI	POST	<code>/api/analysis/</code>	Trigger code analysis
AI	GET	<code>/api/analysis/{id}/</code>	Retrieve analysis result
AI	POST	<code>/api/explanations/</code>	Request AI documentation
AI	POST	<code>/api/conflicts/</code>	Trigger conflict detection
Notif.	POST	<code>/api/notifications/project/</code>	Create notification
Notif.	GET	<code>/api/notifications/</code>	Query user notifications

7.4 Data Access Layer Implementation:

The data access layer forms the foundation of the system's persistence mechanism, bridging the gap between the business logic and the underlying database engines. Given the polyglot persistence strategy adopted by the system, the data access layer varies across services: the UPM_Project service uses the Django ORM with MySQL, the Ai_project service uses Pydantic models with MongoDB and GridFS, and the Notification_Service uses the Django ORM with SQLite. This section details each approach and also examines the cross-service data access patterns that enable communication between the microservices.

7.4.1 Django ORM with MySQL (UPM_Project):

The UPM_Project service leverages the Django ORM as its primary data access mechanism, mapping Python model classes to MySQL 8.0 tables with full support for foreign key relationships, one-to-one mappings, and JSON fields. The ORM provides an abstraction layer that eliminates the need for raw SQL in the majority of operations while maintaining the ability to execute optimized queries when necessary. The model definitions use Django's declarative syntax, with AutoField primary keys, CharField with unique constraints, OneToOneField for profile associations, and ForeignKey for role assignments. Query optimization is achieved through `select_related` for foreign key joins and `prefetch_related` for many-to-many relationships, reducing the number of database round-trips in list and detail views.

```
# Role model - establishes the three-tier role hierarchy
class Role(models.Model):
    role_id = models.AutoField(primary_key=True)
    role_name = models.CharField(max_length=50, unique=True)
    role_type = models.CharField(max_length=20, choices=[
        ('DEVELOPER', 'Developer'),
        ('ADMIN', 'Admin'),
        ('REVIEWER', 'Reviewer')
    ], unique=True, default='DEVELOPER')
    permissions_list = models.JSONField(default=dict)

# UserProfile - extends Django User with role and activity data
class UserProfile(models.Model):
    user = models.OneToOneField(User, on_delete=models.CASCADE,
                               primary_key=True, related_name='profile')
    role = models.ForeignKey(Role, on_delete=models.SET_NULL,
                             null=True, blank=True)
    full_name = models.CharField(max_length=255)
    is_online = models.BooleanField(default=False)
```

7.4.2 Pydantic Models with MongoDB (Ai_project):

The Ai_project service adopts a fundamentally different data access strategy, using Pydantic models for data validation and serialization in conjunction with the pymongo driver for MongoDB operations. This approach was chosen because code analysis results exhibit high schema variability — the structure of an AST, the set of extracted features, and the format of dependency graphs differ significantly across programming languages and analysis types. The CodeFile model stores metadata (filename, file type, source project ID) alongside a GridFS reference (gridfs_id) for the actual source code content, which may exceed MongoDB's 16 MB document size limit. Analysis results are stored as nested documents with flexible schemas, leveraging MongoDB's schema-less nature while Pydantic ensures type safety at the application boundary.


```

# MongoDB query pattern - retrieving and deserializing a CodeFile
db = get_mongo_db()
code_file_data = db[settings.CODE_FILES_COLLECTION].find_one(
    {"_id": ObjectId(code_file_id_str)}
)
code_file = CodeFile(**code_file_data)

# GridFS read for large source files
code_content = read_from_gridfs(code_file.gridfs_id)

# MongoDB insert with Pydantic serialization
result_instance = AnalysisResult(
    code_file_id=ObjectId(code_file_id_str),
    status="COMPLETED",
    ast_structure=analysis_result.get('ast_structure'),
    extracted_features=analysis_result.get('features_extracted'),
)
inserted = db[settings.ANALYSIS_RESULTS_COLLECTION].insert_one(
    result_instance.dict(by_alias=True, exclude_unset=True)
)

```

7.4.3 Django ORM with SQLite (Notification_Service):

The Notification_Service uses Django ORM with SQLite as its persistence mechanism, a choice justified by the service's lightweight nature and single-aggregate domain model. The Notification model includes database-level indexes on frequently queried fields such as user_email, is_read, and task_id, ensuring efficient query performance even as the notification volume grows. The model defines a comprehensive set of choices for notification_type, status, and priority, providing both documentation and validation at the database level. The SQLite database is container-local and does not require a separate database server, simplifying the deployment topology and reducing operational overhead for this non-critical service.

```

class Notification(models.Model):
    NOTIFICATION_TYPES = [
        ('PROJECT_CREATED', 'Project Created'),
        ('CODE_ADDED', 'Code Added'),
        ('DOCUMENTATION_EXPORTED', 'Documentation Exported'),
        ('SYSTEM_ALERT', 'System Alert'),
        ('USER_REGISTERED', 'User Registered'),
        ('REVIEWER_ALERT', 'Reviewer Alert'),
    ]
    user_email = models.EmailField(db_index=True)
    title = models.CharField(max_length=255)
    message = models.TextField()
    notification_type = models.CharField(max_length=50,
                                         choices=NOTIFICATION_TYPES)
    status = models.CharField(max_length=20, default='PENDING')
    priority = models.CharField(max_length=20, default='MEDIUM')
    is_read = models.BooleanField(default=False, db_index=True)
    task_id = models.CharField(max_length=255, blank=True, null=True,
                               db_index=True)

```

7.4.4 Cross-Service Data Access:

One of the most significant implementation challenges in the microservices architecture is cross-service data access. Since each service owns its database exclusively — in accordance with

the database-per-service pattern — any data required from another service must be obtained through inter-service HTTP calls. The UPM_Project service implements this pattern through two dedicated client classes. The ArtifactRepository handles communication with the Ai_project service for external content storage, while the NotificationClient handles communication with the Notification_Service for notification dispatch. Both clients use short timeouts (3–10 seconds) and treat failures as non-critical, ensuring that the primary service remains operational even when a downstream service is temporarily unavailable.

```
class ArtifactRepository:
    def save_content_externally(self, content, project_id,
                               file_name, code_language):
        data = {
            "filename": file_name,
            "file_type": code_language,
            "content": content,
            "source_project_id": str(project_id)
        }
        response = requests.post(AI_CODEFILES_API, json=data,
                                headers=self._get_auth_headers(), timeout=10)
        response.raise_for_status()
        return response.json()['id']

    def fetch_content_externally(self, storage_reference):
        url = f"{AI_CODEFILES_API}{storage_reference}/"
        response = requests.get(url,
                                headers=self._get_auth_headers(), timeout=10)
        response.raise_for_status()
        return response.json()['content']
```

The NotificationClient follows a similar pattern, providing static methods for sending notification requests to the Notification_Service. The `_send_notification_request` helper method constructs the full URL from the service base URL and the endpoint path, sets a 3-second timeout, and returns a boolean indicating success or failure. The calling code in UPM_Project uses this client in a fire-and-forget manner — notification delivery failures are logged but do not cause the primary operation to fail, reflecting the non-critical nature of notification delivery in the overall system workflow.

```
class NotificationClient:
    NOTIFICATION_SERVICE_URL = 'http://notification_django_app:8000'

    @staticmethod
    def _send_notification_request(endpoint, data):
        url = (f"{NotificationClient.NOTIFICATION_SERVICE_URL}"
              f"/api/notifications/{endpoint}/")
        response = requests.post(url, json=data, timeout=3)
        return response.status_code in [200, 201]

    @staticmethod
    def send_project_notification(user_email, action,
                                project_name, **kwargs):
        data = {'user_email': user_email, 'action': action,
                'project_name': project_name, **kwargs}
        return NotificationClient._send_notification_request(
            'project', data)
```

7.5 Business Logic Layer Implementation:

The business logic layer constitutes the core of the system's functional capabilities, implementing the domain rules, workflow orchestration, and computational algorithms that distinguish the application from a simple data repository. This layer is strictly separated from the

API and data access layers, following the Single Responsibility Principle and enabling independent testing and evolution of each concern. The following subsections detail the implementation of the primary business logic modules across all three services.

7.5.1 User Management Logic:

The `UserService` class encapsulates all user-related business logic, including registration with validation, authentication with email-based fallback, and profile management. The class follows the dependency injection pattern, receiving a `UserRepository` instance through its constructor. The `register_new_user` method performs duplicate checking on both username and email before delegating the actual user creation to the repository, raising Django's `ValidationError` if either duplicate is detected. The `authenticate_user` method first attempts standard Django authentication; if that fails and the provided identifier contains an '@' character, it falls back to email-based lookup by first retrieving the `User` object by email and then authenticating with the username. This dual-path authentication strategy improves user experience by allowing login with either username or email address without requiring separate endpoints.

```
class UserService:
    def __init__(self):
        self.user_repo = UserRepository()

    def register_new_user(self, user_data: dict, password: str) -> User:
        if self.user_repo.user_exists_by_username(user_data['username']):
            raise ValidationError("Username already taken.")
        if self.user_repo.user_exists_by_email(user_data['email']):
            raise ValidationError("A user with that email already exists.")
        user = self.user_repo.create_user(
            username=user_data['username'],
            email=user_data['email'],
            password=password
        )
        return user

    def authenticate_user(self, username_or_email: str,
                          password: str) -> User:
        user = authenticate(username=username_or_email, password=password)
        if user is None and '@' in username_or_email:
            user_obj = User.objects.get(email=username_or_email)
            user = authenticate(username=user_obj.username,
                                password=password)
        if user is None or not user.is_active:
            raise PermissionDenied("invalid data")
        return user
```

7.5.2 Project and Artifact Management:

The `ProjectService` and `ArtifactService` classes manage the lifecycle of projects and their associated code artifacts. The `ProjectService` handles project creation with automatic ownership assignment, member management with role-based permission checks, and status transitions that enforce the project lifecycle rules (e.g., a project cannot move from DRAFT directly to ARCHIVED without passing through ACTIVE). The `ArtifactService` implements a more complex workflow that spans both the `UPM_Project` and `Ai_project` services: when a code artifact is uploaded, the service first stores the metadata in MySQL through the `ProjectRepository`, then delegates the actual content storage to the `Ai_project` service through the `ArtifactRepository`'s `save_content_externally` method, and finally updates the artifact's `storage_reference` field with the ID returned by the AI service. This two-phase storage approach ensures that large code files benefit from MongoDB's GridFS while maintaining referential integrity in the relational database.

7.5.3 Code Analysis Pipeline:

The code analysis pipeline is implemented as an asynchronous Celery task workflow that begins when a code file is uploaded to the Ai_project service. The `analyze_code_file_task` is the entry point: it retrieves the `CodeFile` metadata from MongoDB, reads the source content from GridFS, selects the appropriate language processor through the `ProcessorFactory`, creates a `ProjectAnalyzer` instance with the selected processor, and invokes the `analyze_code` method. The `ProjectAnalyzer` orchestrates the three-phase analysis: parsing the source code into an AST representation, extracting structural features (classes, functions, complexity metrics), and generating UML diagram data. The resulting `AnalysisResult` is serialized and inserted into the analysis results collection in MongoDB. The entire process executes asynchronously in a Celery worker, allowing the API to return immediately with a task ID that the client can poll for completion status.

```
@shared_task
def analyze_code_file_task(code_file_id_str):
    db = get_mongo_db()
    code_file_data = db[settings.CODE_FILES_COLLECTION].find_one(
        {"_id": ObjectId(code_file_id_str)})
    code_file = CodeFile(**code_file_data)
    code_content = read_from_gridfs(code_file.gridfs_id)

    processor = (PythonProcessor() if code_file.file_type == 'python'
                 else JavaProcessor())
    analyzer = ProjectAnalyzer(processor)
    analysis_result = analyzer.analyze_code(code_content)

    result_instance = AnalysisResult(
        code_file_id=ObjectId(code_file_id_str),
        status="COMPLETED",
        ast_structure=analysis_result.get('ast_structure'),
        extracted_features=analysis_result.get('features_extracted'),
        dependencies=analysis_result.get('dependencies', []),
        class_diagram_data=analysis_result.get('class_diagram_data'),
    )
    inserted = db[settings.ANALYSIS_RESULTS_COLLECTION].insert_one(
        result_instance.dict(by_alias=True, exclude_unset=True))
    return str(inserted.inserted_id)
```

7.5.4 AI Documentation Generation:

The AI documentation generation workflow is orchestrated by the `DocumentationOrchestrator`, which manages the lifecycle of documentation requests from receipt through generation to caching. When a request arrives, the orchestrator first normalizes the explanation type (`high_level` or `low_level`) and checks MongoDB for a previously generated explanation matching the analysis ID and type. If a cached explanation exists, it is returned immediately without incurring the cost of an LLM API call. If no cached explanation is found, the orchestrator selects the appropriate agent based on two dimensions: the scope (single file versus project-wide) and the level (high-level overview versus low-level detail). For single-file high-level documentation, the `HighLevelAgent` is selected; for single-file low-level documentation, the `LowLevelAgent`; and similarly for project-wide documentation, the `ProjectHighLevelAgent` and `ProjectLowLevelAgent` are used respectively. The selected agent constructs a specialized system prompt, combines it with the source code and analysis context, and delegates the actual generation to the `LLMClient`. The generated documentation is then persisted to MongoDB with a cache TTL of 48 hours, enabling efficient reuse for subsequent identical requests.

```
class DocumentationOrchestrator:
    def get_or_generate_explanation(self, exp_type):
        # Normalize type (high_level / low_level)
        normalized_exp_type = ('high_level' if is_high_level
```

```

        else 'low_level')
# Check MongoDB for cached explanation
existing = self.collection.find_one(
    {"analysis_id": self._query_id,
     "exp_type": normalized_exp_type})
if existing:
    return existing['content'], str(existing['_id'])
# Select agent based on scope and level
if is_high_level:
    agent = (ProjectHighLevelAgent() if is_project
             else HighLevelAgent())
    raw_content = agent.process(code_content,
                               analysis_summary=analysis_summary)
else:
    agent = (ProjectLowLevelAgent() if is_project
             else LowLevelAgent())
    raw_content = agent.process(code_content,
                               detailed_analysis=detailed_analysis)
# Save to MongoDB
result = self.collection.insert_one({
    "analysis_id": self._query_id,
    "exp_type": normalized_exp_type,
    "content": raw_content
})
return raw_content, str(result.inserted_id)

```

7.5.5 Notification Processing:

The notification processing workflow implements a reliable, asynchronous delivery mechanism using Celery tasks. When a notification creation request is received by the Notification_Service, it first persists the notification record with a PENDING status and then spawns a Celery task (send_project_notification_task) that handles the actual email delivery. The task retrieves the notification record from the database, creates a SimpleNotification instance with the recipient details, subject, and body, and invokes the send_email method. If delivery succeeds, the notification status is updated to SENT; if it fails, the status is set to FAILED. This pattern provides several benefits: it decouples the notification creation from the delivery mechanism, it ensures that a record of every notification attempt is maintained regardless of delivery success, and it allows for automatic retries through Celery's built-in retry mechanism. The task_id field in the Notification model provides traceability between the notification record and the Celery task execution.

```

@shared_task
def send_project_notification_task(notification_id, **kwargs):
    notification = Notification.objects.get(id=notification_id)
    notification_obj = SimpleNotification(
        recipient_email=notification.user_email,
        subject=notification.title,
        body=notification.message
    )
    if notification_obj.send_email():
        notification.status = 'SENT'
    else:
        notification.status = 'FAILED'
    notification.save()

```

7.6 User Interface Implementation:

The user interface is implemented as a single-page application using React.js, providing a responsive and interactive experience across desktop devices. The frontend communicates exclusively with the backend services through the Nginx API Gateway, which routes requests to the appropriate microservice based on the URL path prefix. This section provides an overview of the

component architecture, state management strategy, and key user interactions. Detailed screenshots of each interface are provided in the appendix.




Figure 10: Class Diagram Interface




Figure 11: Project File Tree Interface




Figure 12: Conflict Detection Interface




Figure 13: CodVSDoc Detection




Figure 14: CodVSCod Detection




Figure 15: Full Analysis Detection

7.7 Intelligent Algorithms Implementation:

This section details the implementation of the intelligent algorithms that constitute the system's core analytical and generative capabilities. These algorithms span multiple domains including static code analysis for Python and multi-language support via Tree-sitter, UML diagram generation, large language model integration, LangChain and LangGraph orchestration, multi-agent documentation orchestration, topological code processing, intelligent caching, and conflict detection. Each algorithm is described with its design rationale, implementation approach, and representative code excerpts that illustrate the key computational steps.

7.7.1 Python AST Analysis Engine:

The `PythonProcessor` class implements the `ILanguageProcessorStrategy` interface to provide deep structural analysis of Python source code through the built-in `ast` module. The analysis engine operates in three phases: source normalization, AST parsing and serialization, and feature extraction. The `parse_source_code` method first normalizes the input code by removing shebangs, handling encoding declarations, and cleaning trailing whitespace, then parses the normalized code into an AST tree using `ast.parse`. The resulting tree is serialized into a JSON-compatible dictionary structure through a custom `serialize_ast` function that recursively traverses the AST nodes, capturing node types, names, arguments, and child relationships. This serialized representation serves as the foundation for all subsequent analysis operations.

```
class PythonProcessor(ILanguageProcessorStrategy):
    def parse_source_code(self, code_content):
        normalization_result = self._normalize_python_code(code_content)
        tree = ast.parse(normalization_result['normalized_code'])
        return {
            'ast_tree': serialize_ast(tree),
            'code_content': code_content
        }

    def extract_features(self, ast_data):
        code = ast_data.get('code_content', '')
        tree = ast.parse(code)
        num_classes = sum(1 for n in ast.walk(tree)
                           if isinstance(n, ast.ClassDef))
        num_functions = sum(1 for n in ast.walk(tree)
                             if isinstance(n, ast.FunctionDef))
        return {
            'lines_of_code': len(code.splitlines()),
```



```

        'num_classes': num_classes,
        'num_functions': num_functions,
        'complexity_score': ... # cyclomatic complexity
    }

```

The `extract_features` method performs a comprehensive traversal of the AST to compute a set of quantitative metrics that characterize the code's structural complexity. By walking the entire AST tree with `ast.walk`, the method counts class definitions (ClassDef nodes), function definitions (FunctionDef nodes), and async function definitions (AsyncFunctionDef nodes). Additional metrics include total lines of code, import statement count, cyclomatic complexity estimated from control flow structures (if, for, while, try, and boolean operators), and average method length. These metrics are returned as a dictionary stored alongside the AST structure in the AnalysisResult document, providing both human-readable summaries and machine-processable data for downstream algorithms.

7.7.2 Multi-Language Parsing via Tree-sitter:

To support Java and JavaScript codebases beyond Python's native `ast` module, the system integrates Tree-sitter — an incremental, error-tolerant parsing library that generates concrete syntax trees for a wide range of programming languages. The `TreeSitterProcessor` class wraps the Tree-sitter Python bindings and implements the same `ILanguageProcessorStrategy` interface as the Python processor, ensuring a uniform analysis contract across all supported languages. At initialization, the processor loads the appropriate pre-compiled grammar (`tree_sitter_java` or `tree_sitter_javascript`) and instantiates a `Parser` object. The `parse_source_code` method encodes the source string to bytes, invokes `parser.parse()`, and converts the resulting tree into the system's internal JSON-compatible node format through a recursive walk of the Tree-sitter node object.

```

from tree_sitter import Language, Parser

class TreeSitterProcessor(ILanguageProcessorStrategy):
    def __init__(self, language: str):
        lang_obj = Language(f'build/{language}.so', language)
        self._parser = Parser()
        self._parser.set_language(lang_obj)
        self._language = language

    def parse_source_code(self, code_content: str) -> dict:
        tree = self._parser.parse(bytes(code_content, 'utf-8'))
        return {

```

```

        'ast_tree': self._node_to_dict(tree.root_node),
        'code_content': code_content,
        'language': self._language
    }

    def _node_to_dict(self, node) -> dict:
        return {
            'type': node.type,
            'text': node.text.decode('utf-8') if node.is_named else None,
            'start': node.start_point,
            'end': node.end_point,
            'children': [self._node_to_dict(c) for c in node.children]}

```

A key advantage of Tree-sitter over simpler regex-based approaches is its error-recovery capability: the parser produces a valid (partial) syntax tree even when the input code contains syntax errors, allowing the system to extract meaningful structural information from incomplete or temporarily broken files. The feature extraction pipeline for Java and JavaScript mirrors the Python pipeline, traversing the Tree-sitter CST to identify class declarations (class_declaration nodes in Java, class nodes in JavaScript), method definitions, constructor declarations, interface implementations, and inheritance hierarchies. The extracted features feed directly into the UML generation and documentation generation pipelines, enabling the multi-agent architecture to produce accurate documentation for polyglot codebases.

7.7.3 UML Diagram Generation Algorithm:

The `generate_uml_diagram` method transforms the parsed AST into a structured representation suitable for UML class diagram rendering via Mermaid.js. The algorithm operates in two passes: first, it extracts all class definitions along with their attributes, methods (with parameter signatures and return types where available), decorators, and access modifiers; second, it identifies and classifies the relationships between classes. The relationship detection algorithm examines five categories of inter-class relationships: inheritance (identified by ClassDef bases), composition (identified by attributes that are instances of other project classes), aggregation (identified by method parameters of project class types), association (identified by type annotations referencing project classes), and dependency (identified by local variable types and method call targets).

```

def generate_uml_diagram(self, ast_data):
    tree = ast.parse(ast_data.get('code_content', ''))
    classes, relationships = [], []

```

```

# Pass 1: Extract class definitions
for node in ast.walk(tree):
    if isinstance(node, ast.ClassDef):
        class_info = {
            'name': node.name,
            'attributes': self._extract_attributes(node),
            'methods': self._extract_methods(node),
            'bases': [b.id if isinstance(b, ast.Name)
                      else str(b) for b in node.bases]
        }
        classes.append(class_info)

# Pass 2: Detect relationships
for cls in classes:
    for base in cls['bases']:
        relationships.append({
            'source': cls['name'], 'target': base,
            'type': 'INHERITANCE'
        })

# ... Composition, Aggregation, Association detection
return {'classes': classes, 'relationships': relationships}

```

Additionally, the UML generation algorithm includes a design pattern detection heuristic that examines the class structure and relationships for known patterns. The detector identifies Singleton classes (by checking for instance variables and a static creation method), Observer pattern implementations (by checking for listener registration methods and notification calls), Factory patterns (by checking for abstract base classes with concrete implementations that share a creation method), and Strategy patterns (by checking for interface-like base classes with multiple implementations selected at runtime). The detected patterns are included in the UML output as metadata annotations that enhance the AI documentation generator's understanding of the code's architectural intent. The final Mermaid.js-compatible string is assembled from the extracted classes and relationships and stored alongside the AnalysisResult document.

7.7.4 Topological Code Processing:

One of the core novel contributions of this project is the topological code processing methodology, which overcomes a fundamental limitation of applying LLMs to large codebases: context window constraints. Rather than feeding files to the AI in arbitrary or alphabetical order, the system analyzes code based on its dependency graph, processing foundational modules before the

modules that depend on them. This ensures that when the LLM generates documentation for a higher-level module, it already has access to the accurate documentation of all its dependencies, producing coherent cross-file documentation with correct references.

The `DependencyGraphBuilder` class constructs a directed acyclic graph (DAG) of inter-module dependencies by parsing each file's import statements. The topological sort is then performed using Kahn's algorithm, which iteratively removes nodes with zero in-degree (no remaining unprocessed dependencies) and adds them to the processing queue. The algorithm also detects circular dependencies, which are reported as a warning and resolved by breaking the cycle at the lowest-impact edge (determined by the fewest transitive dependents).

```
class DependencyGraphBuilder:
    def build_graph(self, file_map: dict) -> nx.DiGraph:
        G = nx.DiGraph()
        for filepath, content in file_map.items():
            G.add_node(filepath)
            imports = self._extract_imports(content, filepath)
            for dep in imports:
                if dep in file_map:
                    G.add_edge(dep, filepath) # dep must be processed first
        return G

    def topological_order(self, G: nx.DiGraph) -> list:
        try:
            return list(nx.topological_sort(G))
        except nx.NetworkXUnfeasible:
            # Circular dependency detected – break cycle & warn
            cycles = list(nx.simple_cycles(G))
            self._log_cycles(cycles)
            G = self._break_cycles(G, cycles)
            return list(nx.topological_sort(G))

class CrossFileContextBuilder:
    def build_context(self, processed: dict, target_file: str,
                     G: nx.DiGraph) -> str:
        deps = list(nx.ancestors(G, target_file))
        context_parts = []
        for dep in deps:
            if dep in processed:
```

```

# Extract signatures only to minimize token usage
context_parts.append(
    self._extract_signatures(processed[dep]))
return '\n\n'.join(context_parts)

```

The `CrossFileContextBuilder` complements the topological sort by constructing a compact context string for each file being processed. Rather than passing the full source code of every dependency to the LLM (which would exhaust the context window for large projects), it extracts only the public API signatures — class names, method signatures, return types, and one-line docstrings — from already-processed dependencies. This signature extraction reduces context size by up to 85% while preserving the information the LLM needs to understand inter-module relationships. The following table summarizes the processing strategy applied at each project size tier:

Project Size	Files	Strategy	Context Mode
Small	< 20 files	Full source per file	Complete code passed to LLM
Medium	20 – 100 files	Topological + signatures	Dependency signatures only
Large	> 100 files	Topological + chunked	Chunked with sliding window
Extra Large	> 500 files	Module-level grouping	Group summaries as context

7.7.5 LLM-Powered Documentation Generation:

The `LLMClient` class implements the integration with the OpenRouter API, which provides access to multiple large language models through a unified interface. The client manages several concerns critical for production use: model selection and fallback, rate limiting and retry logic, response caching, and error handling. When a request is made, the client first checks a local cache (dictionary-based with a configurable TTL) for an identical prompt, avoiding redundant API calls for frequently requested documentation. If no cached response exists, the client constructs the API payload with the system prompt, user prompt, temperature (0.7 for balanced creativity and consistency), and a maximum token limit of 4,000.

```

class LLMClient:
    MAX_RETRIES = 5

    @staticmethod
    def call_model(system_prompt, user_prompt,
                  use_cache=True, model=None):
        if use_cache:
            cached = cache.get(cache_key)
            if cached: return cached

```

```

url = 'https://openrouter.ai/api/v1/chat/completions'
payload = {
    'model': model,
    'messages': [
        {'role': 'system', 'content': system_prompt},
        {'role': 'user', 'content': user_prompt}
    ],
    'max_tokens': 4000,
    'temperature': 0.7
}
response = requests.post(url, headers=headers,
                        json=payload, timeout=120)
if response.status_code == 429: # rate-limited
    next_model = available_models[
        (current_index + 1) % len(available_models)]
    content = result['choices'][0]['message']['content']
    if use_cache:
        cache.set(cache_key, content, 3600 * 48) # 48-hour TTL
return content

```

The client maintains a list of seven available free models and automatically switches to the next model in the list upon receiving a 429 (rate limit) response, implementing a round-robin fallback strategy that maximizes throughput while respecting API rate limits. The 48-hour TTL caching strategy is central to the system's cost efficiency: since documentation for a given code snapshot rarely needs to be regenerated within a 48-hour window, the cache eliminates redundant API calls and reduces AI API costs by up to 80% in production usage patterns.

7.7.6 Intelligent Caching Strategy:

The intelligent caching system operates at two distinct levels: response-level caching and fingerprint-based staleness detection. Response-level caching stores the raw LLM output keyed by a SHA-256 hash of the combined system prompt, user prompt, and model identifier. This ensures that semantically identical requests — even if submitted by different users — return cached results without incurring additional API costs. The cache is implemented using Django's cache framework with Redis as the backend, providing both high-throughput read access and configurable TTL management.

```

import hashlib
from django.core.cache import cache

```

```

class CacheManager:
    DEFAULT_TTL = 3600 * 48    # 48 hours

    @staticmethod
    def make_key(system_prompt: str, user_prompt: str,
                  model: str) -> str:
        raw = f'{system_prompt}||{user_prompt}||{model}'
        return 'llm:' + hashlib.sha256(raw.encode()).hexdigest()

    @classmethod
    def get(cls, key: str):
        return cache.get(key)

    @classmethod
    def set(cls, key: str, value: str, ttl: int = DEFAULT_TTL):
        cache.set(key, value, timeout=ttl)

class FingerprintBuilder:
    """Detects whether cached documentation is still valid."""

    @staticmethod
    def code_fingerprint(code_content: str) -> str:
        # Normalize whitespace & comments before hashing
        normalized = FingerprintBuilder._strip_comments(code_content)
        return hashlib.sha256(normalized.encode()).hexdigest()

    @staticmethod
    def is_stale(stored_fp: str, current_code: str) -> bool:
        return stored_fp != FingerprintBuilder.code_fingerprint(current_code)

```

The second caching level — fingerprint-based staleness detection — addresses the documentation drift problem. When a code file is modified after its documentation has been generated and cached, the system must detect that the cached documentation is outdated and regenerate it. The `FingerprintBuilder` computes a normalized SHA-256 hash of the source code (with comments stripped and whitespace normalized to eliminate cosmetic-only changes) and stores it alongside the cached documentation. On each analysis request, the current code fingerprint is compared with the stored one; if they differ, the cache entry is invalidated and a fresh generation request is dispatched. This mechanism ensures documentation stays synchronized with the actual codebase without requiring manual invalidation.

7.7.7 LangChain and LangGraph Integration:

The system adopts a LangGraph-inspired orchestration architecture to provide a structured, maintainable, and extensible pipeline for all LLM-driven workflows. The design is modelled on LangGraph's StateGraph pattern, where each agent corresponds to a node and transitions between agents are represented as conditional routing functions operating on a shared Pydantic state object. Importantly, the implementation does not take a hard dependency on the langgraph or langchain packages; instead, a lightweight internal runner (`_SimpleGraph`) replicates the same interface, so the application runs in environments where those libraries are not installed while remaining structurally compatible if they are added in the future.

```
# agents.py – optional LangChain import with graceful fallback
try:
    from langchain.tools import BaseTool
    from langchain_core.prompts import ChatPromptTemplate
    _LANGCHAIN_AVAILABLE = True
except Exception:
    ChatPromptTemplate = None # graceful fallback
    _LANGCHAIN_AVAILABLE = False

# orchestrator.py – lightweight graph runner (mirrors LangGraph StateGraph)
class _SimpleGraph:
    def invoke(self, state: DocumentationState) -> DocumentationState:
        route = _route_generation(state) # -> high_level | low_level | project_*
        state = node_map[route](state) # run chosen agent node
        if _should_verify(state) == 'verify':
            state = node_map['verify'](state)
        return state
```

The internal graph runner mirrors the LangGraph StateGraph pattern. Its nodes correspond to the five agents in the multi-agent system: `HighLevelAgent`, `LowLevelAgent`, `ProjectHighLevelAgent`, `ProjectLowLevelAgent`, and `VerifierAgent`. The shared graph state is modelled as a Pydantic BaseModel (`DocumentationState`) rather than a TypedDict, carrying the source code, AST-derived analysis summaries, intermediate documentation outputs, and any verification feedback. Routing is handled by two plain Python functions: `_route_generation` selects the appropriate agent based on scope and level, while `_should_verify` decides whether to invoke the `VerifierAgent`.

```
from pydantic import BaseModel
from typing import Optional
```



```

class DocumentationState(BaseModel):
    code_content: str = ''
    analysis_summary: Optional[str] = None
    raw_output: Optional[str] = None
    verified_output: Optional[str] = None
    error: Optional[str] = None
    is_project: bool = False
    exp_type: str = 'high_level'

def _route_generation(state: DocumentationState) -> str:
    if state.is_project:
        return 'project_high_level' if state.exp_type == 'high_level' \
            else 'project_low_level'
    return 'high_level' if state.exp_type == 'high_level' else 'low_level'

def _should_verify(state: DocumentationState) -> str:
    return 'done' if state.error or not state.raw_output else 'verify'

```

7.7.8 Multi-Agent Documentation System:

The multi-agent documentation system implements a specialized agent architecture where each agent is designed to produce documentation at a specific level of abstraction and for a specific scope. The base `BaseAgent` class provides the common `ask_ai` method that delegates to the `LLMClient`, while the concrete agent classes define their own system prompts and prompt construction logic. The `HighLevelAgent` produces architectural overviews that focus on the purpose, design patterns, and inter-component relationships, using a system prompt that instructs the LLM to adopt the persona of a Senior Software Architect. The `LowLevelAgent` produces detailed logic analyses examining individual method implementations, data flow, error handling, and edge cases.

```

class HighLevelAgent(BaseAgent):
    def process(self, code_content, class_name=None,
                analysis_summary=None):
        system = ('You are a Senior Software Architect. '
                  'Provide a HIGH-LEVEL overview focused on purpose, '
                  'patterns, and inter-component relationships.')
        prompt = (f'Code:\n{code_content}\n\n'
                  'Provide the high-level architectural report.')
        return self.ask_ai(system, prompt)

class LowLevelAgent(BaseAgent):

```

```

def process(self, code_content, detailed_analysis=None):
    system = ('You are a Senior Technical Architect. '
              'Perform a DEEP LOGIC ANALYSIS of methods, '
              'data flow, error handling, and edge cases.')
    prompt = (f'Context:\n{detailed_analysis}\n\n'
              f'Code:\n{code_content}')
    return self.ask_ai(system, prompt)

class VerifierAgent(BaseAgent):
    def process(self, code_content, raw_documentation):
        system = ('You are a Documentation Quality Auditor. '
                  'Verify the documentation against the actual code. '
                  'Flag missing methods, wrong parameter counts, '
                  'and outdated class relationships.')
        prompt = (f'Code:\n{code_content}\n\n'
                  f'Documentation to verify:\n{raw_documentation}')
        return self.ask_ai(system, prompt)

```

The agent selection logic implemented in the DocumentationOrchestrator follows a two-dimensional matrix: the first dimension is the scope (single file versus project-wide), and the second is the level (high versus low). This results in four agent combinations, with a fifth VerifierAgent available for quality assurance. The VerifierAgent operates by comparing the documentation claims against AST-derived facts, flagging discrepancies such as missing method descriptions, incorrect parameter counts, or outdated class relationships. Comparative internal testing demonstrated that this multi-agent architecture achieves a 73% reduction in technical error rates compared to single-LLM approaches, with format consistency improved to 95%.

Scope	Level	Agent Used	Primary Output
Single File	High	HighLevelAgent	Architectural overview, design patterns
Single File	Low	LowLevelAgent	Method-level analysis, data flow
Project-wide	High	ProjectHighLevelAgent	Cross-module architecture summary
Project-wide	Low	ProjectLowLevelAgent	Full project technical documentation
Any	Verification	VerifierAgent	Corrected & validated documentation

7.7.9 Conflict Detection Algorithm:

The conflict detection algorithm is implemented as a multi-stage pipeline within the conflict_detection subsystem of the Ai_project service. The pipeline consists of five sequential stages, each implemented as an independent module with a well-defined input and output contract. The first stage, AST Unification (ast_unifier.py), normalizes the AST representations of two code

versions into a common intermediate format that enables structural comparison across different programming languages. The second stage, Structural Diff (`structural_diff.py`), performs a tree-diffing algorithm that identifies added, removed, and modified nodes between the two unified ASTs, producing a structural change set.

The third stage, Semantic Analysis (`semantic_analyzer.py`), classifies each structural change by its semantic impact — whether it affects the public API, the internal logic, the data flow, or the control flow — using heuristics derived from the node types and their positions in the AST hierarchy. The fourth stage, Code-Documentation Consistency (`code_doc_consistency.py`), cross-references the detected code changes against the existing AI-generated documentation to identify inconsistencies, using the `FingerprintBuilder` to create content fingerprints and the `RenameDetector` to track variable and method renames across versions.

```
class ConflictDetectionPipeline:
    def __init__(self):
        self.stages = [
            ASTUnifier(),
            StructuralDiff(),
            SemanticAnalyzer(),
            CodeDocConsistency(),
            ReportGenerator()
        ]

    def run(self, old_code: str, new_code: str,
           existing_docs: str) -> ConflictReport:
        state = PipelineState(
            old_code=old_code,
            new_code=new_code,
            existing_docs=existing_docs
        )
        for stage in self.stages:
            try:
                state = stage.process(state)
            except StageException as e:
                # Graceful degradation: return partial results
                state.errors.append(str(e))
                if stage.is_critical:
                    break
        return state.report
```

```

class ReportGenerator:
    SEVERITY_LEVELS = ['critical', 'warning', 'informational']

    def process(self, state: PipelineState) -> PipelineState:
        conflicts = []
        for change in state.semantic_changes:
            severity = self._assess_severity(change)
            conflicts.append(ConflictItem(
                severity=severity,
                location=change.location,
                description=change.description,
                suggested_action=self._suggest_action(change)
            ))
        state.report = ConflictReport(conflicts=conflicts,
                                     errors=state.errors)

        return state

```

The fifth and final stage, Report Generation (`report_generator.py`), aggregates the findings from all previous stages into a structured `ConflictReport` that categorizes each conflict by severity (critical, warning, or informational), provides the specific code locations and documentation references involved, and includes suggested resolution actions. The entire pipeline is orchestrated by the conflict detection orchestrator, which manages the data flow between stages and handles error recovery through a fallback mechanism that returns partial results if any non-critical stage fails, ensuring that the conflict detection system degrades gracefully rather than failing completely.

Stage	Module	Input	Output
1 – AST Unification	<code>ast_unifier.py</code>	Raw source (old & new)	Unified AST pair
2 – Structural Diff	<code>structural_diff.py</code>	Unified AST pair	Change set (added/removed/modified)
3 – Semantic Analysis	<code>semantic_analyzer.py</code>	Change set	Classified semantic impacts
4 – Doc Consistency	<code>code_doc_consistency.py</code>	Semantic impacts + existing docs	Inconsistency list
5 – Report Generation	<code>report_generator.py</code>	Inconsistency list	<code>ConflictReport</code> (severity + actions)

Chapter 7.5: AI Methodology: Prompt Engineering & Evaluation

7.5.1. Introduction:

The research methodology is structured as a four-phase pipeline that transforms raw source code into verified, evaluated documentation. This design addresses three key challenges in AI-generated documentation: factual accuracy (ensuring class and method names are correctly referenced), completeness (covering all code elements), and consistency (detecting conflicts between documentation versions and code changes). The pipeline architecture draws from continuous integration patterns where code changes trigger automated analysis and verification.




Figure 16: Research Methodology-4Phase Pipeline

7.5.2. Phase 1: Code Analysis and Feature Extraction:

The first phase transforms source code into structured representations through Abstract Syntax Tree (AST) analysis. The system implements a factory pattern (`processor_factory.py`) that routes code files to appropriate processors: `PythonProcessor` (Python's `ast` module), `TreeSitterProcessor` (Java, C#, TypeScript via `tree-sitter`), and `GenericProcessor` (fallback). Each processor extracts class names, method signatures, attributes, inheritance relationships, and route definitions.

7.5.2.1 AST Unification and Class Diagram Generation:

The ASTUnifier component (ast_unifier.py) normalizes different AST representations into a unified UnifiedClassDiagram structure, essential for consistent comparison in conflict detection. The ContextBuilder service enriches AST analysis with function_names lists, route definitions, and lines_of_code counts.

7.5.3 Phase 2: Documentation Generation Pipeline:

The second phase implements a multi-agent documentation generation system using five specialized AI agents orchestrated through a LangGraph-style directed graph. The DocumentationOrchestrator routes code files to the appropriate agent based on is_project and exp_type routing decisions. After generation, the _should_verify() function determines whether verification is needed, routing to VerifierAgent when raw_output exists. Auto-trigger evaluation initiates automatically as a non-blocking Celery task.




Figure 17: AI Documentation System Architecture

7.5.3.1 Agent Design and Specialization:

Five agent classes are defined: HighLevelAgent (stakeholder-friendly documentation), LowLevelAgent (deep technical documentation), ProjectHighLevelAgent (multi-file project overview), ProjectLowLevelAgent (comprehensive project-level technical analysis), and VerifierAgent (quality assurance with 50,000 character SIZE_LIMIT). Each agent includes anti-hallucination rules, few-shot examples, and chain-of-thought instructions.

7.5.3.2 LLM Client and Model Management:

The LLMClient provides a unified interface to the OpenRouter API, managing seven free models with intelligent fallback rotation. The default model is openai/gpt-oss-120b:free (131K context), the LLM Judge is meta-llama/llama-3.3-70b-instruct:free, and the Conflict Evaluator is qwen/qwen3-coder:free. Fallback handles 404 errors with model rotation and rate limits with 60-minute blocking periods. Response caching uses Django LocMemCache with MD5 hash keys and 48-hour TTL.

Table 22: Available Free Models on OpenRouter

Model ID	Context	Role
openai/gpt-oss-120b:free	131K	Default Generation
nvidia/nemotron-nano-9b-v2:free	32K	Fast Generation
google/gemma-3-27b-it:free	8K	High Accuracy
meta-llama/llama-3.3-70b-instruct:free	8K	LLM Judge (L3)
meta-llama/llama-3.2-3b-instruct:free	4K	Fast Fallback
mistralai/mistral-7b-instruct:free	32K	General Purpose
mistralai/devstral-small-2505:free	32K	Code-Specialized

7.5.4 Phase 3: Four-Layer Evaluation Framework:

The evaluation framework implements a four-layer assessment system combining automated metrics with LLM-based judgment and optional human review. No single evaluation metric can adequately capture the quality of AI-generated documentation, motivating this multi-layer design.




Figure 18: Layer Evaluation Framework

7.5.4.1 Layer 1: Factual Accuracy (35%);

Verifies factual accuracy by checking whether class and method names from AST analysis appear in the generated documentation. Score = verified names / total names. Testing averaged

0.732 (73.2%), with 30 out of 41 entities correctly referenced. The 35% weight reflects factual accuracy as the most critical quality dimension.

7.5.4.2 Layer 2: Completeness (25%):

Measures documentation completeness as the coverage ratio of documented code elements. Test results show 0.714 (71.4%), with 30 out of 42 elements covered.

7.5.4.3 Layer 3: LLM Judge (40%):

Employs LLM-as-a-Judge using meta-llama/llama-3.3-70b-instruct:free evaluating across five criteria: structure_adherence (8.2/10), factual_accuracy (7.8/10), logic_flow_quality (7.5/10), completeness (7.0/10), and clarity (8.5/10). Overall score: 7.6/10 (0.760 normalized). The 40% weight reflects LLM judgment providing the most nuanced evaluation.

Table 23: LLM Judge Evaluation Results

Criterion	Score (1-10)	Assessment
Structure Adherence	8.2	Well-organized format
Factual Accuracy	7.8	Most entities correct
Logic Flow Quality	7.5	Some gaps in transitions
Completeness	7	Missing some details
Clarity	8.5	Clear and readable
Overall Score	7.6	GOOD quality level

7.5.4.4 Layer 4: Human Review (20%, Optional):

Incorporates human feedback through MongoDB-stored human_feedback collection. Three reviewers provided scores of 0.82, 0.78, and 0.85, averaging 0.817. When human reviews are present, final score formula adjusts to: $L1 \times 0.30 + L2 \times 0.20 + L3 \times 0.30 + \text{Human} \times 0.20$. Without: $L1 \times 0.35 + L2 \times 0.25 + L3 \times 0.40$.

7.5.4.5 Final Score and Verdict:

The verdict system classifies quality into four levels: EXCELLENT (≥ 0.85), GOOD (≥ 0.70), ACCEPTABLE (≥ 0.55), and NEEDS_IMPROVEMENT (< 0.55). Primary evaluation achieved 0.739 (GOOD) without human review and 0.754 with human review.

Test Case	L1	L2	L3	Final	Verdict
HL Single (agents.py)	0.82	0.79	0.74	0.72	GOOD
LL Single (evaluation.py)	0.86	0.82	0.78	0.76	GOOD
Proj HL (Full System)	0.79	0.74	0.71	0.7	ACCEPTABLE
Proj LL (Full System)	0.91	0.88	0.82	0.83	GOOD

LL + Human Review	0.86	0.82	0.78	0.78	GOOD
-------------------	------	------	------	------	------

7.5.5 Phase 4: Conflict Detection System:

The conflict detection system identifies discrepancies between different code versions and between code and documentation. The ConflictOrchestrator implements three analysis modes: `code_vs_code` (structural + semantic comparison), `code_vs_doc` (documentation drift detection), and `full_analysis` (combined with merge and deduplication).




Figure 19: Conflict Detection Pipeline

7.5.5.1 FingerprintBuilder and Rename Detection:

The FingerprintBuilder constructs fingerprints for each method: `params:{count}|ptypes:{types}|return:{type}|abstract:{bool}|constructor:{bool}|visibility:{vis}`. The RenameDetector classifies changes as rename (same fingerprint + different name, confidence 0.92) or logic change (same name + different fingerprint, confidence 0.88) with similarity threshold 0.8.

7.5.5.2 Fusion Reasoning Semantic Analysis:

The SemanticAnalyzer employs Fusion Reasoning, instructing the LLM to analyze each code diagram independently, fuse the analyses, and detect conflicts across 13 semantic categories including method signatures, return types, logic changes, API contract violations, data integrity, performance, security, class removal/addition, and more.

7.5.5.3 Report Generation and Grade System:

The ReportGenerator produces a conflict report with grades A through F: Grade A (0-4 conflicts, 0-2 breaking), B (5-9 or 1-2 breaking), C (10+ or 3-4 breaking), D (3-4 breaking), F (5+ breaking). Each conflict includes type, severity, class_name, member, description, recommendation, and confidence score.

7.5.6 Conflict Detection Evaluation (LLM-as-a-Judge):

A separate LLM-as-a-Judge evaluation pipeline using the ConflictEvaluator class with qwen/qwen3-coder:free model at temperature 0.3 assessed conflict detection results across six criteria on a 0-100 scale.

7.5.6.1 JUDGE_PROMPT Design:

The JUDGE_PROMPT evaluates across six criteria: Correctness, Completeness, Severity Accuracy, Confidence Calibration, Clarity and Usefulness, and Overall Quality. The verdict system uses five levels: Excellent (90-100), Very Good (80-89), Good (70-79), Acceptable (60-69), and Poor (<60).

7.5.6.2 Conflict Evaluation Results:

The ConflictEvaluator on code_vs_code comparing v1.0 and v2.0 produced an overall score of 85/100 (Very Good). The detailed results are presented in Table 4.5.

Table 24: Conflict Evaluation Results (LLM-as-a-Judge)

Criterion	Score (0-100)	Assessment
Correctness	90	Accurate detection of real conflicts
Completeness	80	Missed 3 additions
Severity Accuracy	85	Appropriate severity levels

Confidence Calibration	85	Reasonable confidence scores
Clarity	95	Clear descriptions and recommendations
Usefulness	80	Actionable but incomplete coverage
Overall Score	85	Very Good

7.5.6.3 Strengths and Weaknesses:

Strengths: (1) Accurately detected the breaking change in LLMClient.call_model return type change from dict to LLMResponse; (2) Correctly identified weight formula changes in ExplanationEvaluator; (3) Properly captured threshold change in FingerprintBuilder.similarity (0.7 to 0.8); (4) Clear and actionable recommendations (Clarity score: 95). Weaknesses: Missed three additions – the CONFLICT_MODEL constant, the _rotate_model_on_404 method in LLMClient, and the entirely new VerifierAgent class, reducing Completeness to 80. The system focuses primarily on modifications and removals rather than additions.

Table 25: Missed Conflicts Identified by the Judge

Type	Class	Description
NEW_FEATURE	LLMClient	New CONFLICT_MODEL constant: qwen/qwen3-coder:free
NEW_FEATURE	LLMClient	New _rotate_model_on_404 method for model rotation
NEW_FEATURE	VerifierAgent	Completely new class with verify() and SIZE_LIMIT

7.5.7 System Architecture Overview:

The complete system integrates all four phases into a Django-based web application with Celery task queue support for asynchronous processing. MongoDB serves as the primary data store, while SQLite handles Django admin operations. The technology stack was selected based on cost efficiency (all LLM models are free), scalability (Celery enables horizontal scaling), and maintainability (Django's mature ecosystem).

Table 26: System Technology Stack

Component	Technology	Purpose
Backend	Django + DRF	API endpoints
Task Queue	Celery + Redis	Async processing
Database	MongoDB + SQLite	Data persistence
LLM Provider	OpenRouter API	Free model access
AST Parsing	Python ast + tree-sitter	Code analysis
Validation	Pydantic v2	Data models
Caching	Django LocMemCache	Response caching (MD5, 48h)
Orchestration	LangGraph-style graph	Pipeline routing

7.5.8 Summary:

This chapter presented the four-phase research methodology. Phase 1 (Code Analysis) uses AST parsing with multiple language processors to extract structured representations. Phase 2 (Documentation Generation) employs five specialized AI agents orchestrated through a LangGraph-style pipeline, with auto-trigger evaluation and LLM client fallback management. Phase 3 (Evaluation) implements a four-layer assessment framework achieving 73.9% (GOOD) without human review and 75.4% with human review. Phase 4 (Conflict Detection) provides three analysis modes with Fusion Reasoning across 13 semantic categories and a grade system (A-F).

The conflict detection evaluation using LLM-as-a-Judge achieved 85/100 (Very Good), with strong performance in Correctness (90) and Clarity (95), but identified a gap in Completeness (80) due to missed additions. These results validate the multi-layer evaluation approach and Fusion Reasoning technique, while highlighting opportunities for improvement in addition detection capabilities.

Chapter 8: Testing and Evaluation

8.1 Test Plan:

The test plan defines what will be tested, when, with which tools, and who is responsible for execution.

8.1.1 Test Scope:

- ATDG System — covering UPM (User & Project Management) and AI (Code Analysis & Conflict Detection) services.
- Notification Service — verifying integration with both main services.
- REST API Endpoints — all endpoints exposed by both services.
- Conflict Detection Feature — the core differentiating capability of the system.

8.1.2 What Will Be Tested:

Table 27: Test Plan

Dimension	Details
Unit	Each API endpoint tested independently
Integration	Interaction between UPM, AI, and Notification services
Performance	Response times under 10 / 50 / 100 concurrent users
Security	JWT authentication and RBAC access control
Acceptance	Verification that requirements agreed with supervisor are met

8.1.3 Testing Tools:

Tool	Purpose	Test Type
Locust (Python)	Performance, load, and stress testing	Performance / Load / Stress
Django REST Framework Test Client	Unit and integration testing	Unit / Integration
Postman	Manual endpoint testing and response validation	Manual / Functional
JWT Debugger	Token validity and expiry verification	Security
Docker & Docker Compose	Isolated and reproducible test environment	Environment
pytest	Automated unit testing inside Django	Unit

8.1.4 Test Schedule:

Phase	Test Type	Duration	Expected Outcome
Phase 1	Unit Testing	2 weeks	Coverage $\geq$ 80% per endpoint
Phase 2	Integration Testing	1 week	Data flow verified between services
Phase 3	Performance / Load	3 days	Zero failures at 10–100 users
Phase 4	Security Testing	3 days	All unauthorised access blocked
Phase 5	Acceptance Testing	1 week	Supervisor approves all requirements

8.2 Types of Testing:

8.2.1 Unit Testing:

- POST /api/upm/signup/ — user creation and data validation
- POST /api/token/ — authentication and JWT issuance
- Project CRUD — create, read, update, delete

- POST /api/analysis/detect-conflict/ — conflict detection request reception
- GET /api/analysis/conflict-status/<id>/ — correct status return

8.2.2 Integration Testing:

Table 28: Types of Testing

Integration Scenario	Services Involved	Expected Outcome
User registration triggers notification	UPM → Notification	User created + async notification sent
Project creation triggers notification	UPM → Notification	Project created + supervisor notified
Upload code → analyse → detect conflict	UPM → AI → Celery	Full chain from upload to result
GitHub sync → MongoDB update	UPM → GitHub API → MongoDB	Code fetched and saved to GridFS
Notification service failure non-blocking	UPM (Notification down)	HTTP 201 returned; error logged only

8.2.3 System Testing:

Verifies that the system as a whole satisfies functional requirements end-to-end.

- Register → Login → Create Project → Upload Code → Request Analysis → Receive Results
- Upload two project versions → Request conflict detection → View report
- Admin creates Reviewer → Reviewer evaluates explanation → View statistics

8.2.4 Acceptance Testing:

Acceptance Criterion	Minimum Threshold	Actual Result	Status
Request success rate — Performance	1	100% (546 / 546)	Pass
Request success rate — Load	1	100% (6,066 / 6,066)	Pass
Request success rate — Stress	≥ 95%	100% (19,363 / 19,363)	Pass
UPM median response — Performance	< 500 ms	36–200 ms	Pass
Conflict Detection returns result	Yes	Yes — Celery + LLM	Pass
Zero failures in Performance test	0 failures	0 failures	Pass
RBAC access control enforced	Admin/Dev/Rev	Fully implemented	Pass

8.2.5 Security Testing:

Test Scenario	Expected Response	Status
Request without JWT	HTTP 401 Unauthorized	Pass
Expired JWT token	HTTP 401 — Token expired	Pass
Developer accessing Admin endpoints	HTTP 403 Forbidden	Pass
Admin deactivating own account	HTTP 403 — Cannot deactivate own account	Pass
Admin deactivating another Admin	HTTP 403 — Cannot deactivate another Admin	Pass

User accessing another user's project	HTTP 404 — Not found (filtered by ownership)	Pass
---------------------------------------	--	------

8.2.6 Performance, Load & Stress Testing:

Tool: Locust — Python-based HTTP load testing framework
Environment: Docker Compose on local machine PostgreSQL + MongoDB + Redis + Celery
Definition: Performance = 10 users Load = 50 users Stress = 100 concurrent users

8.3 Test Diagrams:

8.3.1 Test Case Diagrams:

The test suite covers 102 test cases distributed across both services and the Notification Service:

Table 29: Test Diagrams

Section	Test Cases	Service
SEC-01: Authentication & Authorisation	8	UPM
SEC-02: User Registration	5	UPM
SEC-03: Password Management	5	UPM
SEC-04: Project Management	9	UPM
SEC-05: Folder Management	4	UPM
SEC-06: Admin Operations	6	UPM
SEC-07: GitHub Integration	4	UPM
SEC-08: Online Users & Project Tree	4	UPM
AI-01: Code File Management	6	AI
AI-02: Code Analysis	11	AI
AI-03: AI Explanations & Evaluation	8	AI
AI-04: Export & Document Generation	5	AI
AI-05: Conflict Detection ★ (Core)	15	AI
AI-06: File Versions & Project Tree	6	AI
AI-07: Reviewer Statistics	3	AI
NOTIF: Notification Service	3	Notification
Total	102	All

8.3.2 Test Result Diagrams:

The following charts show response time results (Median, P95, Max) across all endpoints for each test type. Yellow bars / highlighted rows denote Conflict Detection endpoints.

● UPM Service — User & Project Management (Port 8001)

Test Conditions: 10 concurrent users | Duration: ~2 min | Total requests: 546 | **Failures: 0** ✓

All endpoints responded within excellent thresholds. Login achieved a 56 ms median; most CRUD operations stayed well below 200 ms. The /tree/ endpoint showed a higher Max (2,496 ms) due to recursive folder traversal — acceptable for that operation's complexity.




Figure 20: UPM Service: Performance Test (10 concurrent users) | Response Times per Endpoint




Figure 21: UPM Service: Performance Test (10 Users) — Response Time & Active Clients

Result: All endpoints rated Excellent (Median < 500 ms). P95 values remained below 530 ms for all standard operations.

Test Conditions: 50 concurrent users | Duration: ~5 min | Total requests: 6,066 | **Failures: 0** ✓

Scaling to 50 simultaneous users, the UPM service maintained zero failures. Median response times remained within acceptable bounds for all endpoints. The signup endpoint (770 ms median) reflects heavier DB operations. P95 values ranged from 1,300–3,400 ms — expected queue build-up under moderate concurrency.




Figure 22: UPM Service: Load Test (50 concurrent users) | Response Times per Endpoint



Figure 23: UPM Service: Load Test (50 Users) — Response Time & Active Clients

Result: Service stable. RPS reached 20.31 req/s. Signup shows expected slowdown under load (Acceptable).

Test Conditions: 100 concurrent users | Duration: ~10 min | Total requests: 19,363 | **Failures: 0** ✓

Even at maximum stress, the UPM service recorded zero failures. Median response times increased proportionally (440–540 ms for most endpoints). Signup at 4,600 ms median reflects expected degradation for heavy operations. The service did not crash or drop any request.




Figure 24: UPM Service: Stress Test (100 concurrent users) | Response Times per Endpoint




Figure 25: UPM Service: Stress Test (100 Users) — Response Time & Active Clients

Result: Service operational with Acceptable latency at extreme concurrency. Peak RPS: 32.30 req/s. No failures recorded.

AI Service — Code Analysis & Conflict Detection (Port 8002)

★ **Conflict Detection:** Endpoints highlighted in yellow throughout the charts belong to the Conflict Detection feature — the core intelligence capability of the system. These endpoints are embedded within the AI service.

* **Note:** The single failure in the Load test occurred at `/project-class-diagram/` due to local Docker resource limits, not application code errors. The AI service processes LLM (Large Language Model) operations which are inherently time-intensive.

Test Conditions: 10 concurrent users | Duration: ~2 min | Total requests: 122 | **Failures: 0** ✓

Under baseline conditions, all 122 requests succeeded. Simple endpoints (`analyze/`, `conflict-status/`) responded in under 100 ms median. LLM-dependent operations such as `generate-explanation/` show expected variance (72 ms median, 39 s P95). Conflict Detection endpoints returned 2,400–3,300 ms median — consistent with LLM inference time at low concurrency.




Figure 26: AI Service: Performance Test (10 concurrent users) | Response Times per Endpoint (yellow = Conflict Detection)




Figure 27: AI Service: Performance Test (10 Users) — Response Time & Active Clients

Result: All 44 AI endpoints responded successfully. Core Conflict Detection operations achieved 2,381–3,300 ms median — within LLM processing expectations.

Test Conditions: 50 concurrent users | Duration: ~5 min | Total requests: 424 | **Failures: 1*** ⚠️ (Docker resource limit)

At 50 concurrent users, response times increased due to the sequential nature of LLM inference. Median times ranged from 14,000–45,000 ms, reflecting queuing when concurrent load exceeds local processing capacity. The single failure (/project-class-diagram/) was caused by a Docker memory constraint — not a code defect. Conflict Detection endpoints remained fully functional throughout.




Figure 28: AI Service: Load Test (50 concurrent users) | Response Times per Endpoint (yellow = Conflict Detection)




Figure 29: AI Service: Load Test (50 Users) — Spike indicates Docker memory constraint at ~0:30

Result: 1 failure attributed to local infrastructure. All Conflict Detection endpoints remained functional. High latency is expected for LLM-backed operations under concurrent load.

Test Conditions: 100 concurrent users | Duration: ~10 min | Total requests: 806 | **Failures: 0** ✓

At maximum stress, the AI service reached the Locust 90-second timeout threshold across most endpoints, reflecting LLM queue saturation on local Docker infrastructure. Critically, zero requests were rejected — the service remained stable and continued processing. Conflict Detection endpoints similarly reached the 90 s timeout but produced no errors.




Figure 30: AI Service: Stress Test (100 concurrent users) | Response Times per Endpoint (yellow = Conflict Detection)




Figure 31: AI Service: Stress Test (100 Users) — Response time reaches Locust 90s timeout ceiling

Result: Zero failures recorded. Response times hit Docker/local LLM throughput ceiling. Service architecture proved resilient — no crashes or dropped connections under extreme concurrency.

8.4 Test Results and Analysis:

8.4.1 UPM Service — Test Summaries:

Table 30: Test Results and Analysis

Test Type	Duration	Total Requests	Failures	Avg (ms)	Median (ms)	P95 (ms)	RPS
Performance (10 users)	~2 min	546	0	82	57	200	4.61
Load (50 users)	~5 min	6066	0	393	140	1600	20.31
Stress (100 users)	~10 min	19363	0	1059	490	4200	32.3

8.4.2 UPM Performance Test — Endpoint Detail:

Endpoint	Method	Requests	Failures	Median (ms)	P95 (ms)	Max (ms)
/api/upm/signup/	POST	10	0	200	520	519
/api/token/ [login]	POST	59	0	56	140	162
/api/upm/projects/ [create]	POST	96	0	89	210	305

/api/upm/projects/ [list]	GET	155	0	48	180	352
/api/upm/projects/<id>/	GET	86	0	36	110	265
/api/upm/projects/<id>/folders/	GET	47	0	56	200	227
/api/upm/projects/<id>/tree/	GET	46	0	48	260	2496
/api/upm/projects/<id>/versions/	GET	47	0	58	130	1073
Aggregated		546	0	57	200	2496

8.4.3 UPM Load Test — Endpoint Detail:

Endpoint	Method	Requests	Failures	Median (ms)	P95 (ms)	Max (ms)
/api/upm/signup/	POST	50	0	770	3400	3687
/api/token/ [login]	POST	617	0	110	1600	6967
/api/upm/projects/ [create]	POST	1126	0	190	1700	8446
/api/upm/projects/ [list]	GET	1677	0	140	1600	8709
/api/upm/projects/<id>/	GET	1057	0	110	1400	8771
/api/upm/projects/<id>/folders/	GET	528	0	120	1300	7696
/api/upm/projects/<id>/tree/	GET	524	0	120	1300	6695
/api/upm/projects/<id>/versions/	GET	487	0	140	2100	8765
Aggregated		6066	0	140	1600	8771

8.4.4 UPM Stress Test — Endpoint Detail:

Endpoint	Method	Requests	Failures	Median (ms)	P95 (ms)	Max (ms)
/api/upm/signup/	POST	100	0	4600	11000	11808
/api/token/ [login]	POST	1856	0	460	4800	11600
/api/upm/projects/ [create]	POST	3590	0	540	4100	11110
/api/upm/projects/ [list]	GET	5214	0	490	3900	11610
/api/upm/projects/<id>/	GET	3479	0	440	4000	8648
/api/upm/projects/<id>/folders/	GET	1747	0	460	4000	8689
/api/upm/projects/<id>/tree/	GET	1705	0	470	4000	8462
/api/upm/projects/<id>/versions/	GET	1672	0	470	4100	8490
Aggregated		19363	0	490	4200	11808

8.4.5 AI Service — Test Summaries:

Test Type	Duration	Total Requests	Failures	Avg (ms)	Median (ms)	P95 (ms)	RPS
Performance (10 users)	~2 min	122	0	6896	3100	27000	1.1
Load (50 users)	~5 min	424	1*	31763	32000	64000	1.42
Stress (100 users)	~10 min	806	0	67405	83000	90000	1.35

* 1 failure in Load = local Docker memory constraint, not application code.

8.4.6 Conflict Detection Endpoints — Cross-Test Results ★:

Endpoint	Method	Perf Median	Load Median	Stress Median	Status
detect-conflict/ [code_vs_code]	POST	3,300 ms	32,000 ms	90,000 ms	Pass
detect-conflict/ [code_vs_doc]	POST	2,381 ms	31,000 ms	68,000 ms	Pass
detect-conflict/ [full_analysis]	POST	2,400 ms	32,000 ms	89,000 ms	Pass
detect-conflict/ [via_versions]	POST	N/A	46,000 ms	90,000 ms	Pass
conflict-status/<id>/	GET	97 ms	32,000 ms	77,000 ms	Pass
conflict-status/ [query_param]	GET	N/A	34,000 ms	69,000 ms	Pass
conflict-result/<id>/	GET	250 ms	29,000 ms	62,000 ms	Pass
conflict-history/ [project_id]	GET	4,115 ms	27,000 ms	82,000 ms	Pass
conflict-history/ [file_id]	GET	15,160 ms	33,297 ms	80,000 ms	Pass
conflict-history/ [version_id]	GET	3,900 ms	45,000 ms	57,000 ms	Pass

Note: 90,000 ms in Stress = Locust timeout ceiling. Requests queued — not rejected. All endpoints remained functional under extreme load.

8.4.7 Result Analysis:

Dimension	Observation	Rating
UPM — Failure Rate	Zero failures across 25,975 requests in all three tests	Excellent
UPM — Response Time	Median 57 ms (Performance) → 490 ms (Stress)	Excellent
AI — Failure Rate	Only 1 failure in Load due to Docker memory — not code	Excellent
AI — Response Time	Tied to LLM inference — normal for AI-backed systems	Acceptable
Conflict Detection	Functional at 2–3 s median under low concurrency	Excellent
Security RBAC	All unauthorised access attempts blocked with 401/403	Excellent
Stress Resilience	No crashes, no dropped connections — queuing only	Excellent

8.5 Quality Metrics:

8.5.1 Test Coverage:

Table 31: Quality Metrics

Service	Total Endpoints	Endpoints Tested	Coverage
UPM Service	21	21	1
AI Service (incl. Conflict Det.)	44	44	1
Notification Service	3	3	1
Total	68	68	1

8.5.2 Defect Density:

Phase	Total Requests	Failures	Failure Rate	Root Cause
Performance (UPM + AI)	668	0	0	None
Load (UPM + AI)	6490	1	0	Docker memory limit (infrastructure)
Stress (UPM + AI)	20169	0	0	None
Overall Total	27327	1	0	Local infrastructure only

8.5.3 Response Time Metrics:

Metric	UPM — Performance	UPM — Load	UPM — Stress	AI — Performance
Median (P50)	57 ms	140 ms	490 ms	3,100 ms
P95	200 ms	1,600 ms	4,200 ms	27,000 ms
P99	430 ms	4,600 ms	7,000 ms	41,000 ms
Max	2,496 ms	8,771 ms	11,808 ms	46,133 ms
Average	82 ms	393 ms	1,059 ms	6,896 ms

8.5.4 Requests Per Second (RPS):

Service	Performance RPS	Load RPS	Stress RPS
UPM Service	4.61	20.31	32.3
AI Service	1.1	1.42	1.35

Note: AI RPS is lower due to LLM processing overhead — expected behaviour for AI-backed systems.

8.6 Usability Testing:

Objective: Verify that the system is usable and comprehensible by real users.

Method: Questionnaire + supervisor review session + manual functional testing.

8.6.1 Manual Functional Test Scenarios:

Table 32: Usability Testing

Scenario	User Role	Task	Outcome
Register new account	Developer	Create account and log in	Completed in < 30 s
Create project and upload code	Developer	Create project and upload .py files	Completed successfully
Request conflict detection	Developer	Compare two project versions	Result returned in 3–5 s locally
Review conflict report	Reviewer	Read report and rate explanation	Clear and comprehensible
Manage users	Admin	Activate / deactivate user account	Completed with single action
View system statistics	Admin	Review reviewer stats and evaluation data	All data displayed correctly

8.6.2 Strengths and Improvement Suggestions:

Aspect	Observation	Recommendation
API Response Time	Very fast in local environment	Deploy to cloud for better production performance
Conflict Detection	Results are accurate and comprehensive	Add graphical UI in future iteration
Error Messages	Clear and specific (400/401/403/404)	Add Arabic language error messages in future
Documentation	Full API docs via Swagger / OpenAPI	Facilitates integration across development teams
LLM Latency	Longer than standard APIs — expected	Add progress indicator for end users

Chapter 9: Results and Analysis

9.1 System Implementation Results:

9.1.1 Functional Results — Core Features:

All 39 functional requirements defined in Chapter 4 were fully implemented and verified through the test suite documented in Chapter 8. The following summarizes the key functional outcomes:

Table 33: System Implementation Results

Feature Area	Status	Key Outcome
User Registration and Authentication	✓ Complete	JWT-based auth with RBAC; three roles enforced; 100% security test pass rate
Project and File Management	✓ Complete	Full CRUD with nested folder structure; GitHub sync operational
Python AST Analysis	✓ Complete	Classes, functions, metrics, dependencies extracted accurately
Java Analysis (Tree-sitter)	✓ Complete	Basic class and method extraction; less deep than Python AST
UML Class Diagram Generation	✓ Complete	Mermaid.js diagrams generated for all test codebases
High-Level Documentation (HighLevelAgent)	✓ Complete	Architectural overviews generated with Senior Architect persona
Detailed Documentation (LowLevelAgent)	✓ Complete	Method-level technical analysis generated
Documentation Verification (VerifierAgent)	✓ Complete	73% reduction in technical errors vs single-LLM baseline
Conflict Detection (5-stage pipeline)	✓ Complete	10 endpoint types; 2.4–3.3 s median under local infrastructure
Documentation Export (PDF/MD/HTML/XML)	✓ Complete	All four formats generated successfully in manual tests
Email Notifications	✓ Complete	Project creation, analysis completion events trigger emails
LLM Caching (48h TTL)	✓ Complete	Cache hit rate >80% in repeated analysis scenarios
Seven-Model Fallback	✓ Complete	Automatic failover tested with rate-limited OpenRouter accounts

9.1.2 Interface and Usage Results:

The web interface was implemented as a React.js SPA with the following key screens: Login/Registration, Project Dashboard, File Upload and Management, Analysis Results View (with UML diagram rendering), Documentation View (High-Level and Low-Level), Conflict Detection

Report, Export Panel, and Admin User Management. Screenshots of all interface screens are provided in Appendix B-4-2. A representative summary of key UI screens is described below:

- **Project Dashboard:** Displays all user projects with status indicators, creation dates, and quick-action buttons for analysis and conflict detection. Projects are paginated with 10 items per page.
- **Analysis Results View:** Presents extracted metrics (lines of code, class count, function count, complexity score) in a structured card layout alongside the interactive Mermaid.js UML diagram.
- **Documentation View:** Displays the HighLevelAgent and LowLevelAgent outputs in a tabbed interface with markdown rendering and an inline rating widget for Reviewers.
- **Conflict Detection Report:** Presents the conflict pipeline results categorized by severity (Critical / Warning / Informational) with expandable detail panels for each detected conflict.

9.2 Comparative Analysis with Existing Systems:

Metric	AutoTest & DocGen Mgr	DocuWriter.ai	Qodo.ai	Mintlify
Documentation error rate	~5% (multi-agent verified)	~15–20% (single-LLM)	N/A (test-only)	~10–15% (estimated)
UML diagram support	✓ Full class diagrams	✗ None	✗ None	✗ None
Conflict detection	✓ 5-stage pipeline	✗	△ Test failures only	✗
Cost optimization	✓ 80% reduction via caching	✗ No caching	△ Limited	✗ No caching
Multi-agent verification	✓ VerifierAgent	✗	✗	✗
API response (standard ops)	36–200 ms median	N/A (SaaS)	N/A (IDE plugin)	N/A (SaaS)
Open academic access	✓ Full source available	✗ Commercial SaaS	✗ Commercial	✗ Commercial

9.3 Objectives Achievement Verification:

Chapter 1 defined six SMART objectives. The following table verifies achievement against each:

Table 34: Objectives Achievement Verification

Objective ID	Objective Description	Target Metric	Achieved	Evidence
PO1	Structural code analysis with class/function extraction	≥ 95% structural accuracy	✓ Yes	Manual verification on 10 test codebases; all structures captured
PO2	Documentation generation	Measurable reduction in	✓ Yes	Full pipeline from upload to

	(Swagger, UML, natural language)	manual effort		documented output in < 5 minutes
SO1	Multi-agent architecture error reduction	$\geq 70\%$ error reduction vs single-LLM	✓ Yes — 73%	Comparative test: 15% single-LLM error rate vs ~5% multi-agent
SO2	Topological code processing for large codebases	Successful analysis beyond context window	✓ Yes	Project-level analysis with dependency-ordered chunking implemented
SO3	Interactive dashboard with $\geq 90\%$ user satisfaction	< 3 s response; $\geq 90\%$ satisfaction	✓ Yes	UPM median 36–490 ms; usability testing shows task completion < 2 min
SO4	CI/CD integration (GitHub trigger)	100% trigger success rate	✓ Partial	GitHub sync implemented; automated trigger on commit pending full CI integration
SO5	Caching reduces AI costs by 80%	$\geq 80\%$ reduction in API calls	✓ Yes	48-hour TTL; cache hit rate measured at > 80% in repeated scenarios
SO6	Security — zero critical vulnerabilities	0 critical vulnerabilities	✓ Yes	All security tests passed; JWT/RBAC fully enforced (Chapter 8 results)

9.4 Strengths and Weaknesses Analysis:

9.4.1 Key Strengths:

- **Multi-Agent Verification:** The HighLevelAgent + LowLevelAgent + VerifierAgent architecture achieves a 73% reduction in documentation errors compared to single-LLM approaches. This is the most significant technical contribution of the project.
- **Zero Failure Resilience:** The UPM service recorded zero failures across 25,975 requests under performance, load, and stress testing conditions. This demonstrates robust backend design suitable for production deployment.
- **Cost-Efficient AI Integration:** The 48-hour caching strategy combined with the seven-model fallback chain enables production-grade LLM integration at near-zero API cost for repeated analysis scenarios.

- **Comprehensive Conflict Detection:** The five-stage conflict detection pipeline provides deeper analysis than any comparable commercial tool, combining structural, semantic, and documentation consistency checks.
- **Clean Microservices Architecture:** The database-per-service pattern with async inter-service communication via Celery ensures high service independence and fault tolerance.

9.4.2 Limitations and Areas for Improvement:

- **LLM Latency Under Concurrent Load:** The AI service shows high median response times (14,000–45,000 ms) under 50 concurrent users, reflecting the inherent sequential nature of LLM inference on local infrastructure. Cloud deployment with GPU-accelerated inference would significantly improve this.
- **Java and JavaScript Support Depth:** While Tree-sitter enables parsing of Java and JavaScript, the feature extraction depth for these languages is less comprehensive than the Python AST implementation. Dedicated processors for Java and JavaScript would improve parity.
- **Frontend Polish:** The React.js frontend is functional but lacks the visual polish and advanced UX features of commercial competitors. Graphical conflict visualization and an interactive UML editor are identified as high-priority future enhancements.
- **CI/CD Integration Completeness:** While GitHub repository synchronization is implemented, the fully automated documentation trigger on commit is partially implemented and requires additional configuration for production use.

9.5 Performance KPIs:

Table 35: Performance KPIs

KPI	Target	Achieved	Status
UPM API error rate (all tests)	0 failures	0 failures (25,975 requests)	✓ Excellent
UPM median response (performance)	< 500 ms	57 ms	✓ Excellent
Conflict detection response (10 users)	< 5 s median	2.4–3.3 s median	✓ Excellent
Test endpoint coverage	≥ 80%	100% (68/68 endpoints)	✓ Excellent
Overall failure rate	< 0.1%	0.004% (1/27,327 requests)	✓ Excellent
Multi-agent error reduction	≥ 70%	0.73	✓ Achieved
Caching cost reduction	≥ 80%	> 80% cache hit rate measured	✓ Achieved
LLM fallback chain uptime	Continuous availability	Zero service interruptions during testing	✓ Achieved

Chapter 10: Conclusion and Recommendations

10.1 Project Summary and Key Achievements:

This thesis has presented the design, implementation, and evaluation of the AutoTest & DocGen Manager — an intelligent, AI-powered web platform that automates software documentation generation and code conflict detection end-to-end. The project successfully delivered a production-quality system that addresses genuine limitations in the current landscape of developer tools.

The most significant technical achievement is the multi-agent documentation architecture, comprising the HighLevelAgent, LowLevelAgent, and VerifierAgent, which achieved a 73% reduction in documentation errors compared to single-LLM baseline approaches. This empirically validates the core hypothesis of the project: that specialized, collaborating AI agents outperform single-model approaches for software documentation tasks.

From an engineering perspective, the system demonstrates excellent operational characteristics: the UPM service achieved zero failures across 25,975 requests under performance, load, and stress testing; the overall system failure rate was 0.004% across all test scenarios; and the intelligent caching strategy reduced LLM API costs by more than 80% in repeated analysis scenarios. These results confirm that the system is not only academically rigorous but also production-viable.

Table 36: Key Achievements

Achievement	Key Metric
Multi-agent error reduction	73% fewer errors vs. single-LLM
UPM service reliability	0 failures in 25,975 requests
AI service reliability	0.004% overall failure rate (27,327 requests)
LLM cost optimization	>80% cache hit rate achieved
Test coverage	100% endpoint coverage (68/68 endpoints)
Conflict detection performance	2.4–3.3 s median (10 concurrent users)
Functional requirements completed	20 of 20 (100%)

10.2 Challenges and Difficulties:

The development of the AutoTest & DocGen Manager encountered several significant challenges that required creative problem-solving and architectural adaptation:

- **LLM Token Limitations:** The most technically challenging problem was analyzing large codebases that exceed the context window of available LLM models. The solution — topological code processing that orders files by dependency relationships and builds context incrementally — required deep understanding of graph algorithms and LLM behavior with long contexts. This challenge ultimately led to one of the project's key innovations.
- **LLM Response Consistency:** Early experiments revealed that single-model LLM responses for documentation generation contained 15–20% error rates — factual inaccuracies, incorrect parameter descriptions, and hallucinated class relationships. The multi-agent architecture with

VerifierAgent was developed specifically to address this challenge, requiring significant additional design effort beyond the initial project scope.

- **Microservices Cross-Service Data Consistency:** Implementing the database-per-service pattern while maintaining referential integrity between UPM (PostgreSQL) and AI service (MongoDB) required careful design of the ArtifactRepository cross-service client, including timeout handling, failure recovery, and two-phase commit-equivalent patterns.
- **Celery and Redis Configuration for LLM Tasks:** Tuning Celery worker concurrency, task routing, and result backend configuration for long-running LLM tasks (which can take 30–90 seconds) required extensive experimentation with timeout settings, visibility windows, and priority queues.
- **Time Constraints on a Four-Person Team:** Delivering a system of this complexity within a 14-week timeline with four part-time developers required aggressive prioritization using MoSCoW methodology and continuous scope management throughout the project.

10.3 Future Development Proposals:

The following enhancements are proposed for future iterations of the AutoTest & DocGen Manager, ordered by estimated impact:

- **Cloud Deployment with GPU Inference:** Deploying the AI service on a cloud provider (AWS, GCP, or Azure) with GPU-accelerated LLM inference would reduce conflict detection response times from 3+ seconds to under 500 ms, enabling real-time documentation feedback during coding sessions.
- **Expanded Language Support:** Adding full-depth support for JavaScript (currently Tree-sitter only) and new languages (TypeScript, Go, Rust) through dedicated ILanguageProcessorStrategy implementations. TypeScript support is particularly high-priority given its widespread adoption in modern web development.
- **Graphical Conflict Visualization:** Replacing the text-based conflict report with an interactive graph visualization (using D3.js or Cytoscape.js) that shows the structural diff between two code versions as an annotated AST diff tree, with color-coded severity indicators.
- **Continuous Documentation Mode:** Completing the CI/CD integration to automatically trigger documentation regeneration and conflict detection on every Git commit, implementing a documentation freshness score that quantifies how synchronized the documentation is with the current codebase.
- **Interactive UML Editor:** Adding a drag-and-drop UML diagram editor built on the generated Mermaid.js output, allowing developers to annotate, reposition, and export UML diagrams in standard formats (PNG, SVG, PlantUML).

- **Edge AI Integration:** Investigating the deployment of smaller, quantized LLMs (Phi-3, Llama 3.1 8B) on local developer machines to provide privacy-preserving documentation generation without any external API calls, critical for proprietary codebases in regulated industries.
- **ML Documentation Specialization:** Extending the documentation generation agents with domain-specific prompts for ML/AI code, capable of documenting model architectures, training pipelines, dataset preprocessing steps, and hyperparameter configurations — a rapidly growing documentation need in AI-driven development teams.

10.4 Recommendations:

10.4.1 Recommendations for Future Student Projects:

- **Architect for Verifiability from Day One:** The most impactful design decision in this project was the VerifierAgent architecture. Future projects using LLMs for content generation should design verification mechanisms before implementation, not after observing quality issues.
- **Use Asynchronous Processing for All LLM Calls:** Any system that integrates LLM APIs must use asynchronous task queues (Celery, RQ, or similar) from the start. Synchronous LLM calls will inevitably cause API timeouts and poor user experience as request volume scales.
- **Implement Caching Early:** LLM API costs accumulate rapidly during development. Implementing response caching in the first sprint, not the last, saves both money and development time by avoiding repeated identical API calls during testing and debugging.
- **Test with Real-Scale Data:** Performance characteristics of AI-powered systems differ dramatically between toy test cases and realistic codebases. Testing with actual project codebases from early sprints reveals bottlenecks — particularly token limits and database query performance — that synthetic tests miss.

10.4.2 Recommendations for Organizations Adopting the Platform:

- **Deploy with Cloud LLM Inference:** For production use, replace the OpenRouter free-tier fallback chain with a dedicated cloud LLM endpoint (OpenAI GPT-4o API, Anthropic Claude API, or a self-hosted model on GPU infrastructure) to achieve consistent response times below 2 seconds for documentation generation.
- **Integrate with Existing CI/CD Pipelines:** Connect the conflict detection API to your existing GitHub Actions or Jenkins pipelines so that documentation freshness is automatically checked on every pull request, preventing documentation drift before it accumulates.
- **Assign Dedicated Reviewer Role:** The Reviewer rating system is most effective when a dedicated team member (technical writer, senior developer, or architect) consistently reviews

and rates generated documentation, providing feedback that can be used to tune agent prompts over time.

The AutoTest & DocGen Manager demonstrates that intelligent, multi-agent AI architectures can deliver measurably superior outcomes in software documentation tasks compared to single-model approaches, while maintaining the reliability and performance standards required for production deployment. The project contributes a validated architecture, a novel cost-optimization strategy, and an open academic implementation that future researchers and practitioners can build upon to advance the state of AI-assisted software engineering.

References

All references are formatted according to APA 7th edition. Website references were last accessed in 2024 unless otherwise noted.

Backend Framework

- Django Software Foundation. (2024). *Django documentation*. <https://docs.djangoproject.com>
- Christie, T. (2024). *Django REST Framework documentation*. <https://www.django-rest-framework.org>

JWT Authentication

- Padilla, J. (2024). *Django REST Framework Simple JWT*. <https://django-rest-framework-simplejwt.readthedocs.io>

Databases

- MongoDB, Inc. (2024). *MongoDB manual*. <https://www.mongodb.com/docs/manual>
- Oracle Corporation. (2024). *MySQL 8.0 reference manual*. <https://dev.mysql.com/doc>

Asynchronous Processing

- Celery Project. (2024). *Celery: Distributed task queue*. <https://docs.celeryq.dev>
- Redis Ltd. (2024). *Redis documentation*. <https://redis.io/docs>

Code Analysis

- Tree-sitter. (2024). *Tree-sitter: An incremental parsing library*. <https://tree-sitter.github.io/tree-sitter>

Containerization

- Docker Inc. (2024). *Docker documentation*. <https://docs.docker.com>
- Docker Inc. (2024). *Docker Compose documentation*. <https://docs.docker.com/compose>

Performance Testing

- Locust Contributors. (2024). *Locust documentation: An open source load testing tool*. <https://docs.locust.io>
- Molyneaux, I. (2014). *The art of application performance testing* (2nd ed.). O'Reilly Media.

Artificial Intelligence & Natural Language Processing

- OpenRouter. (2024). *OpenRouter API documentation*. <https://openrouter.ai/docs>
- Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., & Amodei, D. (2020). Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33, 1877–1901.

Microservices Architecture

- Newman, S. (2021). *Building microservices: Designing fine-grained systems* (2nd ed.). O'Reilly Media.
- Richardson, C. (2018). *Microservices patterns*. Manning Publications.

Frontend Development

- Meta Open Source. (2024). *React documentation*. <https://react.dev>
- Abramov, D., & React Team. (2024). *React hooks reference*. <https://react.dev/reference/react>
- Axios Contributors. (2024). *Axios documentation*. <https://axios-http.com/docs/intro>
- React Router Contributors. (2024). *React Router documentation*. <https://reactrouter.com/en/main>
- Redux Toolkit Contributors. (2024). *Redux Toolkit documentation*. <https://redux-toolkit.js.org>
- Tailwind CSS. (2024). *Tailwind CSS documentation*. <https://tailwindcss.com/docs>
- Vite. (2024). *Vite documentation*. <https://vitejs.dev/guide>